

## Assignment 3

Due Monday, March 27 at 11:59pm

**General Instructions**

- Feel free to talk to other members of the class in doing the homework. You should, however, write down your solutions yourself. *List the names of everyone you worked with at the top of your submission.*
- Keep your solutions brief and clear.
- Please use Piazza if you have questions about the homework but do not post answers. Feel free to use private posts or come to the office hours.

**Homework Submission**

- We DO NOT accept late homework submissions.
- We will be using Compass for collecting the homework assignments. Please submit your answers via [Compass](#). Hard copies are not accepted.
- Contact the TAs if you are having technical difficulties in submitting the assignment; attempt to submit well in advance of the due date/time.
- The homework must be submitted in **pdf** format. Scanned handwritten and/or hand-drawn pictures in your documents won't be accepted.
- Please do not zip the answer document (PDF) so that the graders can read it directly on Compass. You need to submit one answer document, named as **hw3\_netid.pdf**.
- Please see the [assignments](#) page for more details. In particular, we will be announcing errata, if any, on this page.

## 1 SQL queries involving a single relation

Consider a relation with just one attribute,  $R(A)$ . Express the following queries using SQL. Your queries will be graded based on their simplicity, and correctness.

1. [10] Write a query to compute the largest value of  $A$  in  $R$ . You are not allowed to use the **Max** operator. For this question, you may assume that there are no duplicates in  $R$ .
2. [10] Write a query to compute the median value of  $A$  in  $R$ . For this question, you may assume that there are no duplicates in  $R$ , and that  $R$  has odd number of tuples.
3. [10] Write a query to compute the mode of  $A$  in  $R$ . If multiple distinct values of  $A$  are candidates, then all of them should be returned. For this question,  $R$  may contain duplicates tuples; please don't assume otherwise.

## 2 SQL queries involving multiple relations

Consider the following relations:

```
Courses(cid, cname, room, pid)
Professors(pid, pname)
Students(sid, sname)
Enrolled(sid, cid)
```

Express the following queries using SQL. Your queries will be graded based on their simplicity, and correctness.

1. [5] Find the names of students who are not enrolled in any class.
2. [10] Find the names of students who are enrolled in the highest number of classes, i.e., all other students (not outputted) have taken fewer classes.
3. [10] Find the names of professors who teach in every room in which some course is taught.

## 3 Database Manipulation and Views

Consider the relations **Courses**, **Professors**, **Students**, **Enrolled** as described in question 2.

1. [5] Insert the following new tuple in **Students**: John with **sid** as 22222.
2. [5] Remove all professors from **Professors**, who are not teaching any course.
3. [5] Update the name of the student with **sid** as 22222, to Doe.

4. [10] Create a materialized view **BusyProfessors**, containing the names of all professors for whom the number of students they teach across their courses is greater than hundred.

Note: if a student is enrolled in two different courses of the same professor, we count the student only once.

## 4 Constraints and Triggers

Consider the relations **Courses**, **Professors**, **Students**, **Enrolled** as described in question 2.

1. [10] Create the tables for **Courses**, **Professors**, **Students**, **Enrolled** using the SQL **CREATE** statement, with the following constraints:
  - For **Enrolled**, all the student and course ids should be present in the corresponding tables
  - For **Courses**, the professor id must be present in the corresponding table
2. [5] A deletion or update to a **sid** should be handled appropriately in **Enrolled**, due to the constraint as explained in the previous question.
  - (a) Please rewrite the SQL **CREATE** statement such that when an **sid** gets updated/deleted in **Students**, its value is propagated accordingly to **Enrolled**.
  - (b) Also, explain why we can not set **Enrolled.sid** to **NULL** when we delete that **sid** from **Students**.
3. [5] Create a trigger so that an insertion into **Courses** creates an tuple in **Professors** if the **pid** corresponding to the insertion doesn't exist in **Professors**.

## 5 Closure of Functional Dependencies

Consider a relation  $R(A_1, A_2, \dots, A_n)$ , with  $n \geq 3$ .  $R$  satisfies the following functional dependency:  $A_1 \rightarrow A_2$ .

1. [0] How many logically implied functional dependencies are of the form,  $P \rightarrow A_2 A_n$ , where  $P \subseteq \{A_1, A_2, \dots, A_n\}$ .<sup>1</sup>

Note: attempt this problem at your leisure, **do not submit solutions**.

---

<sup>1</sup>based on a [question asked on Piazza](#).