



DNS

[Host Names vs. IP addresses]

■ Host names

- Mnemonic name appreciated by **humans**
- Variable length, full alphabet of characters
- Provide little (if any) information about physical location
- Examples: **www.cnn.com** and **bbc.co.uk**

■ IP addresses

- Numerical address appreciated by **routers**
- Fixed length, binary number
- Hierarchical, related to host location
- Examples: **64.236.16.20** and **212.58.224.131**



[Separating Naming and Addressing]

- Names are easier to remember
 - **cnn.com** vs. **64.236.16.20** (but not shortened urls)
- Addresses can change underneath
 - Move **www.cnn.com** to **4.125.91.21**
 - e.g., renumbering when changing providers



Separating Naming and Addressing

- Name could map to multiple IP addresses
 - www.cnn.com may refer to multiple (8) replicas of the Web site
 - Enables
 - Load-balancing
 - Reducing latency by picking nearby servers
 - Tailoring content based on requester's location/identity
- Multiple names for the same address
 - e.g., aliases like www.cnn.com and cnn.com



Scalable (Name \leftrightarrow Address) Mappings

- Originally: per-host file
 - Flat namespace
 - `/etc/hosts`
 - SRI (Menlo Park) kept master copy
 - Downloaded regularly



Scalable (Name ↔ Address) Mappings

- Why not centralize DNS?
 - Single point of failure
 - Traffic volume
 - Distant centralized database
 - Maintenance
- Doesn't scale!
- Root name server
 - Contacted by local name server that can not resolve name
 - Contacts authoritative name server if mapping not known
 - Gets mapping and returns it to local name server



[Domain Name Service (DNS)]

- Large scale dynamic, distributed application
 - Replaced Network Information Center (NIC)
- RFC 1034 and 1035
- Name space
 - Set of possible names
- Bindings
 - Maps internet domain names into IP addresses
- Name server
 - Resolution mechanism



[Applications' use of DNS]

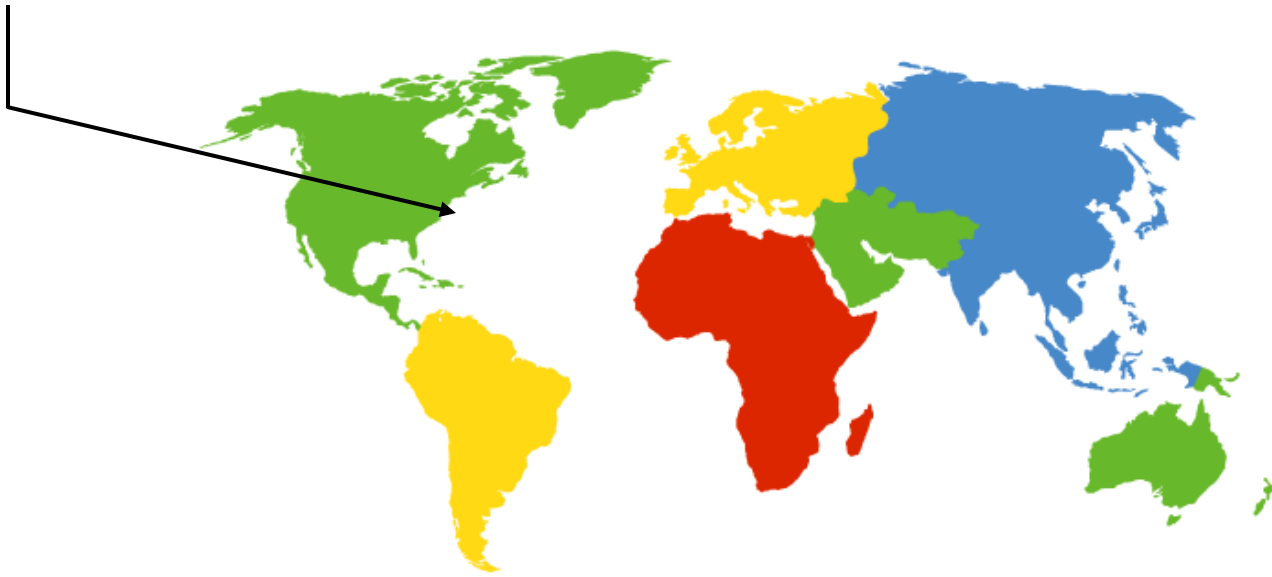
- Local DNS server (“default name server”)
 - Usually near the endhosts that use it
 - Local hosts configured with local server (e.g., `/etc/resolv.conf`) or learn server via DHCP
- Client application
 - Extract server name (e.g., from the URL)
 - Do `getaddrinfo()` to trigger resolver code, sending message to server
- Server application
 - Extract client IP address from socket
 - Optional `getnameinfo()` to translate into name



[DNS Root]

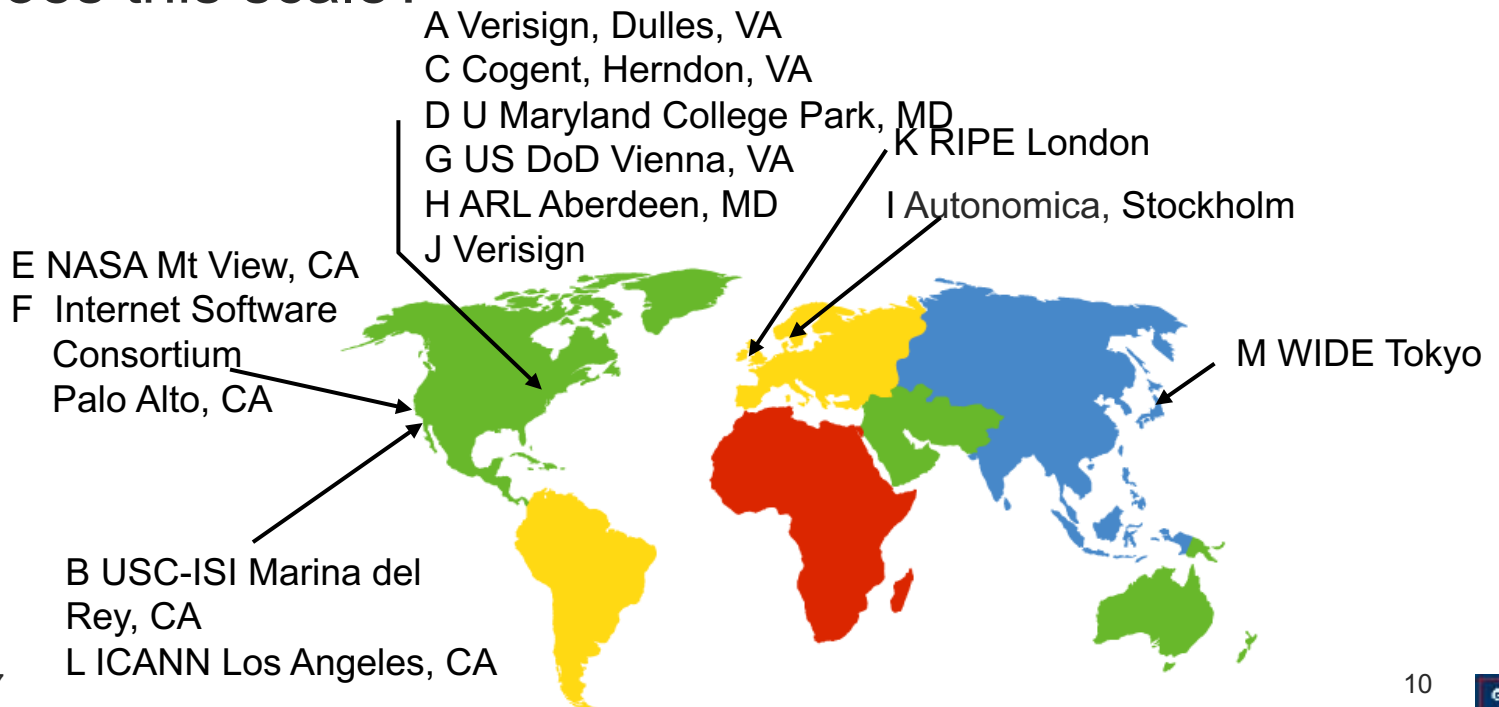
- Located in Virginia, USA
- How do we make the root scale?

Verisign, Dulles, VA



DNS Root Servers

- 13 root servers (see <http://www.root-servers.org/>)
 - Labeled A through M
- Does this scale?



[TLD and Authoritative Servers]

- Top-level domain (TLD) servers
 - Responsible for **com**, **org**, **net**, **edu**, etc, and all top-level country domains **uk**, **fr**, **ca**, **jp**.
 - Network Solutions maintains servers for **com** TLD
 - Educause for **edu** TLD
- Authoritative DNS servers
 - Organization's DNS servers
 - Provide authoritative hostname to IP mappings for organization's servers (e.g., Web, mail).
 - Can be maintained by organization or service provider




[Local Name Server]

- One per ISP (residential ISP, company, university)
 - Also called “default name server”
- When host makes DNS query, query is sent to its local DNS server
 - Acts as proxy, forwards query into hierarchy
 - Reduces lookup latency for commonly searched hostnames



Distributed, Hierarchical Database



- Client wants IP for **www.amazon.com**
 - Client queries a root server to find **com** DNS server
 - Client queries **com** DNS server to get **amazon.com** DNS server
 - Client queries **amazon.com** DNS server to get IP address for **www.amazon.com**



DNS – Name Server

- Host at **cs.uiuc.edu**


- Wants IP address for **gaia.cs.umass.edu**

- Recursive query

- Ask server to get answer for you
- e.g., request 1 and response 8

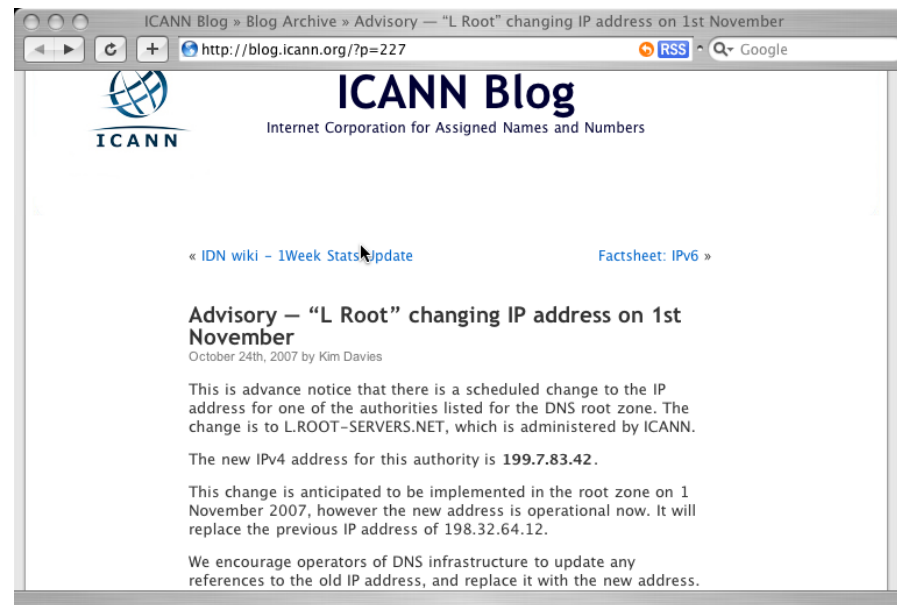
- Iterated query

- Contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”



But how did it know the root server IP?

- Hard-coded
- What if it changes?



[DNS: Caching]

- Performing all these queries takes time
 - And all this before actual communication takes place
 - e.g., 1-second latency before starting Web download
- Caching can greatly reduce overhead
 - The top-level servers very rarely change
 - Popular sites (e.g., www.cnn.com) visited often
 - Local DNS server often has the information cached



[DNS: Caching]

- How DNS caching works
 - DNS servers cache responses to queries
 - Responses include a “time to live” (TTL) field
- Once (any) name server learns mapping, it caches mapping
 - Cache entries timeout (disappear) after some time
 - TLD servers typically cached in local name servers
 - Thus root name servers not often visited



DNS Resource Records

DNS: distributed DB storing resource records (RR)

RR format: (name, value, type, ttl)

- Type=A
 - name is hostname
 - value is IP address
- Type=NS
 - name is domain (e.g. `foo.com`)
 - value is hostname of authoritative name server for this domain
- Type=PTR
 - name is reversed IP quads
 - e.g. `78.56.34.12.in-addr.arpa`
 - value is corresponding hostname
- Type=CNAME
 - name is alias name for some “canonical” name
 - e.g., `www.cs.mit.edu` is really `eeecsweb.mit.edu`
 - value is canonical name
- Type=MX
 - value is name of mailserver associated with name
 - Also includes a weight/preference



DNS Protocol

DNS protocol: *query* and *reply* messages, both with **same message format**

- Message header
- Identification
 - 16 bit # for query, reply to query uses same #
- Flags
 - Query or reply
 - Recursion desired
 - Recursion available
 - Reply is authoritative
- Plus fields indicating size (0 or more) of optional header elements

16 bits	16 bits
Identification	Flags
# Questions	# Answer RRs
# Authority RRs	# Additional RRs
Questions (variable # of resource records)	
Answers (variable # of resource records)	
Authority (variable # of resource records)	
Additional information (variable # of resource records)	



[Reliability]

- DNS servers are replicated
 - Name service available if at least one replica is up
 - Queries can be load-balanced between replicas
- Usually, UDP used for queries
 - Need reliability: must implement this on top of UDP
 - Spec supports TCP too, but not always implemented
- Try alternate servers on timeout
 - Exponential backoff when retrying same server
- Same identifier for all queries
 - Don't care which server responds



Inserting Resource Records into DNS

- Example: just created startup “FooBar”
- Get a block of address space from ISP
 - Say **212.44.9.128/25**
- Register foobar.com at Network Solutions (say)
 - Provide registrar with names and IP addresses of your authoritative name server (primary and secondary)
 - Registrar inserts RR pairs into the **com** TLD server:
 - **(foobar.com, dns1.foobar.com, NS)**
 - **(dns1.foobar.com, 212.44.9.129, A)**
- Put in your (authoritative) server **dns1.foobar.com**:
 - Type A record for **www.foobar.com**
 - Type MX record for **foobar.com**



[Setting up foobar.com]

- In addition, need to provide reverse PTR bindings
 - e.g., `212.44.9.129` → `dns1.foobar.com`
- Normally, these go in `9.44.212.in-addr.arpa`
- Problem
 - You can't run the name server for that domain. Why not?
 - Because your block is `212.44.9.128/25`, not `212.44.9.0/24`
 - Whoever has `212.44.9.0/25` won't be happy with you owning their PTR records
- Solution: ISP runs it for you
 - Now it's more of a headache to keep it up-to-date :-(



DNS Measurements (MIT data from 2000)

- What is being looked up?
 - ~60% requests for A records
 - ~25% for PTR records
 - ~5% for MX records
 - ~6% for ANY records
- How long does it take?
 - Median ~100msec (but 90th percentile ~500msec)
 - 80% have no referrals; 99.9% have fewer than four
- Query packets per lookup: ~2.4



DNS Measurements (MIT data from 2000)

- Top 10% of names accounted for ~70% of lookups
 - Caching should really help!
- 9% of lookups are unique
 - Cache hit rate can never exceed 91%
- Cache hit rates ~ 75%
 - But caching for more than 10 hosts doesn't add much



DNS Measurements (MIT data from 2000)

- Does DNS give answers?
 - ~23% of lookups fail to elicit an answer!
 - ~13% of lookups result in NXDOMAIN (or similar)
 - Mostly reverse lookups
 - Only ~64% of queries are successful!
 - How come the web seems to work so well?
- ~ 63% of DNS packets in unanswered queries!
 - Failing queries are frequently retransmitted
 - 99.9% successful queries have ≤ 2 retransmissions



[Moral of the Story]

- If you design a highly resilient system, many things can be going wrong without you noticing it!



Security Analysis of DNS

- What security issues does the design & operation of the Domain Name System raise?



16 bits	16 bits
Identification	Flags
# Questions	# Answer RRs
# Authority RRs	# Additional RRs
Questions (variable # of resource records)	
Answers (variable # of resource records)	
Authority (variable # of resource records)	
Additional information (variable # of resource records)	



Security Problem #1: Starbucks (and China...)

- As you sip your latte and surf the Web, how does your laptop find **google.com**?
- Answer: it asks the local name server per Dynamic Host Configuration Protocol (DHCP) ...
 - ... which is run by Starbucks or their contractor
 - ... and can return to you **any answer they please**
 - ... including a “man in the middle” site that forwards your query to Google, gets the reply to forward back to you, yet can **change anything** they wish in **either** direction
- How can you know you’re getting correct data?
 - Today, you can’ t. (Though if site is HTTPS, that helps)
 - One day soon: **DNSSEC** extensions to DNS



Security Problem #2: Cache Poisoning

- Suppose you are a Bad Guy and you control the name server for **foobar.com**. You receive a request to resolve **www.foobar.com** and reply:

```
;; QUESTION SECTION:
;www.foobar.com.                IN      A

;; ANSWER SECTION:
www.foobar.com.                300     IN      A      212.44.9.144

;; AUTHORITY SECTION:
foobar.com.                    600     IN      NS      dns1.foobar.com.
foobar.com.                    600     IN      NS      google.com.

;; ADDITIONAL SECTION:
google.com.                    5       IN      A      212.44.9.155
```

Evidence of the attack disappears 5 seconds later!



[Cache Poisoning]

- Okay, but how do you get the victim to look up `www.foobar.com` in the first place?
- Perhaps you connect to their mail server and send
 - `HELO www.foobar.com`
 - Which their mail server then looks up to see if it corresponds to your source address (anti-spam measure)
- Note, with compromised name server we can also lie about PTR records (address → name mapping)
 - e.g., for `212.44.9.155 = 155.44.9.212.in-addr.arpa` return `google.com` (or `whitehouse.gov`, or `whatever`)
 - If our ISP lets us manage those records as we see fit, or we happen to directly manage them



[Cache Poisoning]

- Suppose Bad Guy is at Starbucks and they can **sniff** (or even **guess**) the identification field the **local server** will use in its next request.
- They:
 - Ask local server for a (recursive) lookup of **google.com**
 - Locally **spoof** subsequent reply from correct name server using the identification field
 - Bogus reply arrives **sooner** than legit one
- Local server duly caches the bogus reply!
 - Now: **every** future Starbucks customer is served the bogus answer out of the local server's cache
 - In this case, the reply uses a **large** TTL



[Summary]

- Domain Name System (DNS)
 - Distributed, hierarchical database
 - Distributed collection of servers
 - Caching to improve performance
- DNS currently lacks authentication
 - Can't tell if reply comes from the correct source
 - Can't tell if correct source tells the truth
 - Malicious source can insert extra (mis)information
 - Malicious bystander can spoof (mis)information
 - Playing with caching lifetimes adds extra power to attacks

