





# Reliable Transmission

# [ Reliable Transmission ]



# [ Reliable Transmission ]



# [ Reliable Transmission ]

- Suppose error protection identifies valid and invalid packets
  - How?
- Can we make the channel appear reliable?
  - Insure packet delivery
  - Maintain packet order
  - Provide reliability at full link capacity



# [ Reliable Transmission Outline ]

- Fundamentals of Automatic Repeat reQuest (**ARQ**) algorithms
  - A family of algorithms that provide reliability through retransmission
- ARQ algorithms (simple to complex)
  - stop-and-wait
  - concurrent logical channels
  - sliding window
    - go-back-n
    - selective repeat
- Alternative: forward error correction (**FEC**)



# [Terminology]

- Acknowledgement (**ACK**)
  - Receiver tells the sender when a frame is received
    - Selective acknowledgement (**SACK**)
      - Specifies set of frames received
    - Cumulative acknowledgement (**ACK**)
      - Have received specified frame and all previous
    - Negative acknowledgement (**NAK**)
      - Receiver refuses to accept frame now, *e.g.*, when out of buffer space



# [ Terminology ]

- Timeout (TO)
  - Sender decides the frame (or ACK) was lost
  - Sender can try again



# [ Stop-and-Wait ]


- Basic idea

1. Send a frame
2. Wait for an ACK or TO
3. If TO, go to 1
4. If ACK, get new frame, go to 1






# Stop-and-Wait: Success






# [ Stop-and-Wait: Lost Frame ]

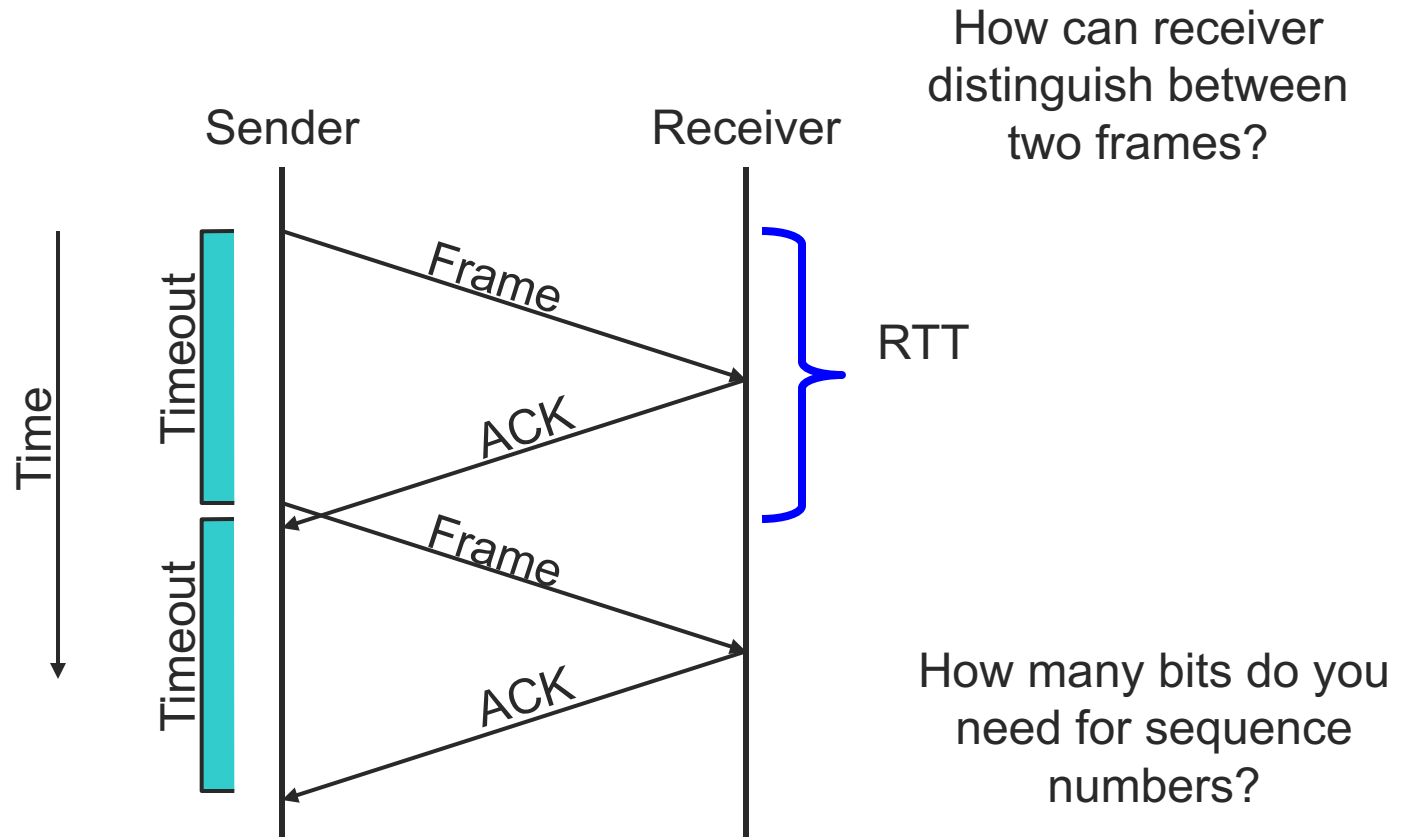




# [ Stop-and-Wait: Lost ACK ]



# Stop-and-Wait: Delayed Frame



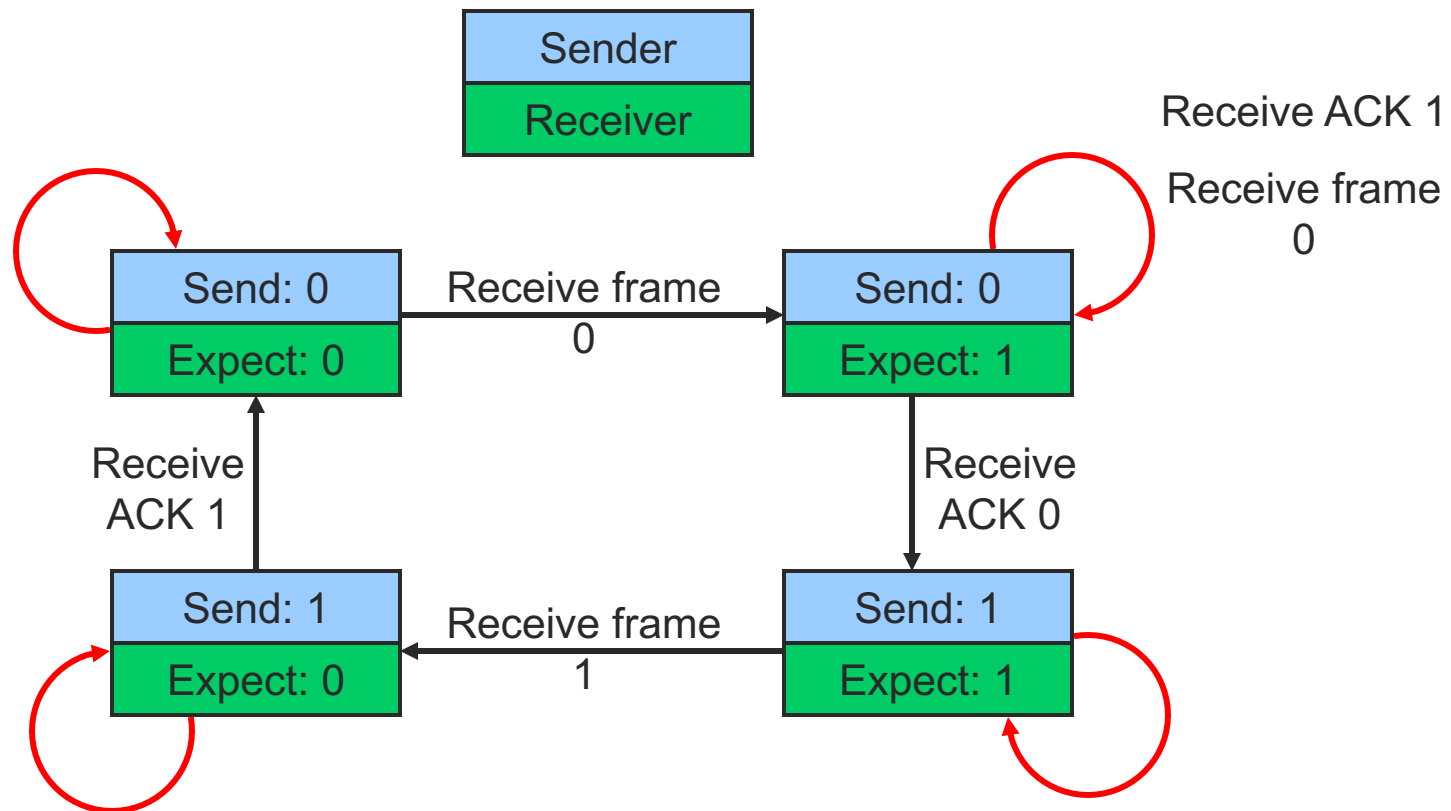


# [ Stop-and-Wait ]

- Goal
  - Guaranteed, at-most-once delivery
- Protocol Challenges
  - Dropped frame/ACK
  - Duplicate frame/ACK
- Requirements
  - 1-bit sequence numbers (if physical network maintains order)
    - sender tracks frame ID to send
    - receiver tracks next frame ID expected



# Stop-and-Wait State Diagram



# [ Stop-and-Wait ]

- We have achieved
  - Frames delivered reliably and in order
  - Is that enough?
- Problem
  - Only allows one outstanding frame
    - Does not keep the pipe full
  - Example
    - 100ms RTT
    - One frame per RTT = 1KB
    - $1024 \times 8 \times 10 = 81920$  kbps
    - Regardless of link bandwidth!



# [ Concurrent Logical Channels ]

- Used in ARPANET IMP-IMP protocol
- Idea
  - Multiplex logical channels over a physical link
    - Include channel ID in header
  - Use stop-and-wait for each channel
- Result
  - Each channel is limited to stop-and-wait bandwidth
  - Aggregate bandwidth uses full physical channel
  - Supports multiple communicating processes
  - Can use more than one channel per process





# [ Concurrent Logical Channels ]

- Problem

- Bandwidth

- Use of a single channel per process may waste BW

- Ordering

- Use of multiple channel per process does not maintain packet ordering across channels!
    - If application has  $n$  channels, and one needs a retransmission, it will always be one packet behind the other channels





# [ ARQ: Where are We? ]

- Goals for reliable transmission
  - Make channel appear reliable
  - Maintain packet order (usually)
  - Impose low overhead/allow full use of link
- Stop-and-Wait
  - Provides reliable in-order delivery
  - Sacrifices performance
- Concurrent Logical Channels
  - Provides reliable delivery at full link bandwidth
  - Sacrifices packet ordering
- Sliding Window Protocol
  - Achieves all three!




# [ Sliding Window Protocol ]

- Most important and general ARQ algorithm
- Used by TCP
- Outline
  - Concepts
  - Terminology (from P&D)
  - Details
  - Code example
  - Proof of eventual in-order delivery
  - Classification scheme
    - (go-back-n, selective repeat)

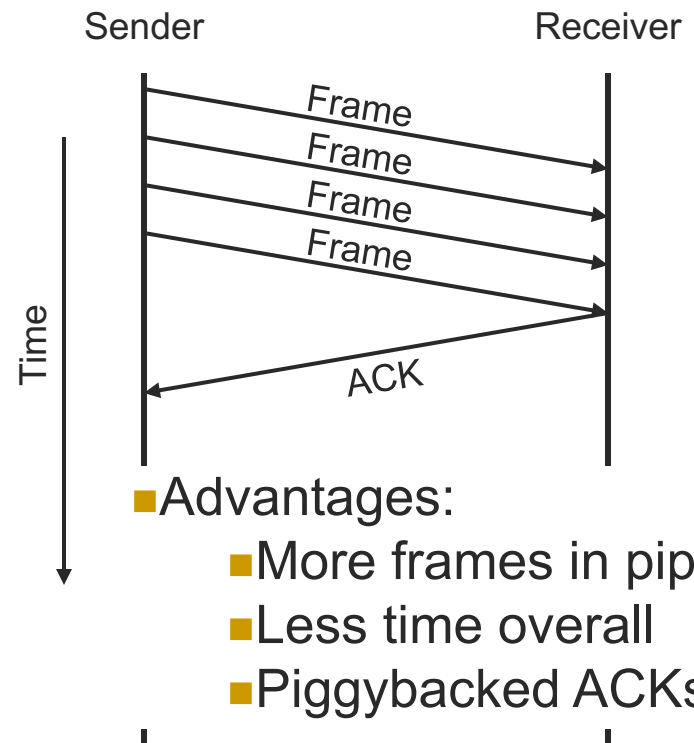


# Keeping the Pipe Full

## Stop-and-Wait



## Goal




### Advantages:

- More frames in pipe
- Less time overall
- Piggybacked ACKs




# [ Concepts ]

- Consider an ordered stream of data frames
- Stop-and-Wait
  - Window of one frame
  - Slides along stream over time



# [ Concepts ]


- Sliding Window Protocol
  - Multiple-frame send window
  - Multiple frame receive window



# [ Sliding Window ]

## ■ Send Window

- Fixed length
- Starts at earliest unacknowledged frame
- Only frames in window are active



# [ Sliding Window ]

- Receive Window
  - Fixed length (unrelated to send window)
  - Starts at earliest frame not received
  - Only frames in window accepted






# Sliding Window Terminology

## ■ Sender Parameters


- Send Window Size (**SWS**)
- Last Acknowledgement Received (**LAR**)
- Last Frame Sent (**LFS**)



# Sliding Window Terminology


## ■ Receiver Parameters

- Receive Window Size (**RWS**)
- Next Frame Expected (**NFE**)
- Last Frame Acceptable (**LFA**)



# Sliding Window Details

- Sender Tasks
  - Assign sequence numbers
  - On ACK Arrival
    - Advance LAR
    - Slide window




# Sliding Window Details

## ■ Receiver Tasks

### ○ On Frame Arrival (N)

- Silently discard if outside of window
  - $N < \text{NFE}$  (NACK possible, too)
  - $N \geq \text{NFE} + \text{RWS}$
- Send cumulative ACK if within window



Receive Frame 4

Send ACK 7




# Sliding Window Details

## ■ Receiver Tasks

### ○ On Frame Arrival (N)

- Silently discard if outside of window
  - $N < \text{NFE}$  (NACK possible, too)
  - $N \geq \text{NFE} + \text{RWS}$
- Send cumulative ACK if within window



# [ Sliding Window Details ]

- Sequence number space
  - Finite number, so wrap around
  - Need space larger than SWS (outstanding frames)
    - In fact, need twice as large
- Example
  - 3-bit sequence numbers (0-7)
  - $RWS = SWS = 7$



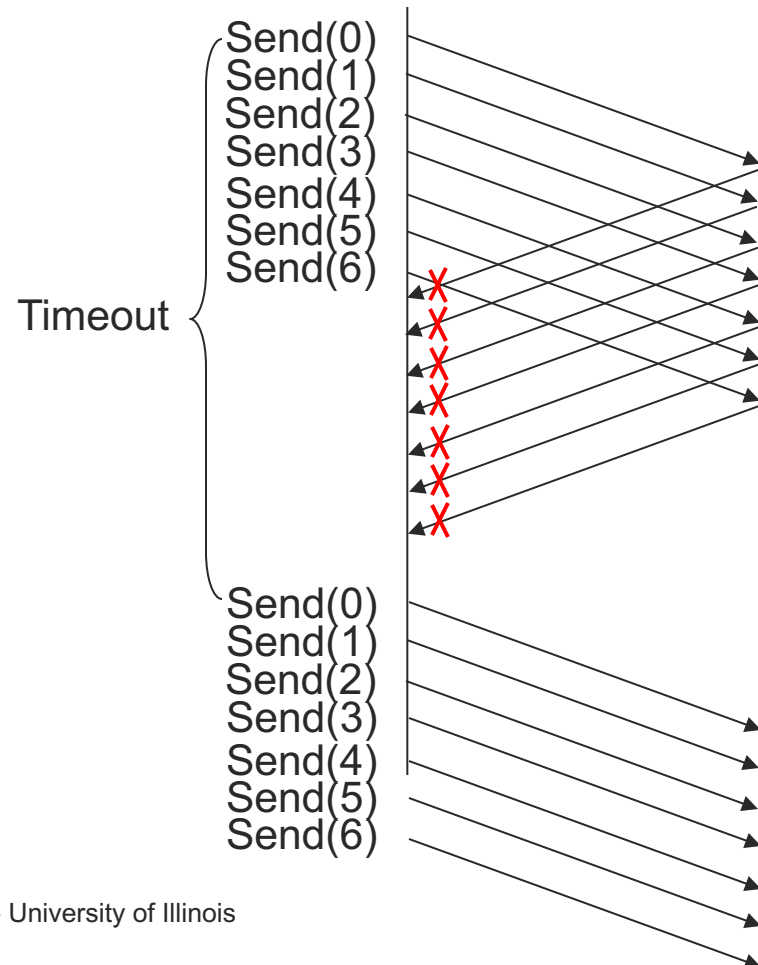
# [ Sliding Window Details ]

- Is  $\log_2(\text{SWS}+1)$  bits enough?
  - No. Example:
  - 3-bit sequence numbers (0-7)
  - $\text{RWS} = \text{SWS} = 7$
  - Why isn't 3 bits enough (can you think of an example where it doesn't work?)



# Sliding Window Details

- Example of incorrect behavior
  - 3-bit sequence numbers 0-7
  - $RWS = SWS = 7$
  - Sender transmits 0-6
  - All arrive, but ACK's lost
  - Sender retransmits
  - Receiver accepts as second incarnation of 0-6





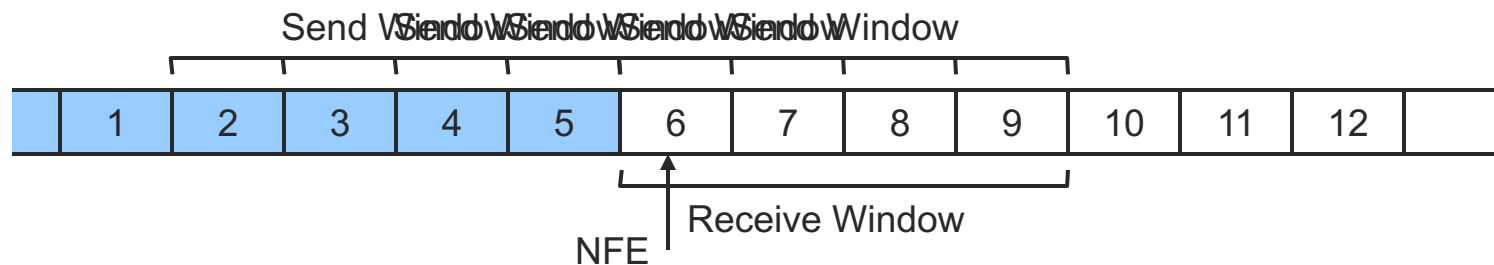
# Sliding Window Sequence Numbers

- How many sequence numbers are necessary?
  - Key questions
    - Where can the send window be?
    - What frame can be received next?



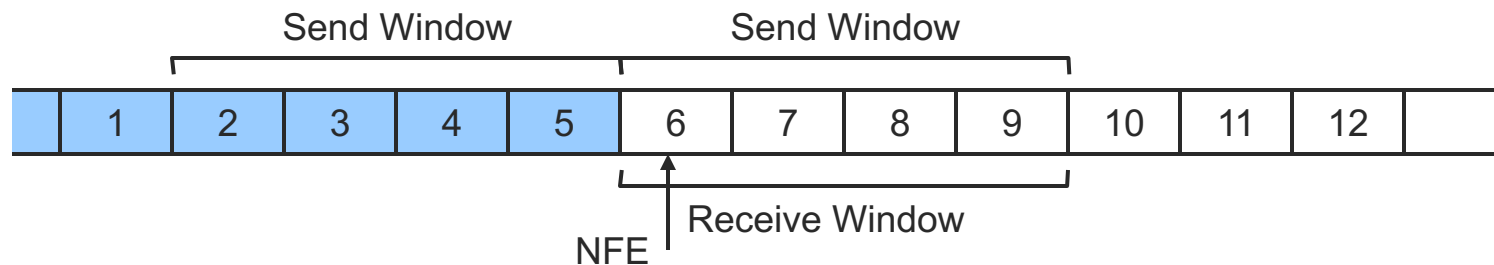
# Sliding Window Sequence Numbers

- Assume **SWS = RWS** (simplest, and typical)
- Sender transmits full SWS
- Two extreme cases:
  - None received (waiting for  **$0 \dots \text{SWS} - 1$** )
  - All received (waiting for  **$\text{SWS} \dots 2 \text{ SWS} - 1$** )
- All possible packets must have unique sequence numbers



# Sliding Window Sequence Numbers

- Extreme Locations for SWS
- Requirements
  - If a received packet is not in the receive window with no wrap, then it must not be in the receive window with wrap!
- Correctness condition:
  - **Number of Sequence Numbers  $\geq$  SWS + RWS**
  - Alternates between two halves of the sequence number space



# Sliding Window Sequence Numbers

## ■ Example

- If  $SWS = RWS = 8$
- At least 16 sequence numbers are needed
- A 4-bit sequence number space is enough

## ■ Warning

- P&D sometimes uses the variable `Max_Seq_Num` for the number of sequence numbers and sometimes for the maximum sequence number (these differ by one!)
- Use `Num_Seq_Num` for the number of sequence numbers:  $0, 1, \dots, \text{Num\_Seq\_Num} - 1$



# [ Window Sizes ]

- How big should we make SWS?
  - Compute from delay x bandwidth
- How big should we make RWS?
  - Depends on buffer capacity of receiver



# Delay x Bandwidth Product - Revisited

- Amount of data in “pipe”
  - channel = pipe
  - delay = length
  - bandwidth = area of a cross section
  - bandwidth x delay product = volume



# [ Delay x Bandwidth Product ]


- Pipe

- Half of data that must be buffered before sender responds to slowdown request



# Delay x Bandwidth Product

- Bandwidth x delay product
  - How many bits the sender must transmit before the first bit arrives at the receiver if the sender keeps the pipe full
  - Takes another one-way latency to receive a response from the receiver





# Sliding Window Protocol Code Example

## ■ Parameters

- last acknowledgement received (**LAR**)
- last frame sent (**LFS**)
- next frame expected (**NFE**)
- last frame acceptable (**LFA**)



# Sliding Window Protocol Code Example

## ■ Constants

- Rend/receive window size (**SWS/RWS**)
- Maximum sequence number (**MAX\_SEQ\_NO**)
- Frame size (**FRAME\_SIZE**, constant for simplicity)



# Sliding Window Protocol Code Example

- Data structures
  - Next frame expected (an integer)
  - One frame buffer for each entry in receive window
  - One presence bit for each entry
- Receive window cycles through
  - Sequence numbers
  - Data structures (thus RWS must divide MAX\_SEQ\_NO)



# Sliding Window Protocol Code Example

```
#define RWS            8          /* receive window size      */
#define MAX_SEQ_NO    16          /* max. sequence number+1 */
                                  /* (must be multiple of    */
                                  /* RWS for this code)     */
#define FRAME_SIZE    1000       /* constant for simplicity*/

char buf[RWS][FRAME_SIZE];      /* RWS frame buffers      */
int present[RWS];               /* are frame buffers full?*/
                                  /* (initialized to 0's)   */
int NFE = 0;                    /* next frame expected    */

extern void send_ack (int seq_no);
extern void pass_to_app (char* data);
void recv_frame (char* data, int seq_no);
```



# Sliding Window Protocol Code Example

```
void recv_frame (char* data, int seq_no)
{
    int idx;           /* index into data structures */
    int i;             /* loop index */

    /* Map sequence numbers NFE...predecessor (NFE)
       into 0...MAX_SEQ_NO - 1, then see if seq_no
       falls within the receive window. */

    if (seq_no - NFE < RWS)

        /* Frames outside the window */
        /* are ignored. (but an ACK */
        /* is sent; why?) */
}
```



# Sliding Window Protocol Code Example

```
/* Calculate index into data structures. */  
idx = (seq_no % RWS);  
  
if (!present[idx]) { /* frame is not dup */  
    present[idx] = 1; /* mark received */  
    memcpy (buf[idx], data, FRAME_SIZE);  
    /* copy data into buf */  
}
```



# Sliding Window Protocol Code Example

```
/* Got a new frame; pass frames up to host? */
for (i = 0; i < RWS; i++) {
    idx = (i + NFE) % RWS;    /* Re-use idx.*/
    /* first missing frame becomes NFE */
    /* after this loop terminates */
    if (!present[idx]) break;

    /* Frame is present—send it up! */
    pass_to_app (buf[idx]);
    present[idx] = 0; /* Mark buffer empty. */
}
/* Advance NFE to first missing frame. */
NFE = NFE + i;
```



# Sliding Window Protocol Code Example

```
        /* Frame handled (might have */  
        /* been duplicate). */  
    } /* (Send ACK for any frame received */  
  
    /* Now send acknowledgement for */  
    /* predecessor (NFE). */  
    send_ack (NFE - 1);  
}
```





# [Correctness]

## ■ Claim

- A sliding window protocol leads to in-order delivery of all frames

## ■ Assumptions

- All sequence numbers are different
- Frames can be lost
- Frames can be delayed an arbitrarily finite amount of time
- Frames are not reordered
- Frames can arrive with detectable errors

## ■ Are these assumption adequate?




# Sliding Window Protocol Correctness

- Need one more assumption
  - Any given frame is received without errors after a finite number of retransmissions
- Proof in two steps
  - Establish correctness assuming infinite sequence number space
  - Show that finite sequence number space does not affect result as long as it has  $\geq 2 \max(\text{SWS}, \text{RWS})$  possible numbers



# Sliding Window Protocol Correctness

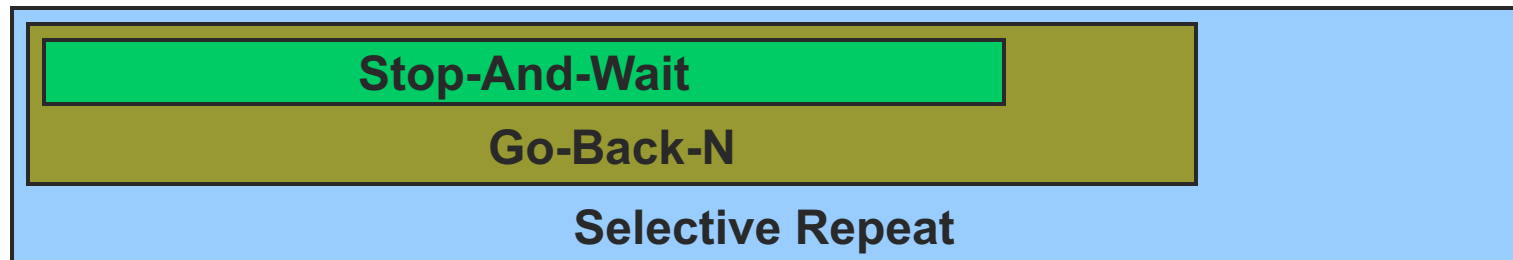
- Step 1: establish correctness assuming infinite sequence number space
  - Use induction on  $k$  with invariant “the  $k^{\text{th}}$  frame is eventually received”
- Step 2: show that finite sequence number space does not affect result as long as it has  $\geq 2 \max(\text{SWS}, \text{RWS})$  possible numbers



# [ ARQ Algorithm Classification ]

- Three Types:

- Stop-and-Wait:       $SWS = 1$        $RWS = 1$
- Go-Back-N:       $SWS = N$        $RWS = 1$
- Selective Repeat:       $SWS = N$        $RWS = M$ 
  - Usually  $M = N$




# Sliding Window Variations: Go-Back-N

- $SWS = N, RWS = 1$
- Receiver only buffers one frame
- If a frame is lost, the sender may need to retransmit up to N frames
  - i.e., sender “goes back” N frames
- Variations
  - How long is the frame timeout?
  - Does receiver send NACK for out-of-sequence frame?



# Go-Back-N: Cumulative ACKs




# Sliding Window Variations: Selective Repeat

- $SWS = N$ ,  $RWS = M$
- Receiver buffer  $M$  frames
- If a frame is lost, sender must only resend
  - Frames lost within the receive window
- Variations
  - How long is the frame timeout?
  - Use cumulative or per-frame ACK?
  - Does protocol adapt timeouts?
  - Does protocol adapt SWS and/or RWS?



# Selective Repeat





# [ Roles of a Sliding Window Protocol ]

- Reliable delivery on an unreliable link
  - Core function
- Preserve delivery order
  - Controlled by the receiver
- Flow control
  - Allow receiver to throttle sender
- Separation of Concerns
  - Must be able to distinguish between different functions that are sometimes rolled into one mechanism



# [ Forward Error Correction (FEC) ]

- Alternative to ARQ algorithms
- Idea
  - Error correction instead of error detection
  - Send extra information to avoid retransmission (i.e., fix errors first/forward rather than afterward/backward)
- Why
  - Very high latency connections
  - Difficult for retransmission

