

CS447: Natural Language Processing

<http://courses.engr.illinois.edu/cs447>

# Lecture 2: Finite-state methods for morphology

Julia Hockenmaier

*[juliahmr@illinois.edu](mailto:juliahmr@illinois.edu)*

3324 Siebel Center

# Last lecture

## The NLP pipeline:

tokenization — POS tagging — syntactic parsing  
— semantic analysis — coreference resolution

## Why is NLP difficult?

ambiguity  
coverage

## Course admin:

HW0 will be out later today (due Friday, Sep 15)

Office hours

Homework policies (no late submissions...)

Midterm and final exams

Projects and Literature surveys

# Compass and enrollment...

I spoke too soon in the last lecture....

We have 100 seats in this class(room).

I received ~40 requests to be added to Compass.

(Apologies if I haven't replied to your email...)

That is simply not feasible.

We are not able to grade more than 100 assignments.

I am not allowed to let significantly more students into this classroom (fire code).

# Compass and enrollment...

Lecture slides and the PDFs for the assignments will always be posted on the class website.

You don't need to be on Compass for that.

Piazza is also available to everybody.

If you are planning to drop this class, please do so ASAP, so that others can take your spot.

# DRES accommodations

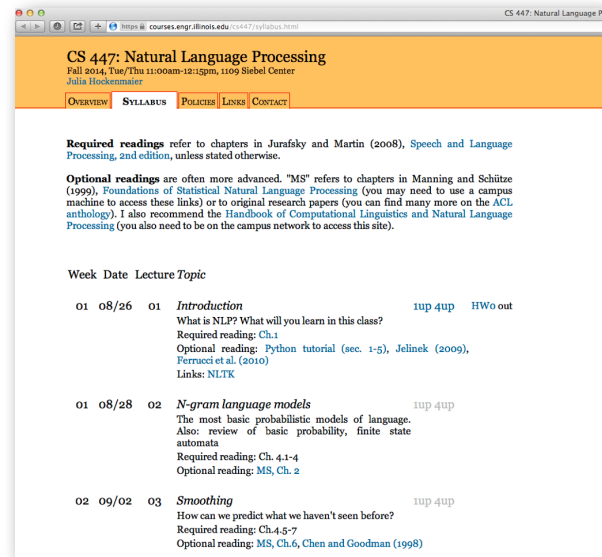

If you need any disability related accommodations, talk to DRES (<http://disability.illinois.edu>, [disability@illinois.edu](mailto:disability@illinois.edu), phone 333-4603)

If you are concerned you have a disability-related condition that is impacting your academic progress, there are academic screening appointments available on campus that can help diagnosis a previously undiagnosed disability by visiting the DRES website and selecting “Sign-Up for an Academic Screening” at the bottom of the page.”

Come and talk to me as well, especially once you have a letter of accommodation from DRES.

Do this early enough so that we can take your requirements into account for exams and assignments.

# Reading



The textbook: <https://web.stanford.edu/~jurafsky/slp3/>  
Jurafsky and Martin, **Speech and Language Processing**  
(3rd edition PDFs in prep.; 2nd edition, 2008 in print)

Other material (Slides, additional reading):

Posted on course website

For some assignments:

The NLTK book (<http://www.nltk.org/book>)

# Assessment

If you take this class for **3 hours credit**:

- 1/3 homework assignments

- 1/3 midterm exam

- 1/3 final exam

If you take this class for **4 hours credit**:

- 1/4 homework assignments

- 1/4 midterm exam

- 1/4 final exam

- 1/4 literature reviews

We reserve the right to improve your grade by up to 5% depending on your class participation. If you're in between grades, but attended class and participated frequently and actively in in-class discussions etc., we will give you the higher grade.

# Homework assignments

## Schedule:

Week 1: Friday, 09/01 HW0 out (today!)

Week 3: Friday, 09/15 HW0 due, HW1 out

Week 6: Friday, 10/06 HW1 due, HW2 out

Week 9: Friday, 10/27 HW2 due, HW3 out

Week 12: Friday, 11/17 HW3 due, HW4 out

Week 15: Wednesday, 12/13 HW4 due (last lecture)

## Points per assignment:

HW0 = 2 points

(Did you submit (on time)? Was it in the right format?)

HW1, HW2, HW3, HW4 = 10 points per assignment



# Homework assignments

For now, we will use **Enthought Canopy Python (2.7.6)**

This is available on the EWS linux machines

```
~> ssh linux.ews.illinois.edu  
[juliahr@linux-a2 ~]$ module load canopy  
[juliahr@linux-a2 ~]$ python  
Enthought Canopy Python 2.7.6 | 64-bit |  
(default, Jun  4 2014, 16:32:15)  
...
```

# 4<sup>th</sup> credit hour: Research Projects

## What?

You need to read and describe a few (2-3) NLP papers on a particular task, implement an NLP system for this task and describe it in a written report.

## Why?

To make sure you get a deeper knowledge of NLP by reading original papers and by building an actual system.

## When?

**Oct 20 (Wk 8):** Proposal due (What topic? What papers will you read?)

**Nov 15 (Wk 12):** Progress report due (Are your experiments on track?)

**Dec 14 (Reading Day):** Final report due (Summary of papers, your system)

# 4<sup>th</sup> credit hour: Literature Survey

## What?

You need to read and describe several (5-7) NLP papers on a particular task or topic, and produce a written report that compares and critiques these approaches.

## Why?

To make sure you get a deeper knowledge of NLP by reading original papers, even if you don't build an actual system.

## When?

**Oct 20 (Wk 8):** Proposal due (What topic? What papers will you read?)

**Nov 15 (Wk 12):** Progress report due (Is your paper on track?)

**Dec 14 (Reading Day):** Final report due (Summary of papers)

# Course Outline (tentative)

- Lectures 2–5: Words and strings  
(morphology, language models)
- Lectures 7–10: Sequence labeling (POS tagging etc.)
- Lectures 11–12: Lexical similarities, word clustering
- Lecture 13: Review for midterm

— — — — — **Midterm exam** — — — — —

- Lecture 14–21: Syntax and Parsing
- Lecture 22–24: Machine Translation
- Lecture 25–28: Semantics, Discourse
- Lecture 29: Review for Final Exam

<http://courses.engr.illinois.edu/cs447/syllabus.html>

# Exams

## What?

Midterm exam: Thursday, Oct 12 (Week 7), 6:30pm DCL 1320

Final exam: TBD [between Dec. 15 and Dec. 21]

(based on material after first midterm)

## Why?

To make sure you understand what you learned well enough to explain and apply it.

## How?

Essay questions and problem questions

Closed-book (no cheatsheets, no electronics, etc.)

Will be based on lectures and readings

# Today's lecture

What is the **structure of words**?

(in English, Chinese, Arabic,...)

**Morphology**: the area of linguistics that deals with this.

How can we identify the structure of words?

We need to build a **morphological analyzer** (parser).

We will use **finite-state transducers** for this task.

Finite-State Automata and Regular Languages  
(Review)

NB: No probabilities or machine learning yet.

We're thinking about (symbolic) representations today.

# Morphology: What is a word?

# A Turkish word

uygarlaştıramadıklarımızdanmışsınızcasına  
uygar\_laş\_tır\_ama\_dık\_lar\_ımız\_dan\_mış\_sınız\_casına

*“as if you are among those whom we were not able to civilize  
(=cause to become civilized)”*

*uygar: civilized*

*\_laş: become*

*\_tır: cause somebody to do something*

*\_ama: not able*

*\_dık: past participle*

*\_lar: plural*

*\_ımız: 1st person plural possessive (our)*

*\_dan: among (ablative case)*

*\_mış: past*

*\_sınız: 2nd person plural (you)*

*\_casına: as if (forms an adverb from a verb)*

*K. Oflazer pc to J&M*



# Basic word classes (parts of speech)

## Content words (open-class):

Nouns: student, university, knowledge,...

Verbs: write, learn, teach,...

Adjectives: difficult, boring, hard, ....

Adverbs: easily, repeatedly,...

## Function words (closed-class):

Prepositions: in, with, under,...

Conjunctions: and, or,...

Determiners: a, the, every,...

# Words aren't just defined by blanks

## Problem 1: Compounding

“ice cream”, “website”, “web site”, “New York-based”

## Problem 2: Other writing systems have no blanks

*Chinese:* 我开始写小说 = 我 开始 写 小说  
*I start(ed) writing novel(s)*

## Problem 3: Clitics

English: “doesn’t”, “I’m”,

Italian: “dirglielo” = dir + gli(e) + lo  
*tell + him + it*

# How many words are there?

Of course he wants to take the advanced course too.  
He already took two beginners' courses.

This is a bad question. Did I mean:

How many **word tokens** are there?

(16 to 19, depending on how we count punctuation)

How many **word types** are there?

(i.e. How many different words are there?

Again, this depends on how you count, but it's  
usually much less than the number of tokens)

# How many words are there?

Of **course** he wants to **take** the advanced **course** **too**.  
He already **took** **two** beginners' **courses**.

The same (underlying) word can take different forms:  
course/courses, take/took

We distinguish concrete **word forms** (take, taking)  
from abstract **lemmas** or dictionary forms (take)

Different words may be spelled/pronounced the same:  
of course vs. advanced course  
two vs. too

# How many different words are there?

**Inflection** creates different forms of the same word:

Verbs: to be, being, I am, you are, he is, I was,

Nouns: one book, two books

**Derivation** creates different words from the same lemma:

grace  $\Rightarrow$  disgrace  $\Rightarrow$  disgraceful  $\Rightarrow$  disgracefully

**Compounding** combines two words into a new word:

cream  $\Rightarrow$  ice cream  $\Rightarrow$  ice cream cone  $\Rightarrow$  ice cream cone bakery

**Word formation is productive:**

New words are subject to all of these processes:

Google  $\Rightarrow$  Googler, to google, to ungoogle, to misgoogle, googlification, ungooglification, googlified, Google Maps, Google Maps service,...

# Inflectional morphology in English

## Verbs:

Infinitive/present tense: walk, go

3rd person singular present tense (s-form): walks, goes

Simple past: walked, went

Past participle (ed-form): walked, gone

Present participle (ing-form): walking, going

## Nouns:

Common nouns inflect for number:

singular (book) vs. plural (books)

Personal pronouns inflect for person, number, gender, case:

I saw him; he saw me; you saw her; we saw them; they saw us.

# Derivational morphology

## Nominalization:

V + -ation: computerization

V+ -er: killer

Adj + -ness: fuzziness

## Negation:

un-: undo, unseen, ...

mis-: mistake,...

## Adjectivization:

V+ -able: doable

N + -al: national

# Morphemes: stems, affixes

**dis-grace-ful-ly**  
**prefix-stem-suffix-suffix**

Many word forms consist of a **stem** plus a number of **affixes** (*prefixes* or *suffixes*)

*Infixes* are inserted inside the stem.

*Circumfixes* (German gesehen) surround the stem

**Morphemes**: the smallest (meaningful/grammatical) parts of words.

*Stems* (grace) are often **free morphemes**.

Free morphemes can occur by themselves as words.

*Affixes* (dis-, -ful, -ly) are usually **bound morphemes**.

Bound morphemes have to combine with others to form words.



# Morphemes and morphs

There are many *irregular word forms*:

Plural nouns add -s to singular: book-books,  
but: box-boxes, fly-flies, child-children

Past tense verbs add -ed to infinitive: walk-walked,  
but: like-liked, leap-leapt

One morpheme (e.g. for plural nouns) can be realized as *different surface forms (morphs)*:

-s/-es/-ren

Allomorphs: two different realizations (-s/-es/-ren)  
of the same underlying morpheme (plural)

# Morphological parsing and generation

# Morphological parsing

disgracefully			
dis	grace	ful	ly
<i>prefix</i>	<i>stem</i>	<i>suffix</i>	<i>suffix</i>
<i>NEG</i>	grace+N	<i>+ADJ</i>	<i>+ADV</i>

# Morphological generation

We cannot enumerate all possible English words, but we would like to capture the rules that define whether a string *could* be an English word or not.

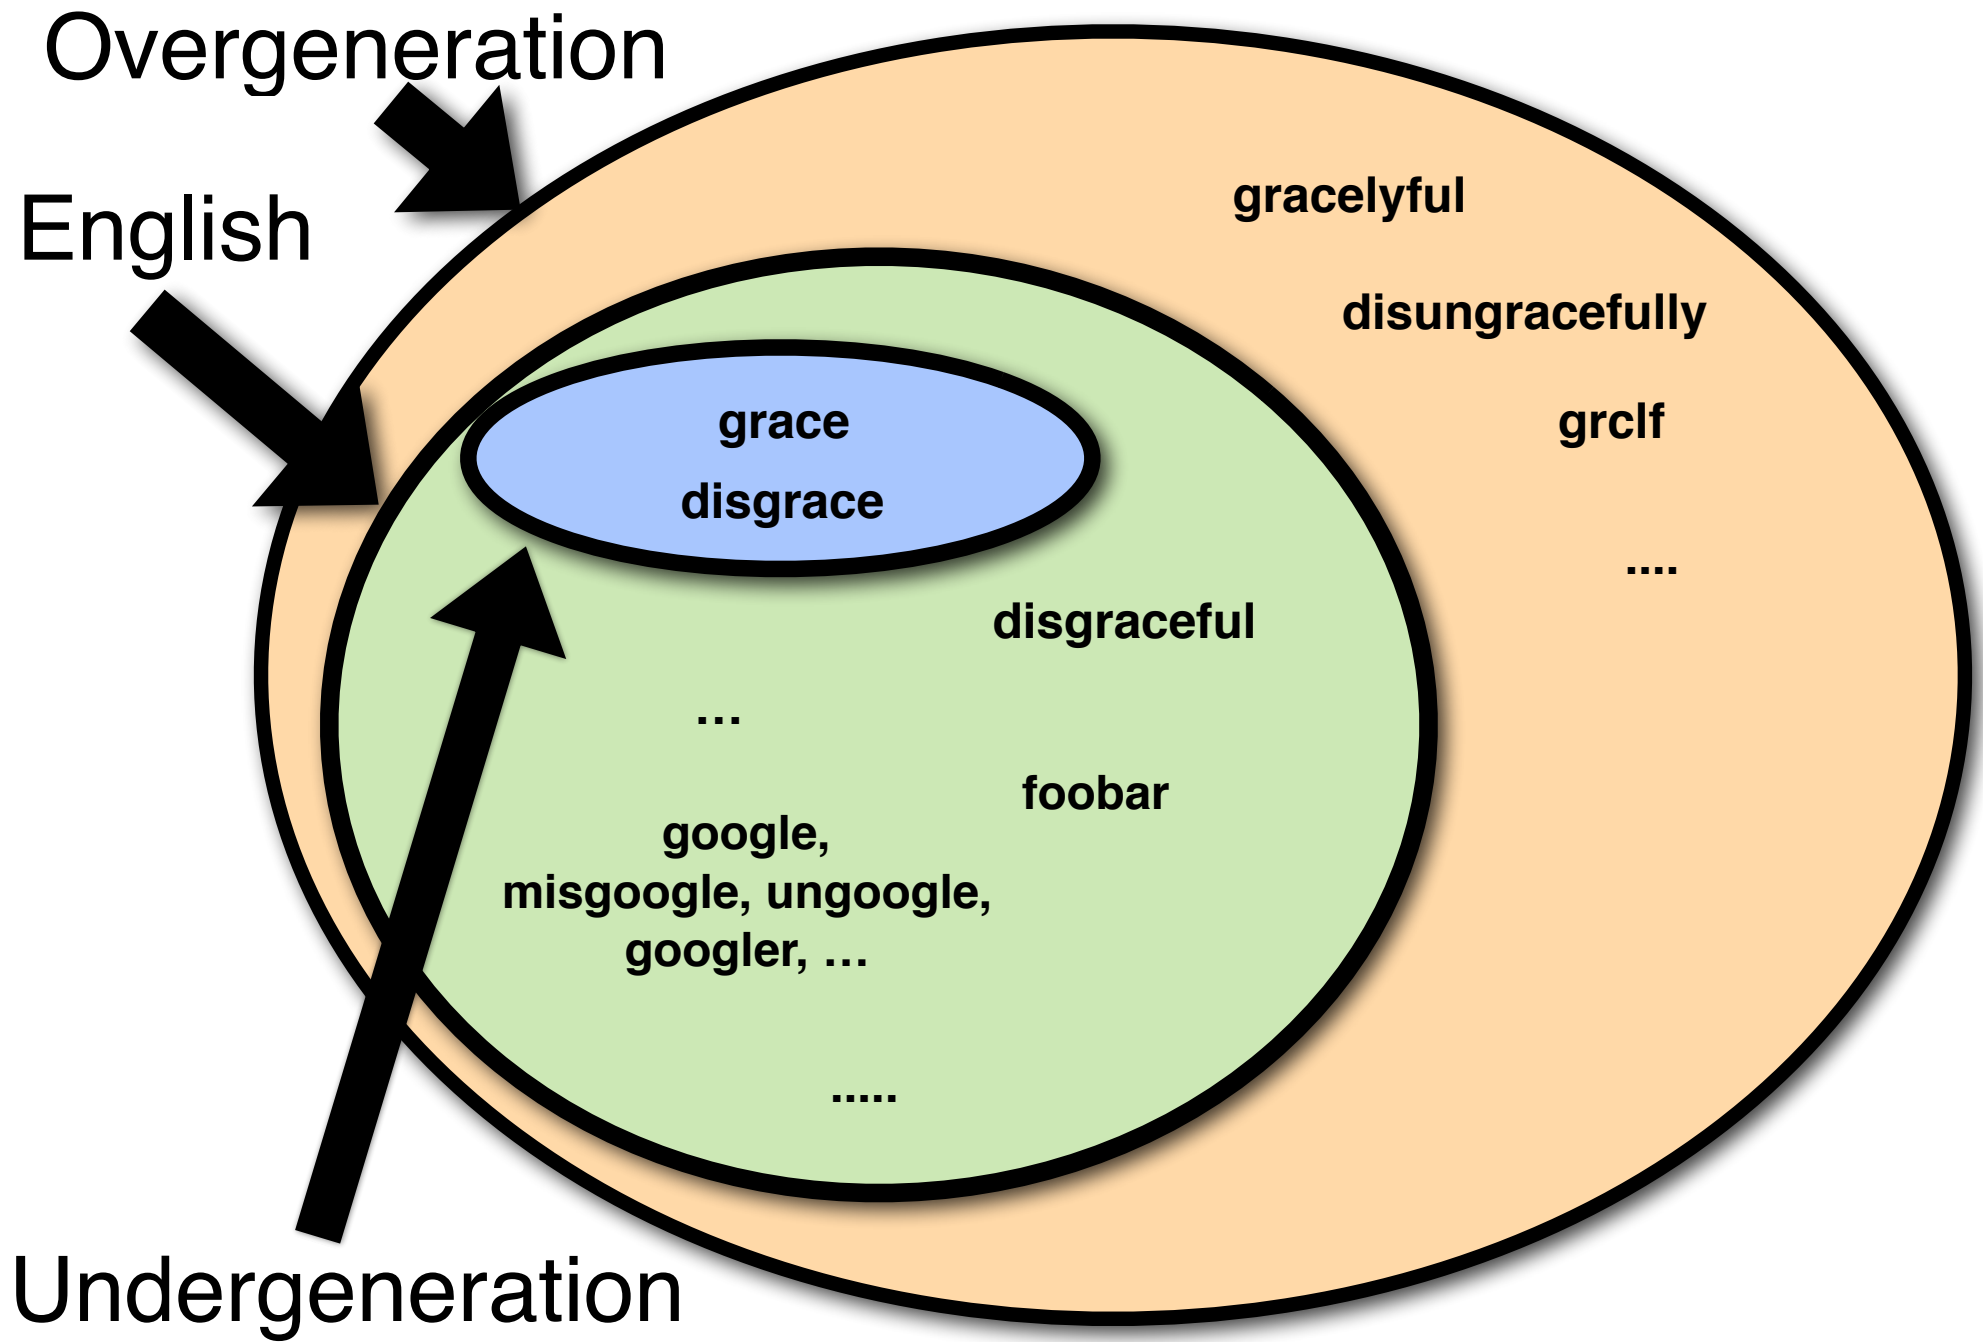
That is, we want a procedure that can generate (or accept) possible English words...

grace, graceful, gracefully  
disgrace, disgraceful, disgracefully,  
ungraceful, ungracefully,  
undisgraceful, undisgracefully,...

without generating/accepting impossible English words

\*gracelyful, \*gracefully, \*disungracefully,...

NB: \* is linguists' shorthand for "this is ungrammatical"



# Review: Finite-State Automata and Regular Languages

# Formal languages

An **alphabet**  $\Sigma$  is a **set of symbols**:

e.g.  $\Sigma = \{a, b, c\}$

A **string**  $\omega$  is a **sequence of symbols**, e.g.  $\omega = abcb$ .

The **empty string**  $\varepsilon$  consists of zero symbols.

The Kleene closure  $\Sigma^*$  (**'sigma star'**) is the **(infinite) set of all strings** that can be formed from  $\Sigma$ :

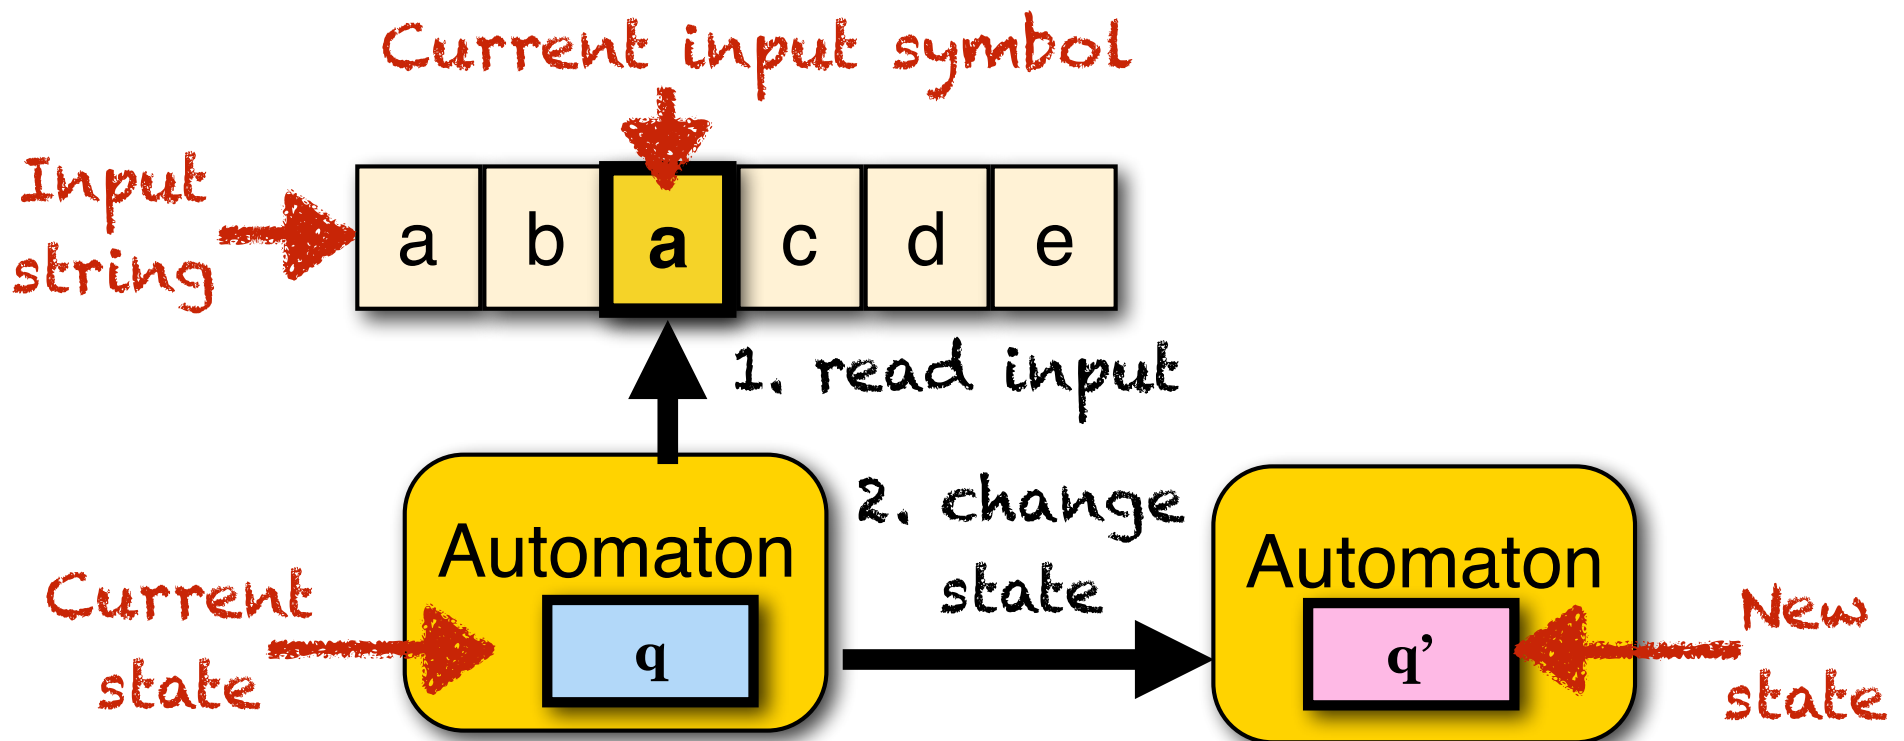
$\Sigma^* = \{\varepsilon, a, b, c, aa, ab, ba, aaa, \dots\}$

A **language**  $L \subseteq \Sigma^*$  over  $\Sigma$  is also a set of strings.

Typically we only care about **proper subsets of  $\Sigma^*$**  ( $L \subset \Sigma^*$ ).

# Automata and languages

An **automaton** is an abstract model of a computer. It *reads* an **input string** symbol by symbol. It *changes* its **internal state** depending on the **current input symbol** and its **current internal state**.






# Automata and languages

The automaton either **accepts** or **rejects** the input string.

**Every automaton defines a language**

(the set of strings it accepts).

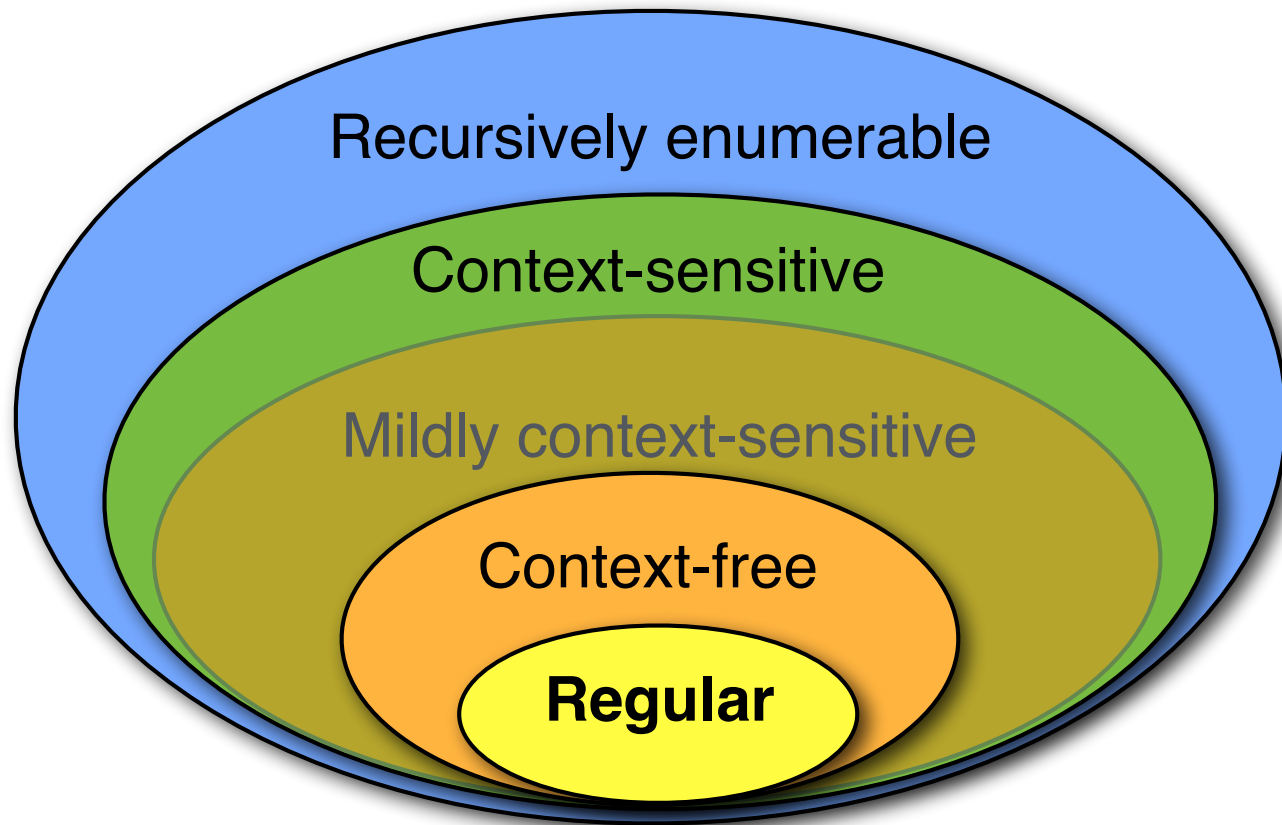


# Automata and languages

Different types of automata define different language classes:

- **Finite-state** automata define **regular** languages
- **Pushdown** automata define **context-free** languages
- **Turing machines** define **recursively enumerable** languages

# The Chomsky Hierarchy









The structure of English words can be described by a regular (= finite-state) grammar.

# Finite-state automata




A (deterministic) finite-state automaton (FSA) consists of:

- a **finite set of states**  $Q = \{q_0, \dots, q_N\}$ , including a **start state**  $q_0$  and one (or more) **final (=accepting) states** (say,  $q_N$ )
- a (**deterministic**) transition function  $\delta(q, w) = q'$  for  $q, q' \in Q, w \in \Sigma$





Start in  $q_0$




Accept!

We've reached the end of the string,  
and are in an accepting state.





# Rejection: Automaton does not end up in accepting state

b




b



# Rejection: Transition not defined



Start in  $q_0$



Reject!  
(There is no  
transition  
labeled 'c')

# Finite State Automata (FSAs)

A finite-state automaton  $M = \langle Q, \Sigma, q_0, F, \delta \rangle$  consists of:

- A finite set of states  $Q = \{q_0, q_1, \dots, q_n\}$
- A finite alphabet  $\Sigma$  of input symbols (e.g.  $\Sigma = \{a, b, c, \dots\}$ )
- A designated start state  $q_0 \in Q$
- A set of final states  $F \subseteq Q$
- A transition function  $\delta$ :
  - The transition function for a **deterministic (D)FSA**:  $Q \times \Sigma \rightarrow Q$   
$$\delta(q, w) = q' \quad \text{for } q, q' \in Q, w \in \Sigma$$


If the current state is  $q$  and the current input is  $w$ , go to  $q'$
  - The transition function for a **nondeterministic (N)FSA**:  $Q \times \Sigma \rightarrow 2^Q$   
$$\delta(q, w) = Q' \quad \text{for } q \in Q, Q' \subseteq Q, w \in \Sigma$$

If the current state is  $q$  and the current input is  $w$ , go to any  $q' \in Q'$



# Finite State Automata (FSAs)

Every NFA can be transformed into an equivalent DFA:



Recognition of a string  $w$  with a DFA is linear in the length of  $w$

Finite-state automata define the class of **regular languages**

$L_1 = \{ a^n b^m \} = \{ ab, aab, abb, aaab, abb, \dots \}$  is a regular language,

$L_2 = \{ a^n b^n \} = \{ ab, aabb, aaabbb, \dots \}$  is not (it's context-free).

You cannot construct an FSA that accepts all the strings in  $L_2$  and nothing else.

# Regular Expressions

Regular expressions can also be used to define a regular language.

## Simple patterns:

- **Standard characters** match themselves: `'a'`, `'1'`
- **Character classes**: `'[abc]'`, `'[0-9]'`, **negation**: `'[^aeiou]'`  
(Predefined: `\s` (whitespace), `\w` (alphanumeric), etc.)
- **Any character** (except newline) is matched by `'.'`


## Complex patterns: (e.g. `^[A-Z]([a-z])+\s`)

- **Group**: `'(...)'`
- **Repetition**: 0 or more times: `'*'`, 1 or more times: `'+'`
- **Disjunction**: `'...|...'`
- **Beginning of line** `'^'` and **end of line** `'$'`


# Finite-state methods for morphology

# Finite state automata for morphology


grace:




dis-grace:



grace-ful:




dis-grace-ful:



# Union: merging automata

grace,  
dis-grace,  
grace-ful,  
dis-grace-ful



# Stem changes


Some irregular words require stem changes:

Past tense verbs:

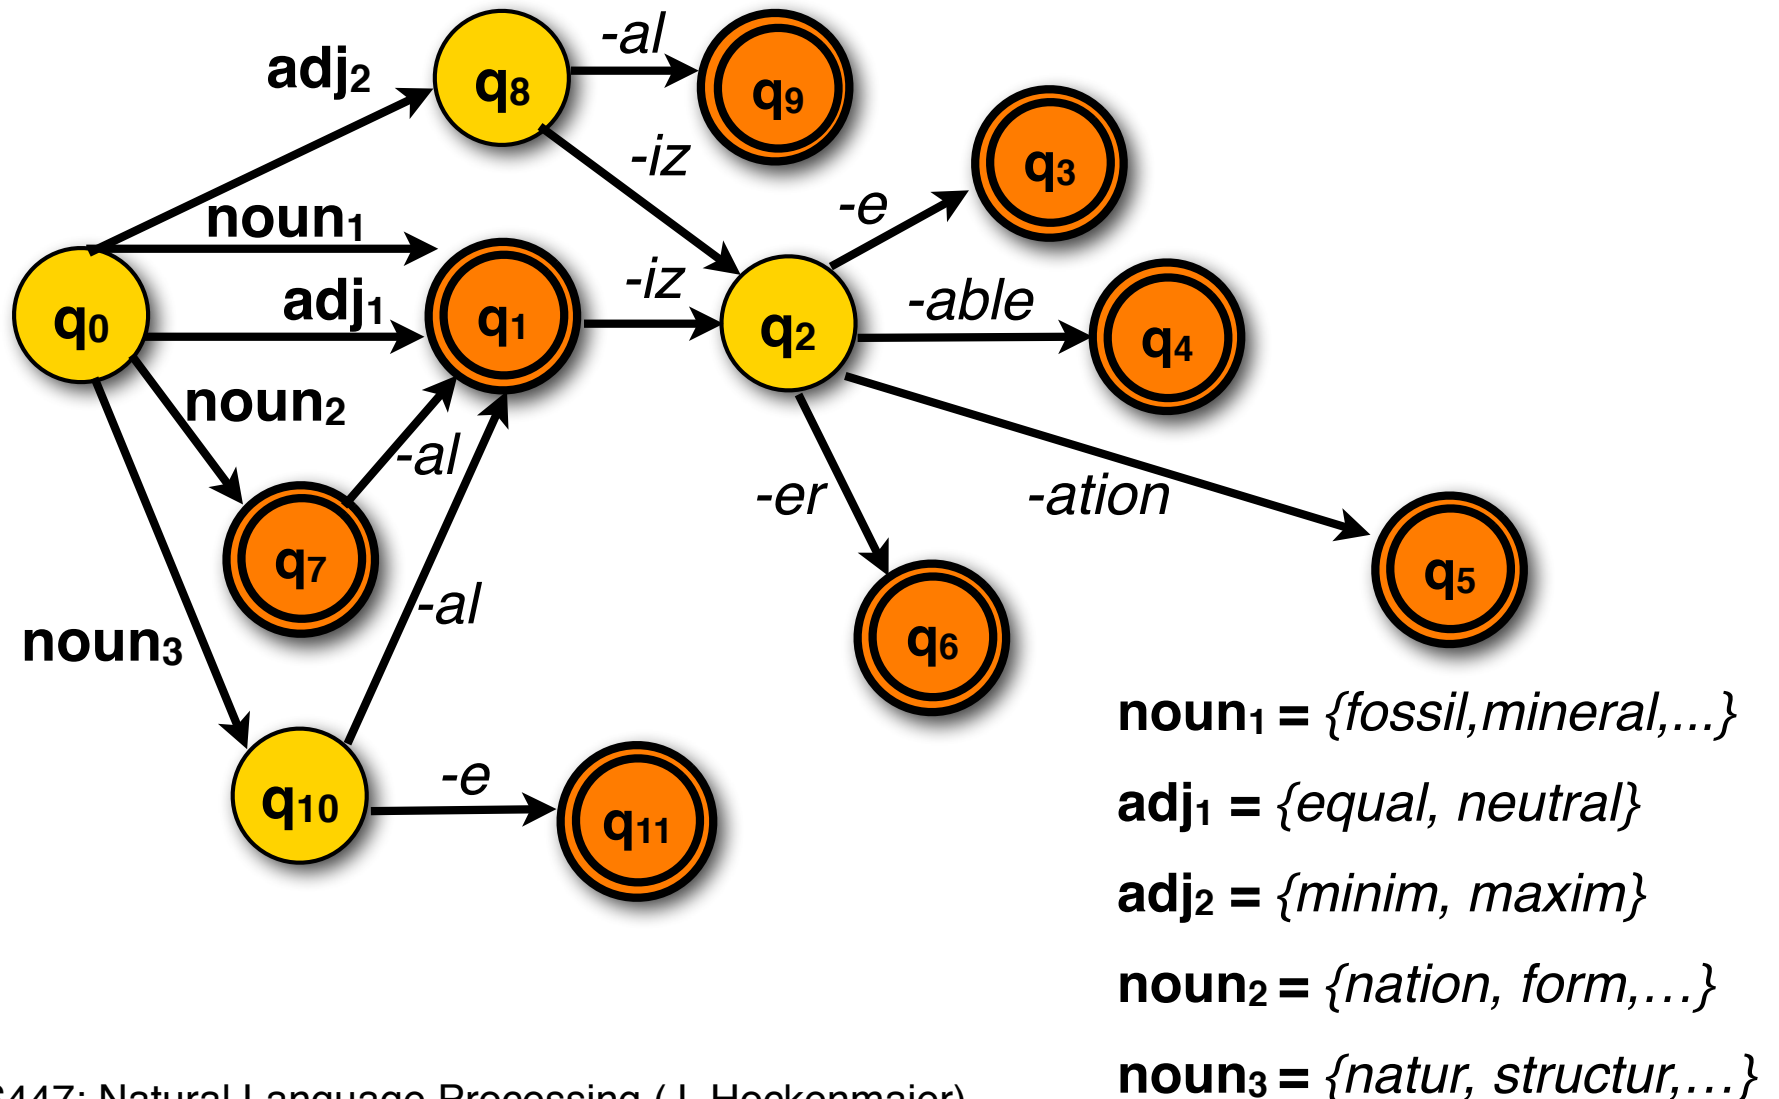
teach-taught, go-went, write-wrote

Plural nouns:

mouse-mice, foot-feet, wife-wives




# FSAs for derivational morphology



# Recognition vs. Analysis

FSAs can recognize (**accept**) a string, but they don't tell us its internal structure.

We need is a machine that maps (**transduces**) the input string into an output string that encodes its structure:





# Finite-state transducers

A **finite-state transducer**  $T = \langle Q, \Sigma, \Delta, q_0, F, \delta, \sigma \rangle$  consists of:

- A finite **set of states**  $Q = \{q_0, q_1, \dots, q_n\}$
- A finite alphabet  $\Sigma$  of **input symbols** (e.g.  $\Sigma = \{a, b, c, \dots\}$ )
- A finite alphabet  $\Delta$  of **output symbols** (e.g.  $\Delta = \{+N, +pl, \dots\}$ )
- A designated **start state**  $q_0 \in Q$
- A set of **final states**  $F \subseteq Q$
- A **transition function**  $\delta: Q \times \Sigma \rightarrow 2^Q$   
 $\delta(q, w) = Q' \quad \text{for } q \in Q, Q' \subseteq Q, w \in \Sigma$
- **An output function**  $\sigma: Q \times \Sigma \rightarrow \Delta^*$   
 $\sigma(q, w) = \omega \quad \text{for } q \in Q, w \in \Sigma, \omega \in \Delta^*$

If the current state is  $q$  and the current input is  $w$ , write  $\omega$ .


(NB: Jurafsky&Martin define  $\sigma: Q \times \Sigma^* \rightarrow \Delta^*$ . Why is this equivalent?)

# Finite-state transducers

An FST  $T = L_{in} \times L_{out}$  defines a **relation between two regular languages**  $L_{in}$  and  $L_{out}$ :

$L_{in} = \{\mathbf{cat}, \mathbf{cats}, \mathbf{fox}, \mathbf{foxes}, \dots\}$

$L_{out} = \{cat+N+sg, cat+N+pl, fox+N+sg, fox+N+pl \dots\}$



$T = \{ \langle \mathbf{cat}, cat+N+sg \rangle, \\ \langle \mathbf{cats}, cat+N+pl \rangle, \\ \langle \mathbf{fox}, fox+N+sg \rangle, \\ \langle \mathbf{foxes}, fox+N+pl \rangle \}$

# Some FST operations

## Inversion $T^{-1}$ :

The inversion ( $T^{-1}$ ) of a transducer switches input and output labels.

*This can be used to switch from **parsing** words to **generating** words.*

## Composition ( $T \circ T'$ ): (Cascade)

Two transducers  $T = L_1 \times L_2$  and  $T' = L_2 \times L_3$  can be composed into a third transducer  $T'' = L_1 \times L_3$ .

*Sometimes **intermediate representations** are useful*

# English spelling rules

Peculiarities of English spelling (orthography)

The same underlying morpheme (e.g. *plural-s*) can have different orthographic “surface realizations” (-s, -es)

This leads to spelling changes at morpheme boundaries:

E-insertion: fox +s = foxes

E-deletion: make +ing = making

# Side note: “Surface realization”?

This terminology comes from Chomskyan Transformational Grammar.

Dominant early approach in theoretical linguistics, superseded by other approaches (“minimalism”).

Not computational, but has some historical influence on computational linguistics (e.g. Penn Treebank)

“**Surface**” = standard English (Chinese, Hindi, etc.).

“Surface string” = a written sequence of characters or words  
vs. “**Deep**”/“**Underlying**” structure/representation:

A more abstract representation.

Might be the same for different sentences with the same meaning.

# Intermediate representations

English plural -s: **cat**  $\Rightarrow$  **cats**   **dog**  $\Rightarrow$  **dogs**

but: fox  $\Rightarrow$  foxes, bus  $\Rightarrow$  buses    buzz  $\Rightarrow$  buzzes

We define an **intermediate representation** to capture **morpheme boundaries (^)** and **word boundaries (#)**:

*Lexicon:*

*cat+N+PL*    *fox+N+PL*

**⇒ Intermediate representation:**

**cat<sup>s</sup>#**

**fox<sup>^</sup>s#**

$\Rightarrow$  *Surface string:*


***cats***

**foxes**


## Intermediate-to-Surface Spelling Rule:

If plural **'s'** follows a morpheme ending in **'x'**, **'z'** or **'s'**, insert **'e'**.

# FST composition/cascade:



# T<sub>lex</sub>: Lexical to intermediate level



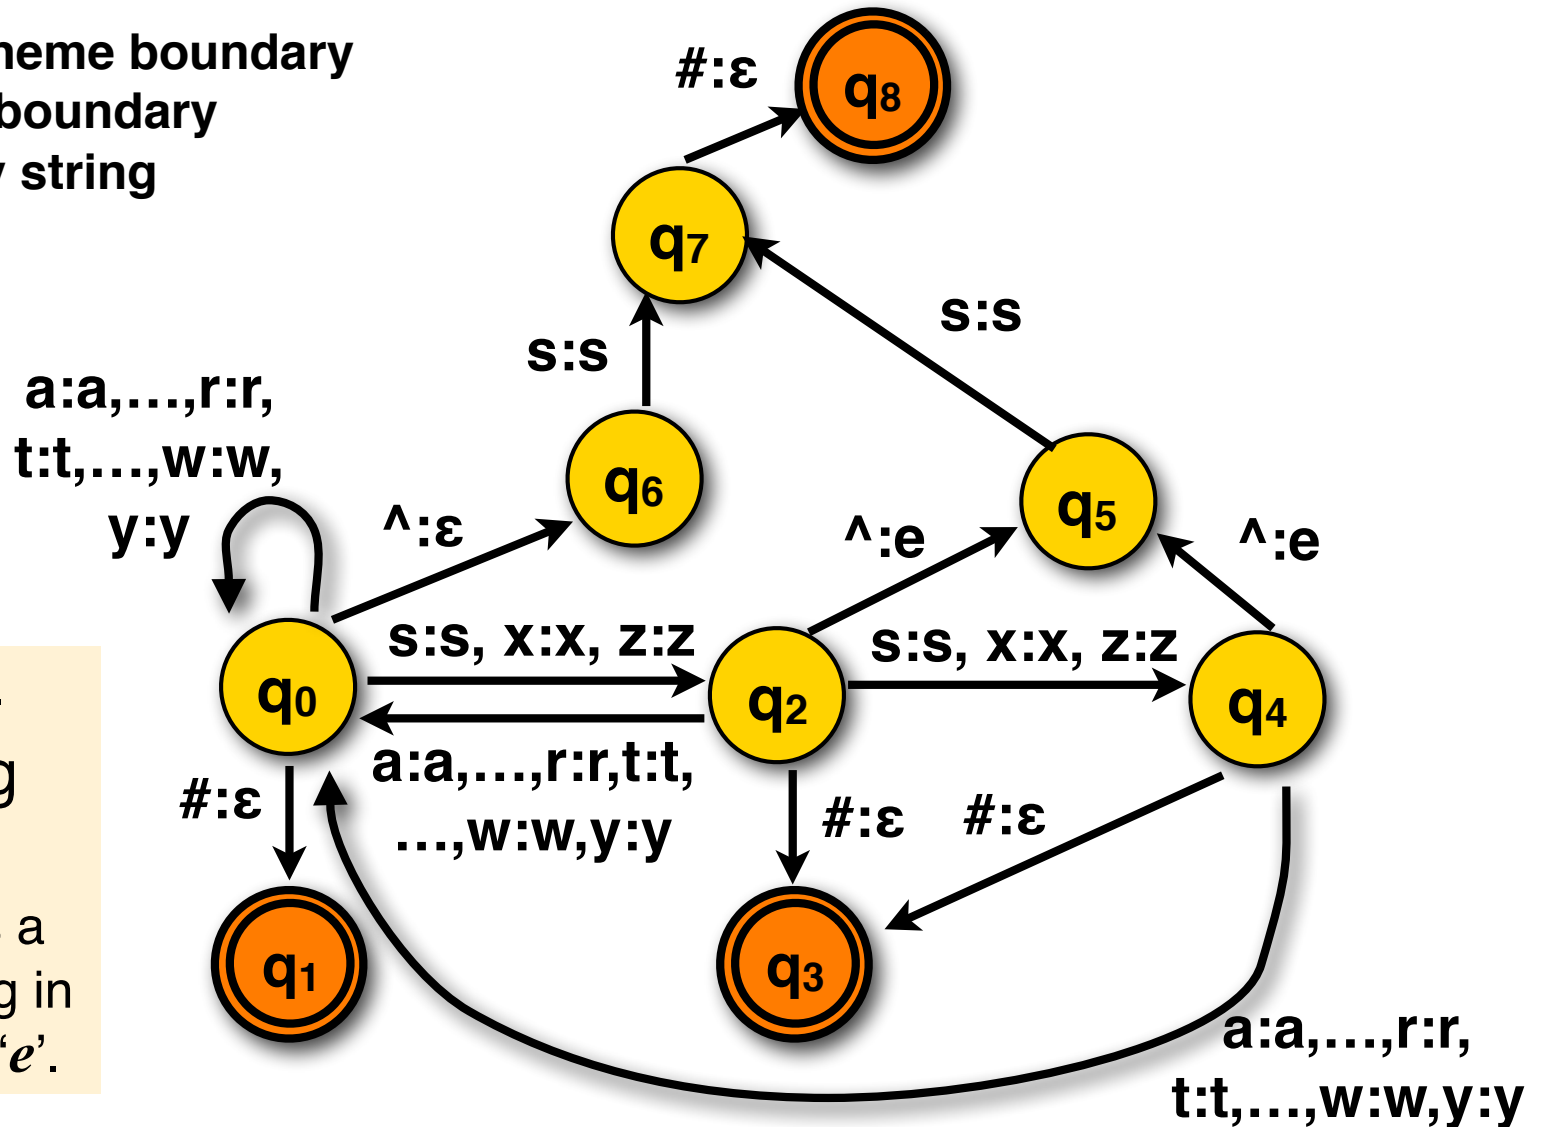


# T<sub>e</sub>-insert: intermediate to surface level

$\wedge$  = morpheme boundary

# = word boundary

$\varepsilon$  = empty string



Intermediate-to-Surface Spelling Rule:

If plural 's' follows a morpheme ending in 'x', 'z' or 's', insert 'e'.

# Dealing with ambiguity

*book: book +N +sg or book +V?*

Generating words is generally unambiguous,  
but **analyzing** words often requires disambiguation.

We need a **nondeterministic FST**.

- Efficiency problem: Not every nondeterministic FST can be translated into a deterministic one!

We also need a **scoring function** to identify which analysis is more likely.

- We may need to know the **context** in which the word appears: (**I read a book** vs. **I book flights**)

# Other applications of FSTs

# A computer therapist?

**Computer:** TELL ME MORE ABOUT YOUR FAMILY

**Human:** My mother takes care of me.

**Computer:** WHO ELSE IN YOUR FAMILY TAKES CARE OF YOU

**Human:** My father.

**Computer:** YOUR FATHER

**Human:** You are like my father in some ways.

**Computer:** WHAT RESEMBLANCE DO YOU SEE

**Human:** You are not very aggressive but I think you don't want me to notice that.

**Computer:** WHAT MAKES YOU THINK I AM NOT VERY AGGRESSIVE

**Human:** You don't argue with me.

**Computer:** WHY DO YOU THINK I DON'T ARGUE WITH YOU

**Human:** You are afraid of me.

*Weizenbaum (1966), ELIZA.*

# ELIZA as a FST cascade

**Human:** *You don't argue with me.*

**Computer:** *WHY DO YOU THINK I DON'T ARGUE WITH YOU*

1. Replace **you** with **I** and **me** with **you**:

*I don't argue with you.*

2. Replace **<...>** with **Why do you think <...>**:

*Why do you think I don't argue with you.*

What about other NLP tasks?

Could we write an FST for machine translation?

# What about compounds?

Semantically, compounds have hierarchical structure:

((ice cream) cone) bakery)  
not (ice ((cream cone) bakery))

((computer science) (graduate student))  
not (computer ((science graduate) student))

We will need context-free grammars to capture this underlying structure.

# Today's key concepts

Morphology (word structure): stems, affixes

Derivational vs. inflectional morphology

Compounding

Stem changes

Morphological analysis and generation

Finite-state automata

Finite-state transducers

Composing finite-state transducers

# Today's reading

This lecture follows closely  
Chapter 3.1-7 in J&M 2008

Optional readings (see website)

Karttunen and Beesley '05, Mohri (1997), the Porter stemmer, Sproat et al. (1996)