

CS447: Natural Language Processing

<http://courses.engr.illinois.edu/cs447>

Lecture 6:

Part-of-speech tagging

Julia Hockenmaier

juliahmr@illinois.edu

3324 Siebel Center

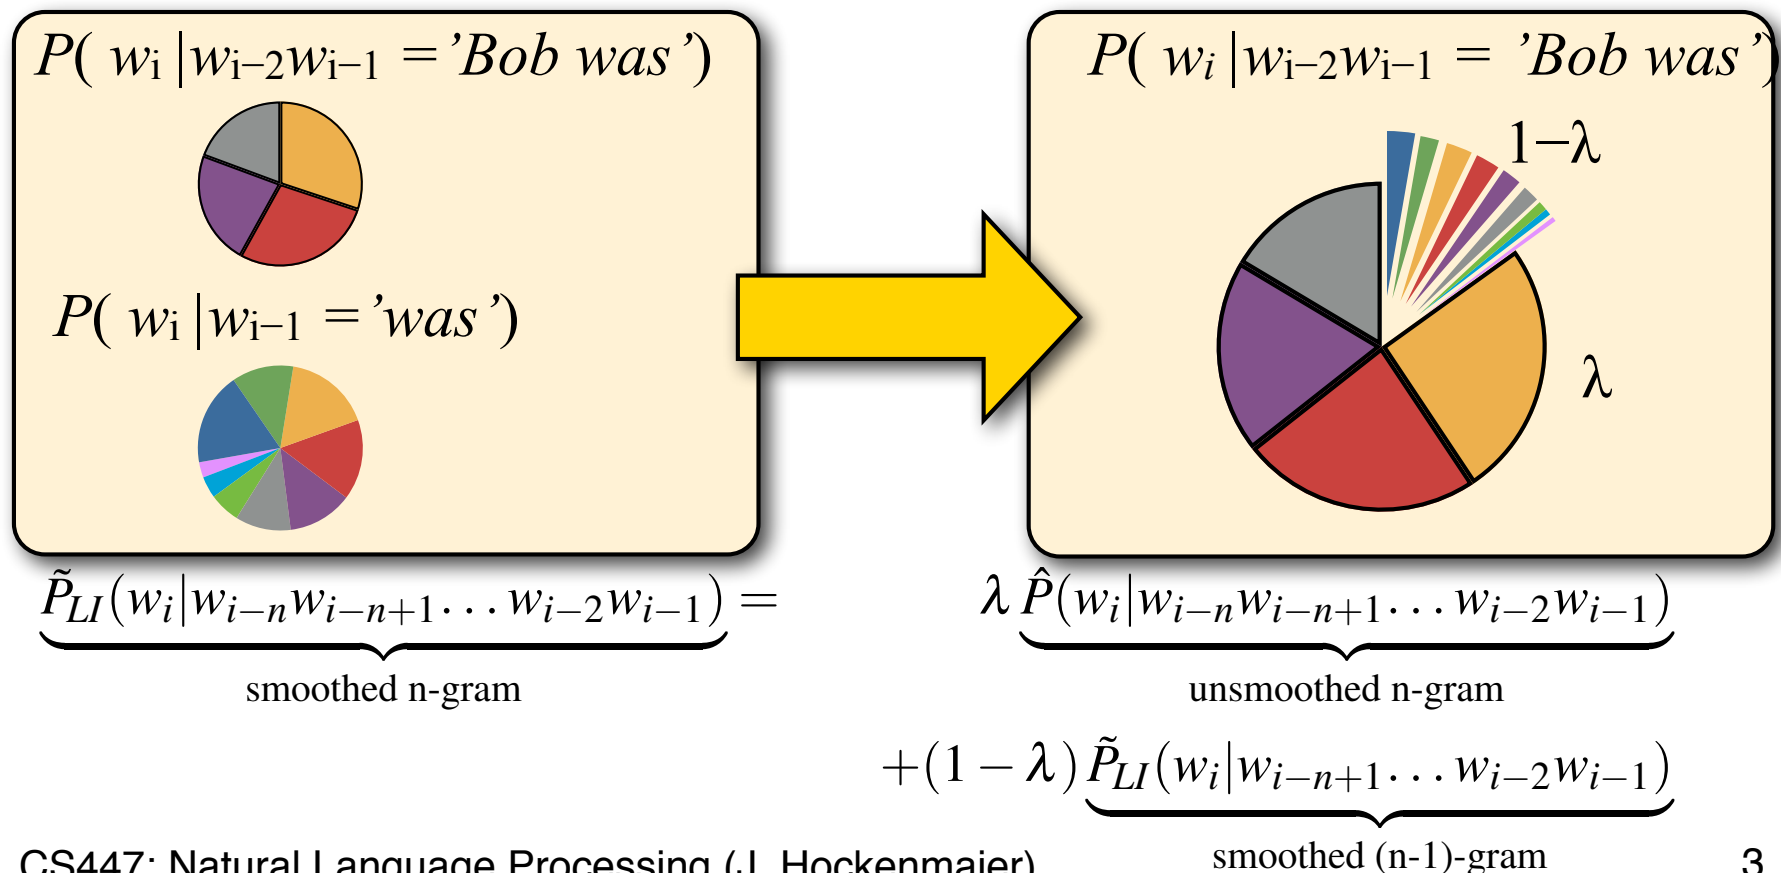
Smoothing: Reserving mass in $P(X|Y)$ for unseen events

Linear Interpolation (Recap)

We don't see “*Bob was reading*”, but we see “__ *was reading*”.

We estimate $P(\text{reading} \mid \text{'Bob was'}) = 0$ but $P(\text{reading} \mid \text{'was'}) > 0$


Use $(n-1)$ -gram probabilities to smooth n -gram probabilities:



What happens to $P(w \mid \dots)$?

The smoothed probability $P_{\text{smoothed-tri}}(w_i \mid w_{i-2} w_{i-1})$ is a linear combination of $P_{\text{unsmoothed-tri}}(w_i \mid w_{i-2} w_{i-1})$ and $P_{\text{bi}}(w_i \mid w_{i-1})$:

$$P_{\text{smoothed-tri}}(w_i \mid w_{i-2} w_{i-1}) = (1-\lambda) P_{\text{unsmoothed-tri}}(w_i \mid w_{i-2} w_{i-1}) + \lambda P_{\text{bi}}(w_i \mid w_{i-1})$$



Absolute discounting

Subtract a constant factor $D < 1$ from each nonzero n -gram count, and interpolate with $P_{AD}(w_i | w_{i-1})$:

non-zero if trigram $w_{i-2}w_{i-1}w_i$ is seen

$$P_{AD}(w_i | w_{i-1}, w_{i-2}) = \frac{\max(C(w_{i-2}w_{i-1}w_i) - D, 0)}{C(w_{i-2}w_{i-1})} + (1 - \lambda)P_{AD}(w_i | w_{i-1})$$

If S seen word types occur after $w_{i-2} w_{i-1}$ in the training data, this reserves the probability mass $P(U) = (S \times D)/C(w_{i-2}w_{i-1})$ to be computed according to $P(w_i | w_{i-1})$. Set:

$$(1 - \lambda) = P(U) = \frac{S \cdot D}{C(w_{i-2}w_{i-1})}$$

N.B.: with N_1, N_2 the number of n -grams that occur once or twice, $D = N_1/(N_1+2N_2)$ works well in practice

Kneser-Ney smoothing

Observation: “*San Francisco*” is frequent, but “*Francisco*” only occurs after “*San*”.

Solution: the **unigram probability** $P(w)$ should not depend on the frequency of w , but on the **number of contexts in which w appears**

$N_{+1}(\bullet w)$: number of contexts in which w appears
= number of word types w' which precede w

$$N_{+1}(\bullet\bullet) = \sum_{w'} N_{+1}(\bullet w')$$

Kneser-Ney smoothing: Use absolute discounting, but use $P(w) = N_{+1}(\bullet w)/N_{+1}(\bullet\bullet)$

Modified Kneser-Ney smoothing: Use different D for *bigrams and trigrams* (Chen & Goodman '98)

Class Admin

Homework assignments

Schedule:

Week 1: Friday, 09/25 HW0 out (today!)

Week 3: Friday, 09/15 HW0 due, HW1 out

Week 6: Friday, 10/06 HW1 due, HW2 out

Week 9: Friday, 10/27 HW2 due, HW3 out

Week 12: Friday, 11/17 HW3 due, HW4 out

Week 15: Wednesday, 12/13 HW4 due (last lecture)

Points per assignment:

HW0 = 2 points

(Did you submit (on time)? Was it in the right format?)

HW1, HW2, HW3, HW4 = 10 points per assignment

Homework assignments

HW0 is due **at 10pm** today.

HW1 will go out this evening.

HW1 is due **at 10pm** on Friday, Oct 6.

We use Compass.

Email us **ASAP** if you cannot access the Compass page for our class.

4th credit hour

Two choices:

- **Research project** (alone, or with one other student)
- **Literature survey** (alone)

Deadlines:

- **Before Oct 20: Check your idea with me**
- **Oct 20 (Wk 8): Proposal due**
(What topic? What papers will you read?)
- **Nov 15 (Wk 12): Progress report due**
(Are your experiments on track? Is your paper on track?)
- **Dec 14 (Reading Day): Final report due**
(Summary of papers, your system)

4th credit hour: Research Projects

What?

You need to read and describe a few (2-3) NLP papers on a particular task, implement an NLP system for this task and describe it in a written report.

Why?

To make sure you get a deeper knowledge of NLP by reading original papers and by building an actual system.

When?

Oct 20 (Wk 8): Proposal due (What topic? What papers will you read?)

Nov 15 (Wk 12): Progress report due (Are your experiments on track?)

Dec 14 (Reading Day): Final report due (Summary of papers, your system)

4th credit hour: Literature Survey

What?

You need to read and describe several (5-7) NLP papers on a particular task or topic, and produce a written report that compares and critiques these approaches.

Why?

To make sure you get a deeper knowledge of NLP by reading original papers, even if you don't build an actual system.

When?

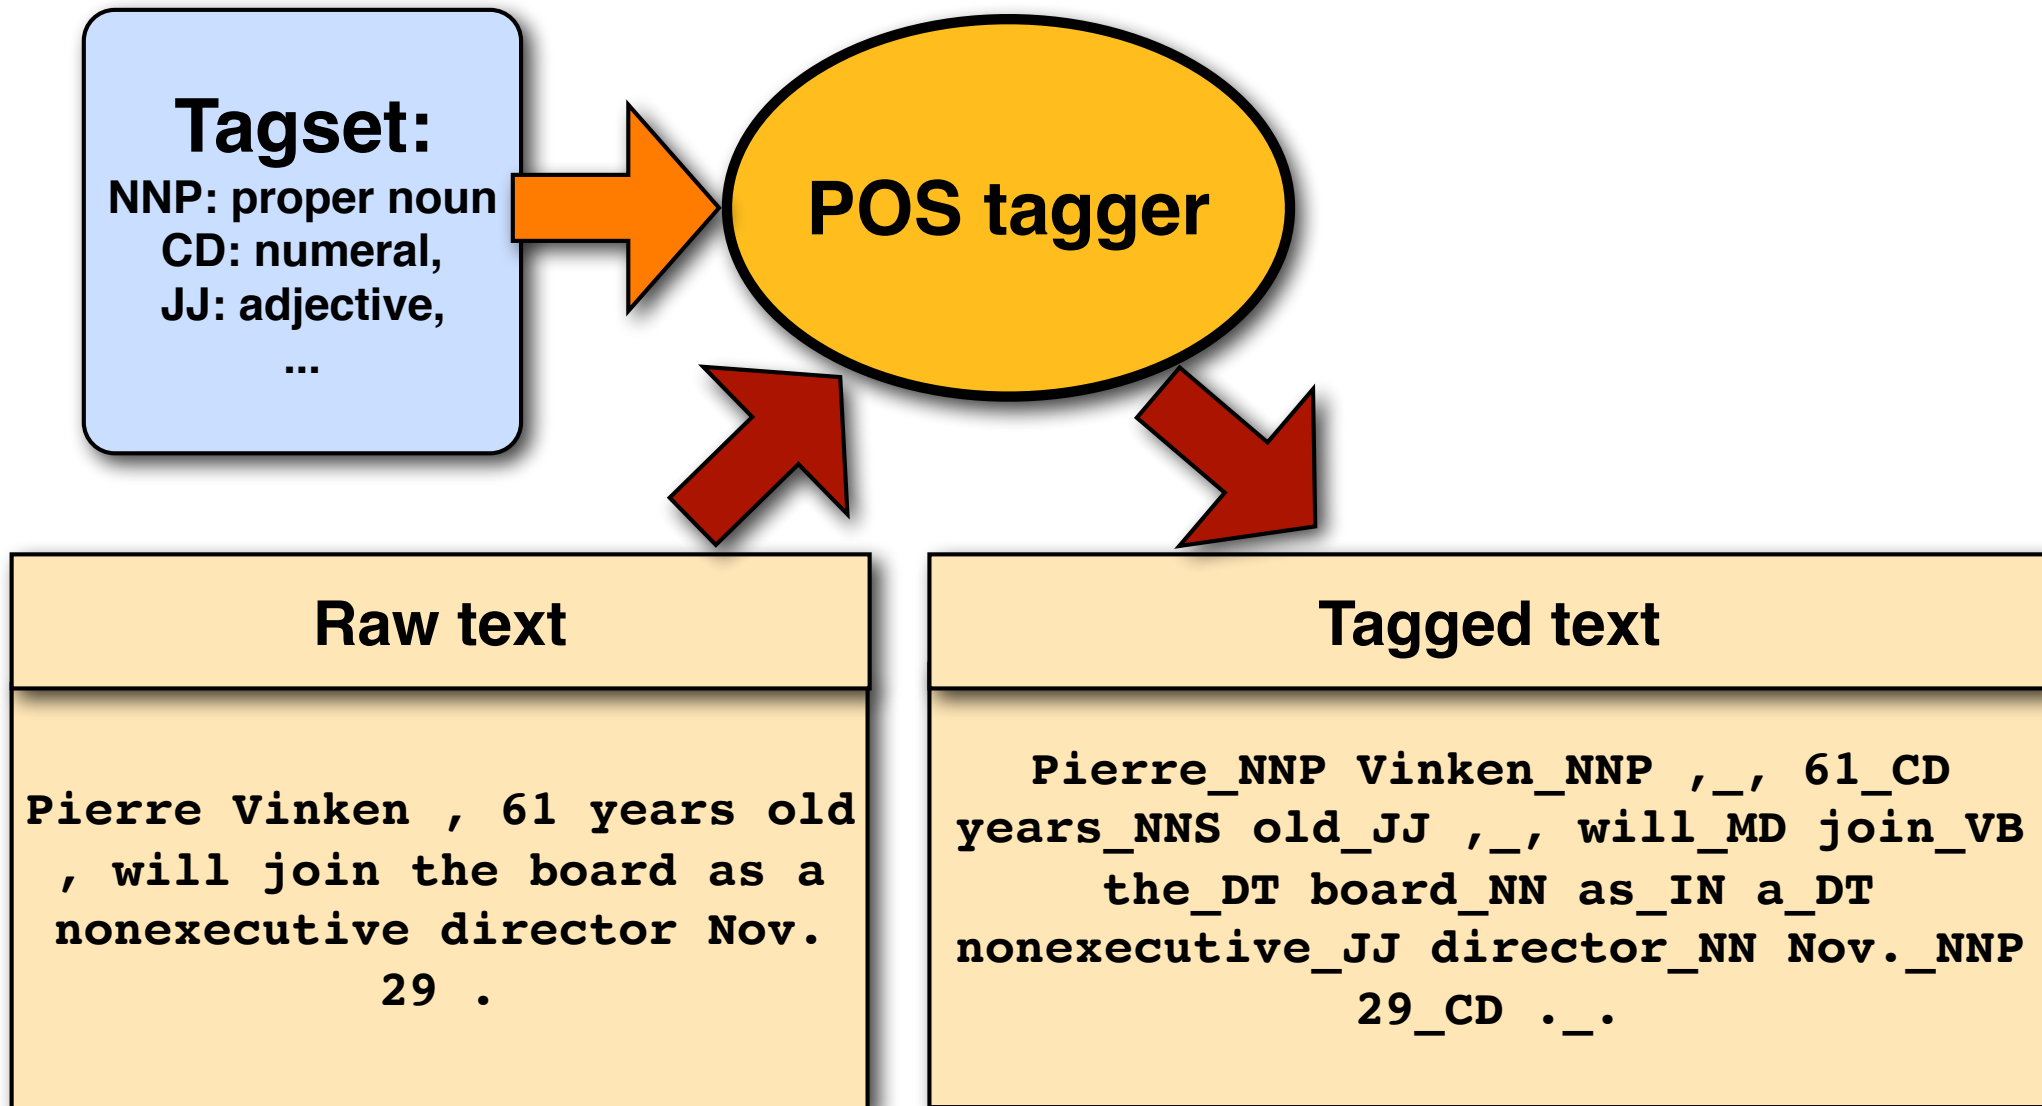
Oct 20 (Wk 8): Proposal due (What topic? What papers will you read?)

Nov 15 (Wk 12): Progress report due (Is your paper on track?)

Dec 14 (Reading Day): Final report due (Summary of papers)

Part-of-speech (POS) tagging

POS tagging



What is POS tagging?

POS = part-of-speech:

- Broad word classes:

Noun, Verb, Adjective, Adverb, Preposition, ...

- More fine-grained distinctions, e.g.:

Common noun, proper noun, pronoun, infinitive, past participle, ...

POS tag:

A label for a particular part-of-speech

(e.g. NN = common noun (singular), VB = infinitive verb, ...)

POS tagging:

Determine for each word in a sentence

which part-of-speech tag it has in that context.

POS Tagging: Ambiguity

Words often have more than one POS:

The back door	(adjective)
On my back	(noun)
Win the voters back	(particle)
Promised to back the bill	(verb)

The POS tagging task is to determine the POS tag for a particular instance of a word.

Since there is ambiguity, we cannot simply look up the correct POS in a dictionary.

These examples from Dekang Lin

Why POS tagging?

POS tagging is a prerequisite for further analysis:

- Parsing:

POS tags give a good indication of possible grammatical analyses.

- Information extraction:

Finding names, relations, etc.

- Machine Translation:

The noun “content” may have a different translation from the adjective

- Speech synthesis:

- How to pronounce “lead”?
- INsult or inSULT, OBject or obJECT, OVERflow or overFLOW, DIScount or disCOUNT, CONtent or conTENT

Defining a tagset

Word classes

Open classes:

Nouns, Verbs, Adjectives, Adverbs

Closed classes:

Auxiliaries and modal verbs

Prepositions, Conjunctions

Pronouns, Determiners

Particles, Numerals

(see Appendix for details)

Defining a tagset

Tagsets have different granularities:

- Brown corpus (Francis and Kucera 1982): 87 tags
 - Penn Treebank (Marcus et al. 1993): 45 tags
- Simplified version of Brown tag set; de facto standard for English now:

NN: common noun (singular or mass): water, book

NNS: common noun (plural): books

- Prague Dependency Treebank (Czech): 4452 tags

Complete morphological analysis:

AAFP3----3N----: (nejnezajímavějším) [Hajic 2006, VMC tutorial]

Adjective Regular Feminine Plural Dative....Superlative

Tagsets for English

We have to agree on a standard inventory of word classes.

Most taggers rely on statistical models; therefore the tagsets used in large corpora become de facto standard.

Tagsets need to capture semantically or syntactically important distinctions that can easily be made by trained human annotators.

How much ambiguity is there?

How many tags does each word type have?

(Original Brown corpus: 40% of tokens are ambiguous)

		87-tag Original Brown	45-tag Treebank Brown
Unambiguous (1 tag)		44,019	38,857
Ambiguous (2–7 tags)		5,490	8844
Details:	2 tags	4,967	6,731
	3 tags	411	1621
	4 tags	91	357
	5 tags	17	90
	6 tags	2 (<i>well, beat</i>)	32
	7 tags	2 (<i>still, down</i>)	6 (<i>well, set, round, open, fit, down</i>)
	8 tags		4 (<i>'s, half, back, a</i>)
	9 tags		3 (<i>that, more, in</i>)

NB: These are just the tags that words appeared with in the corpus.

There are always unseen word/tag combinations.

Evaluating POS taggers

Evaluation metric: accuracy

How many words in the unseen test data can you tag correctly?

State of the art on Penn Treebank: around 97%.

Compare your model against a baseline

Standard: assign to each word its most likely tag
(use training corpus to estimate $P(t|w)$)

Baseline performance on Penn Treebank: around 93.7%

... and a (human) ceiling

How often do human annotators agree on the same tag?

Penn Treebank: around 97%

Is POS-tagging a solved task?

Penn Treebank POS-tagging accuracy
 \approx human ceiling

Yes, but:


Other languages with more complex morphology
need much larger tagsets for tagging to be useful,
and will contain many more distinct word forms
in corpora of the same size

They often have much lower accuracies

Evaluating POS taggers


Evaluation setup:

- Split data into training (+development) and separate test sets.



Better setup: n -fold cross validation:

- Split data into n sets of equal size
- Run n experiments, using set $i=1 \dots n$ to test and remainder to train. This gives average, maximal and minimal accuracies



When comparing two taggers:

- Use the same test and training data with the same tagset

Qualitative evaluation

Generate a **confusion matrix** (for development data):
How often was tag i mistagged as tag j:

		Predicted tag						
		IN	JJ	NN	NNP	RB	VBD	VBN
Correct tag	IN	—	.2			.7		
	JJ	.2	—	3.3	2.1	1.7	.2	2.7
	NN		8.7	—				.2
	NNP	.2	3.3	4.1	—	.2		
	RB	2.2	2.0	.5		—		
	VBD		.3	.5			—	4.4
	VBN		2.8				2.6	—

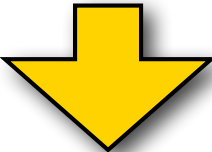
Percent of overall tagging error

See what errors are causing problems:

- Noun (NN) vs ProperNoun (NNP) vs Adj (JJ)
- Preterite (VBD) vs Participle (VBN) vs Adjective (JJ)

Building a POS tagger

Statistical POS tagging

	she ₁	promised ₂	to ₃	back ₄	the ₅	bill ₆
w =	w ₁	w ₂	w ₃	w ₄	w ₅	w ₆
						
t =	t ₁	t ₂	t ₃	t ₄	t ₅	t ₆
	PRP ₁	VBD ₂	TO ₃	VB ₄	DT ₅	NN ₆

What is the most likely sequence of tags **t** for the given sequence of words **w** ?

Statistical POS tagging

What is the most likely sequence of tags \mathbf{t} for the given sequence of words \mathbf{w} ?

$$\begin{aligned}\operatorname{argmax}_{\mathbf{t}} P(\mathbf{t}|\mathbf{w}) &= \operatorname{argmax}_{\mathbf{t}} \frac{P(\mathbf{t}, \mathbf{w})}{P(\mathbf{w})} \\ &= \operatorname{argmax}_{\mathbf{t}} P(\mathbf{t}, \mathbf{w}) \\ &= \operatorname{argmax}_{\mathbf{t}} P(\mathbf{t})P(\mathbf{w}|\mathbf{t})\end{aligned}$$

$P(\mathbf{t})$: Generating $\mathbf{t}=t_1 \dots t_n$

We make the same Markov assumption as in language modeling:

$$\begin{aligned} P(t_1 t_2 \dots t_{n-1} t_n) &= P(t_1) P(t_2 | t_1) \dots P(t_n | t_1 \dots t_{n-1}) \\ &= \prod_{i=1}^n P(t_i | t_1 \dots t_{i-1}) \\ &:=_{def} \prod_{i=1}^n P(t_i | t_{i-1} \dots t_{i-1+n}) \end{aligned}$$

We define an n-gram model over POS tags

$P(\mathbf{w}|\mathbf{t})$: Generating $\mathbf{w}=w_1\dots w_n$

We assume that words are independent of each other, and depend only on their own POS-tag:

$$\begin{aligned} P(w_1 w_2 \dots w_{n-1} w_n | t_1 t_2 \dots t_{n-1} t_n) &= \prod_i P(w_i | w_1 \dots w_{i-1}, t_1 \dots t_i \dots t_n) \\ &:= \text{def} \prod_i P(w_i | t_i) \end{aligned}$$

Hidden Markov Models

HMM models are **generative models** of $P(\mathbf{w}, \mathbf{t})$

$P(\mathbf{w}, \mathbf{t})$ describes a stochastic process that “generates” the data.

HMMs decompose $P(\mathbf{t}, \mathbf{w})$ as $P(\mathbf{t})P(\mathbf{w} | \mathbf{t})$

HMMs make **two independence assumptions**:

- a) approximate $P(\mathbf{t})$ with an N-gram model
- b) assume that each word depends only on its POS tag

An example HMM

Transition Matrix A


	D	N	V	A	.
D		0.8		0.2	
N		0.7	0.3		
V	0.6				0.4
A		0.8		0.2	
.					

Emission Matrix B

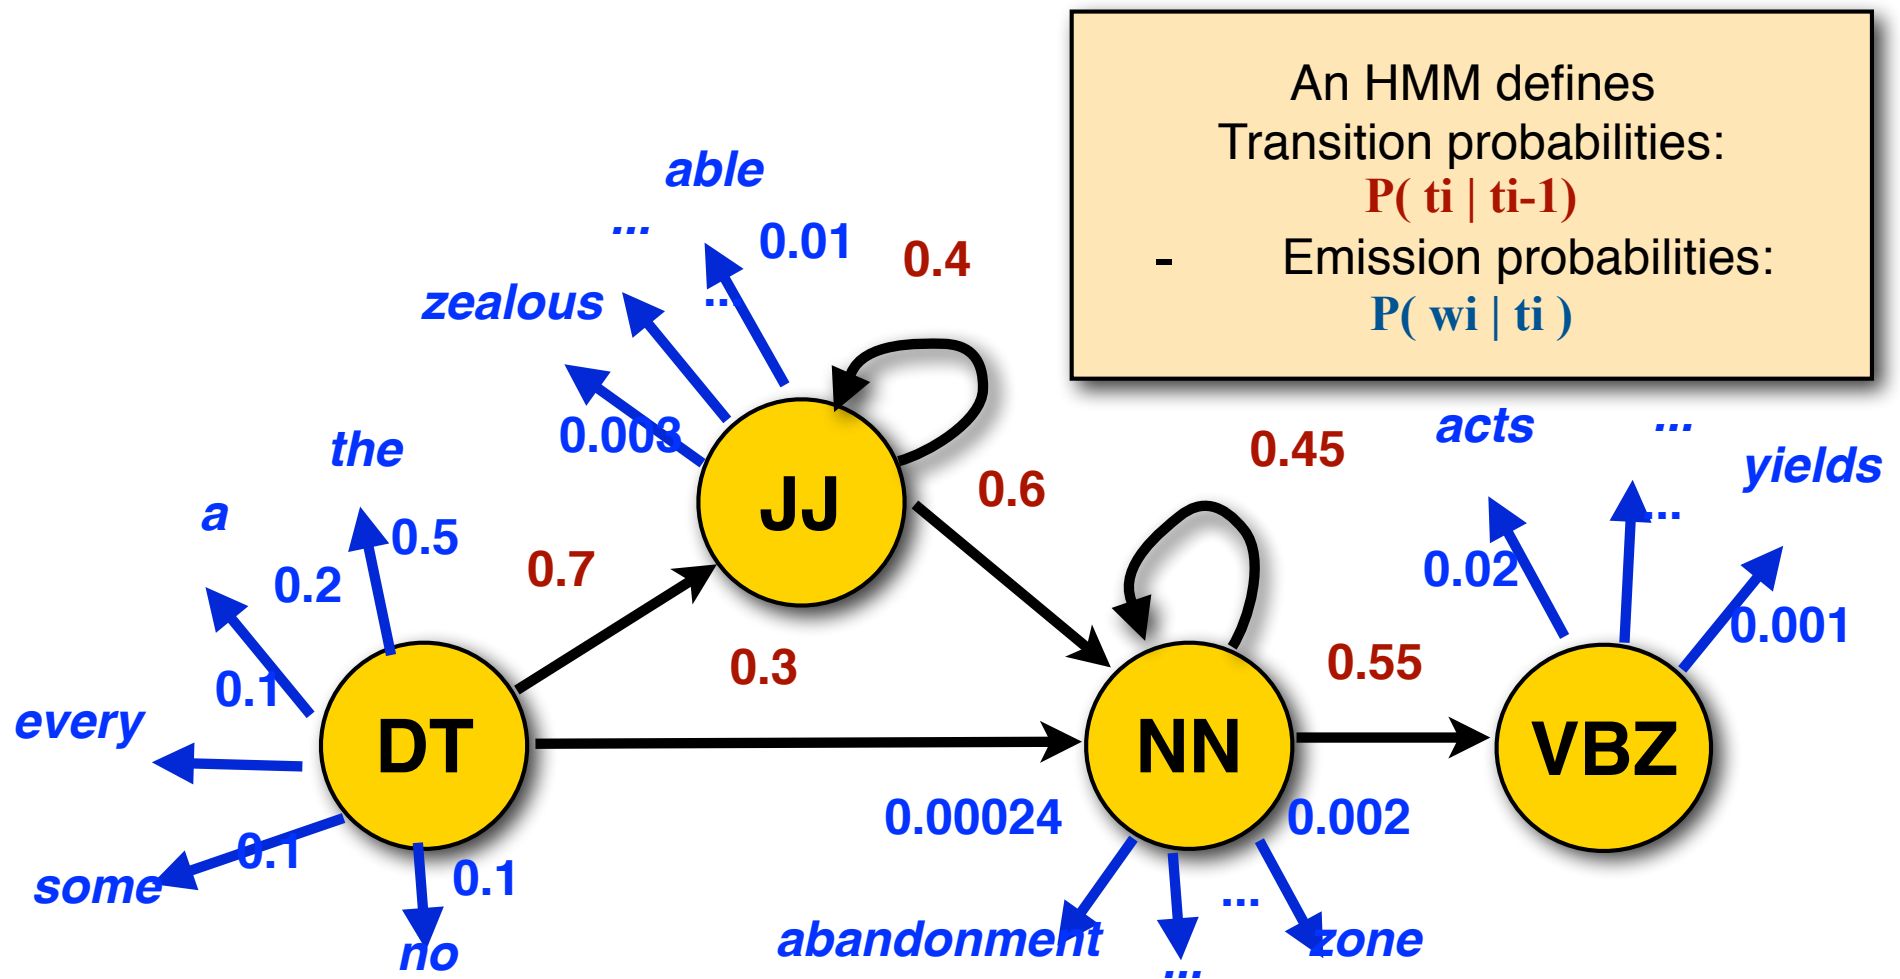
	<i>the</i>	<i>man</i>	<i>ball</i>	<i>throws</i>	<i>sees</i>	<i>red</i>	<i>blue</i>	.
D	1							
N		0.7	0.3					
V				0.6	0.4			
A						0.8	0.2	
.								1

Initial state vector π

	D	N	V	A	.
π	1				




HMMs as probabilistic automata



Using HMMs for tagging

- The input to an HMM tagger is a sequence of words, \mathbf{w} .
The output is the most likely sequence of tags, \mathbf{t} , for \mathbf{w} .
- For the underlying HMM model, \mathbf{w} is a sequence of output symbols, and \mathbf{t} is the most likely sequence of states (in the Markov chain) that generated \mathbf{w} .


$$\begin{aligned}
 \operatorname{argmax}_{\mathbf{t}} P(\underbrace{\mathbf{t}}_{\text{Output}_{\text{tagger}}} \mid \underbrace{\mathbf{w}}_{\text{Input}_{\text{tagger}}}) &= \operatorname{argmax}_{\mathbf{t}} \frac{P(\mathbf{w}, \mathbf{t})}{P(\mathbf{w})} \\
 &= \operatorname{argmax}_{\mathbf{t}} P(\mathbf{w}, \mathbf{t}) \\
 &= \operatorname{argmax}_{\mathbf{t}} P(\underbrace{\mathbf{w}}_{\text{Output}_{\text{HMM}}} \mid \underbrace{\mathbf{t}}_{\text{States}_{\text{HMM}}}) P(\underbrace{\mathbf{t}}_{\text{States}_{\text{HMM}}})
 \end{aligned}$$


How would the
automaton for a trigram
HMM with transition probabilities
 $P(t_i \mid t_{i-2}t_{i-1})$ look like?


What about unigrams
or n-grams?

???

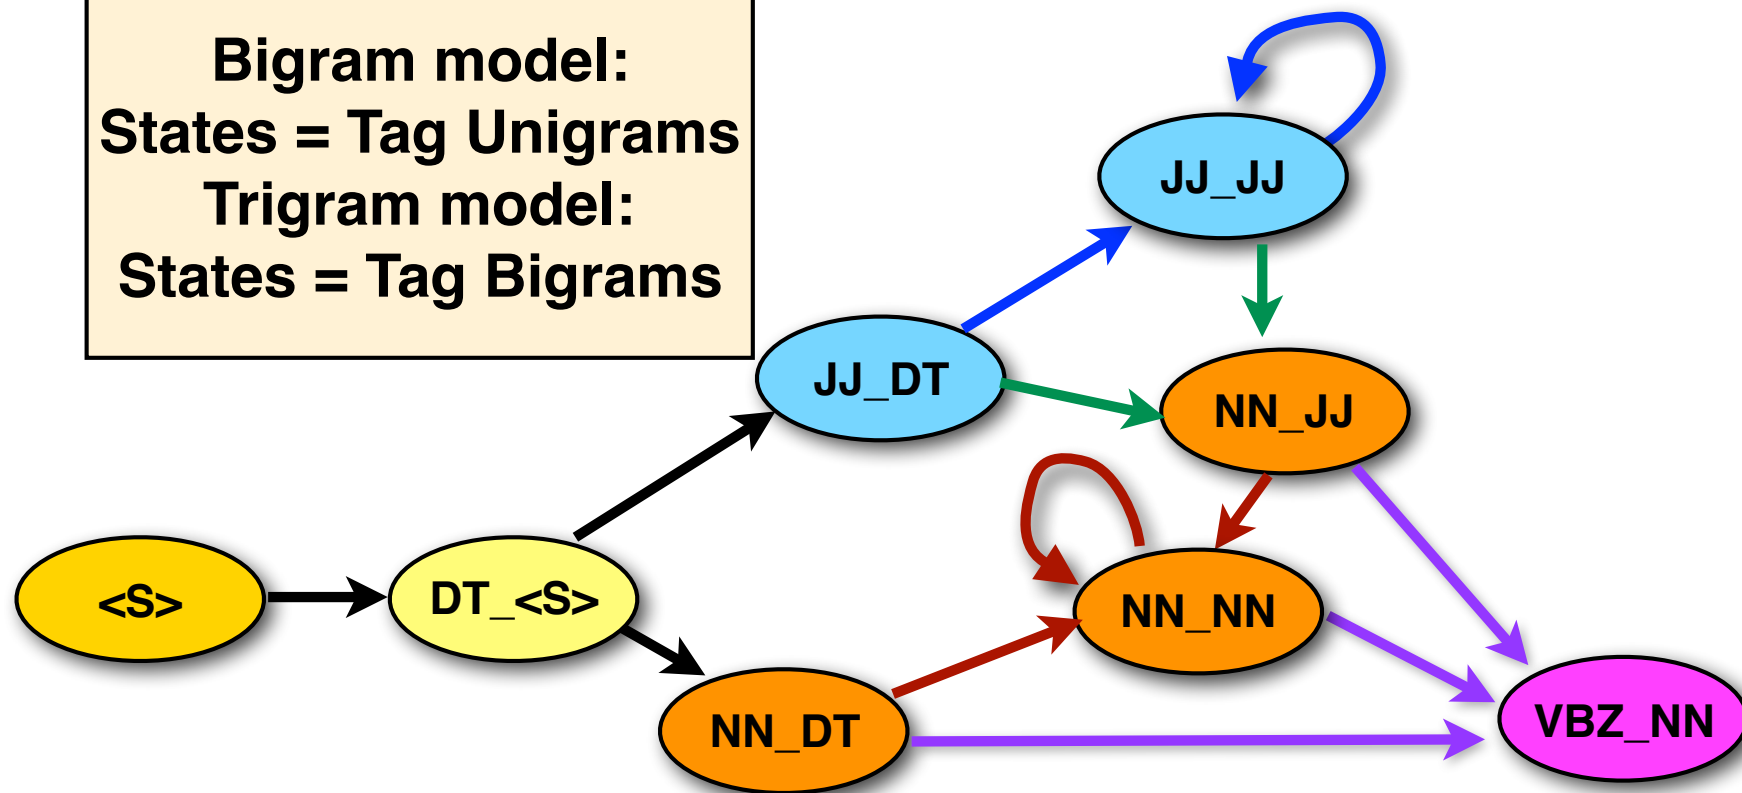
???



Encoding a trigram model as FSA



Bigram model:
States = Tag Unigrams
Trigram model:
States = Tag Bigrams



HMM definition

A HMM $\lambda = (A, B, \pi)$ consists of

- a set of N **states** $Q = \{q_1, \dots, q_N\}$
with $Q_0 \subseteq Q$ a set of **initial states**
and $Q_F \subseteq Q$ a set of **final (accepting) states**
- an **output vocabulary** of M items $V = \{v_1, \dots, v_m\}$
- an $N \times N$ **state transition probability matrix** A
with a_{ij} the probability of moving from q_i to q_j .
($\sum_{j=1}^N a_{ij} = 1 \ \forall i; \ 0 \leq a_{ij} \leq 1 \ \forall i, j$)
- an $N \times M$ **symbol emission probability matrix** B
with b_{ij} the probability of emitting symbol v_j in state q_i
($\sum_{j=1}^M b_{ij} = 1 \ \forall i; \ 0 \leq b_{ij} \leq 1 \ \forall i, j$)
- an **initial state distribution vector** $\pi = \langle \pi_1, \dots, \pi_N \rangle$
with π_i the probability of being in state q_i at time $t = 1$.
($\sum_{i=1}^N \pi_i = 1 \ 0 \leq \pi_i \leq 1 \ \forall i$)

Learning an HMM

Where do we get the transition probabilities $P(t_j | t_i)$ (matrix A) and the emission probabilities $P(w_j | t_i)$ (matrix B) from?

Case 1: We have a POS-tagged corpus.

- This is learning from labeled data, aka “supervised learning”

```
Pierre_NNP Vinken_NNP ,_, 61_CD years_NNS  
old_JJ ,_, will_MD join_VB the_DT board_NN  
as_IN a_DT nonexecutive_JJ director_NN Nov._NNP  
29_CD ._. 
```

Case 2: We have a raw (untagged) corpus and a tagset.

- This is learning from unlabeled data, aka “unsupervised learning”

```
Pierre Vinken , 61 years old , will  
join the board as a nonexecutive  
director Nov. 29 .
```

Tagset:
NNP: proper noun
CD: numeral,
JJ: adjective,...

Learning an HMM from *labeled* data

```
Pierre_NNP Vinken_NNP ,_, 61_CD years_NNS  
old_JJ ,_, will_MD join_VB the_DT board_NN  
as_IN a_DT nonexecutive_JJ director_NN Nov._NNP  
29_CD ._.
```

We **count** how often we see t_it_j and $w_j_t_i$ etc. in the data (use relative frequency estimates):

Learning the transition probabilities:

$$P(t_j|t_i) = \frac{C(t_it_j)}{C(t_i)}$$

Learning the emission probabilities:

$$P(w_j|t_i) = \frac{C(w_j_t_i)}{C(t_i)}$$

We might use some smoothing, but this is pretty trivial...

Outlook: Dynamic Programming for HMMs

The three basic problems for HMMs

We observe an **output sequence** $w = w_1 \dots w_N$:

$w = \text{"she promised to back the bill"}$

Problem I (Likelihood): find $P(w \mid \lambda)$

Given an HMM $\lambda = (A, B, \pi)$, compute the likelihood of the observed output, $P(w \mid \lambda)$

Problem II (Decoding): find $Q = q_1 \dots q_T$

Given an HMM $\lambda = (A, B, \pi)$, what is the most likely sequence of states $Q = q_1 \dots q_N \approx t_1 \dots t_N$ to generate w ?

Problem III (Estimation): find $\operatorname{argmax}_{\lambda} P(w \mid \lambda)$

Find the parameters A, B, π which maximize $P(w \mid \lambda)$

How can we solve these problems?

I. Likelihood of the input:

Compute $P(\mathbf{w} \mid \lambda)$ for the input \mathbf{w} and HMM λ

II. Decoding (= tagging) the input:

Find the best tags $\mathbf{t}^* = \operatorname{argmax}_{\mathbf{t}} P(\mathbf{t} \mid \mathbf{w}, \lambda)$ for the input \mathbf{w} and HMM λ

III. Estimation (= learning the model):

Find the best model parameters $\lambda^* = \operatorname{argmax}_{\lambda} P(\mathbf{t}, \mathbf{w} \mid \lambda)$
for the training data \mathbf{w}

These look like hard problems: With T tags, every input string $\mathbf{w}_{1\dots n}$ has T^n possible tag sequences

Can we find efficient (polynomial-time) algorithms?

Dynamic programming

We will use a general technique called **dynamic programming** to solve these problems.

- We will **recursively decompose** each of these problems into smaller subproblems that can be solved efficiently
- There is only a **polynomial number of subproblems**.
- We will **store the solution of each subproblem** in a common data structure
- Processing this data structure takes polynomial time

Solution: Dynamic programming

I. Likelihood of the input:

Compute $P(\boldsymbol{w} | \lambda)$ for the input \boldsymbol{w} and HMM λ

⇒ **Forward algorithm**

II. Decoding (=tagging) the input:

Find best tags $\boldsymbol{t}^* = \operatorname{argmax}_{\boldsymbol{t}} P(\boldsymbol{t} | \boldsymbol{w}, \lambda)$ for the input \boldsymbol{w} and HMM λ

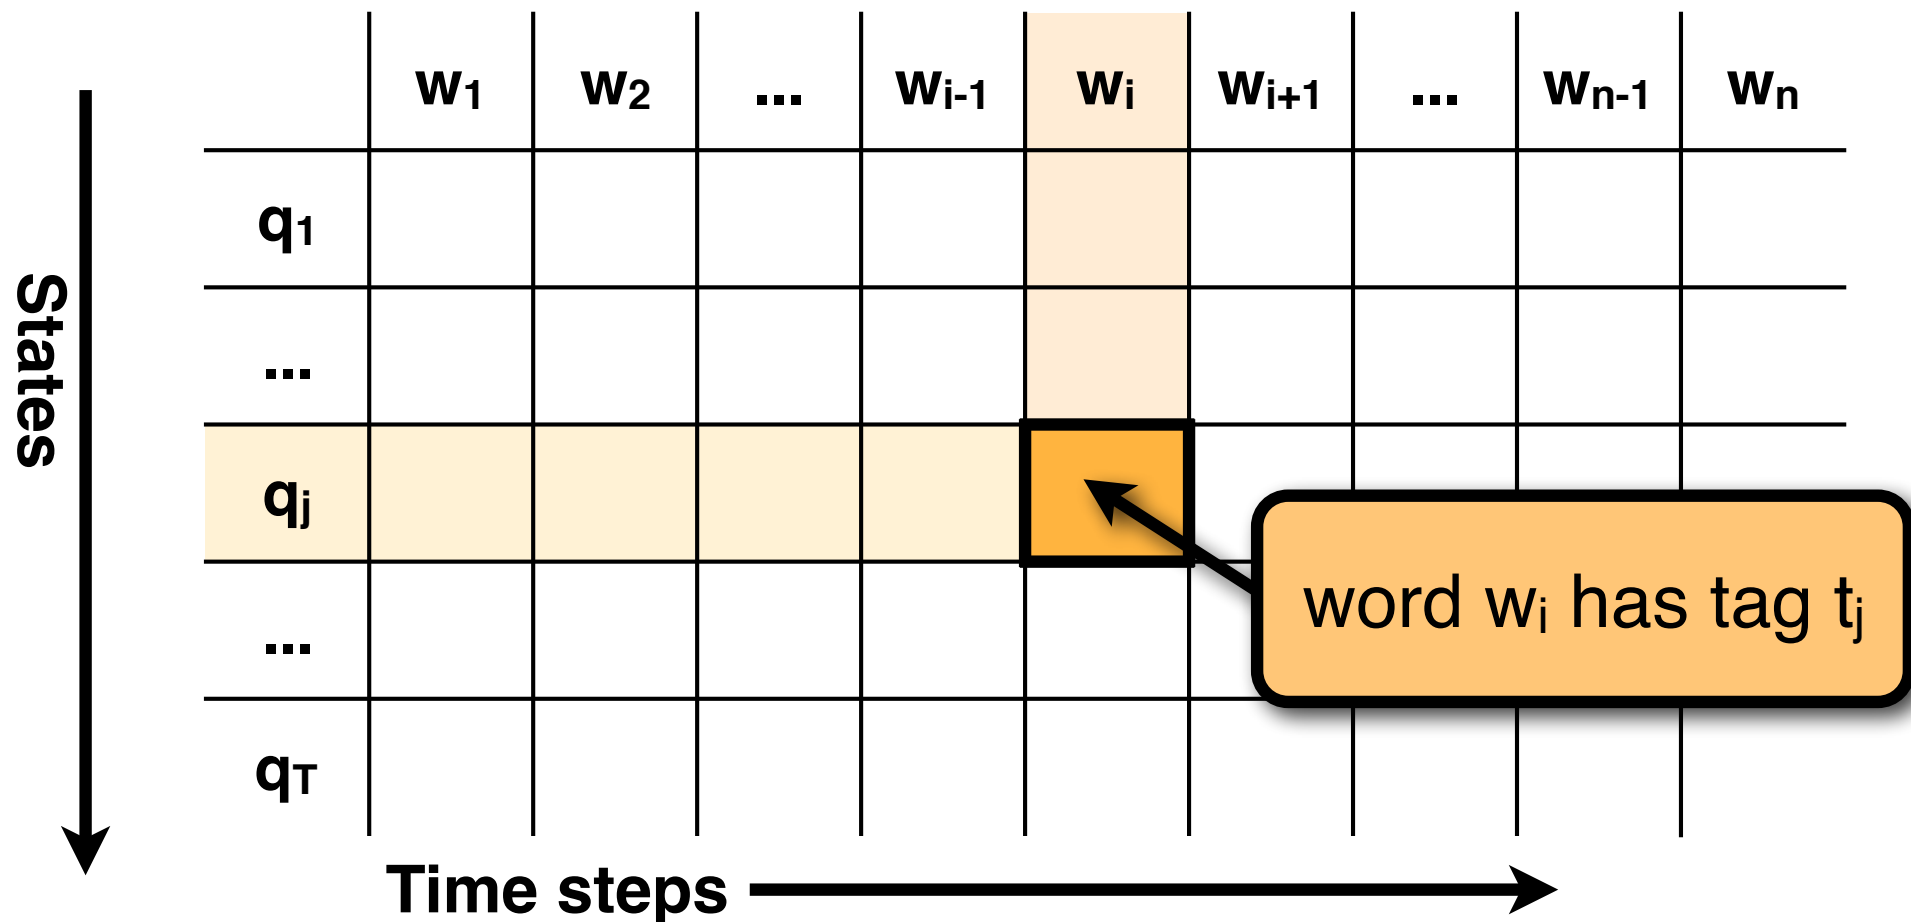
⇒ **Viterbi algorithm**

III. Estimation (=learning the model):

Find best model parameters $\lambda^* = \operatorname{argmax}_{\lambda} P(\boldsymbol{t}, \boldsymbol{w} | \lambda)$ for training data \boldsymbol{w}

⇒ **Forward-Backward algorithm**


Bookkeeping: the trellis



We use a $n \times T$ table ("**trellis**") to keep track of the HMM.

At every one of the n time steps (=words), the HMM can be in one of T states (=tags)

Computing $P(\mathbf{t}, \mathbf{w})$ for one tag sequence



- One path through the trellis = one tag sequence
- We just multiply the probabilities as before

$$P(\mathbf{t}, \mathbf{w}) = P(t_1)P(w_1|t_1) \prod_{i=2}^T P(t_i|t_{i-1})P(w_i|t_i)$$

Appendix:

English parts of speech

Nouns

Nouns describe entities and concepts:

Common nouns: dog, bandwidth, dog, fire, snow, information

- Count nouns have a plural (dogs) and need an article in the singular (the dog barks)
- Mass nouns don't have a plural (*snows) and don't need an article in the singular (snow is cold, metal is expensive). But some mass nouns can also be used as count nouns: Gold and silver are metals.

Proper nouns (Names): Mary, Smith, Illinois, USA, France, IBM

Penn Treebank tags:

NN: singular or mass

NNS: plural

NNP: singular proper noun

NNPS: plural proper noun

(Full) verbs

Verbs describe activities, processes, events:

eat, write, sleep,

Verbs have different morphological forms:

infinitive (to eat), present tense (I eat), 3rd pers sg. present tense (he eats),
past tense (ate), present participle (eating), past participle (eaten)

Penn Treebank tags:

VB: infinitive (base) form

VBD: past tense

VBG: present participle

VBD: past tense

VCN: past participle

VBP: non-3rd person present tense

VBZ: 3rd person singular present tense

Adjectives

Adjectives describe properties of entities:

blue, hot, old, smelly,...

Adjectives have an...

... attributive use (modifying a noun):

the blue book

... and a predicative use (e.g. as argument of be):

The book is blue.

Many gradable adjectives also have a...

...comparative form: greater, hotter, better, worse

...superlative form: greatest, hottest, best, worst

Penn Treebank tags:

JJ: adjective JJR: comparative JJS: superlative

Adverbs

Adverbs describe properties of events/states.

- Manner adverbs: slowly (slower, slowest) fast, hesitantly,...
- Degree adverbs: extremely, very, highly....
- Directional and locative adverbs: here, downstairs, left
- Temporal adverbs: yesterday, Monday,...

Adverbs modify verbs, sentences, adjectives or other adverbs:

Apparently, the very ill man walks extremely slowly

NB: certain temporal and locative adverbs (yesterday, here)
can also be classified as nouns

Penn Treebank tags:

RB: adverb RBR: comparative adverb RBS: superlative adverb

Auxiliary and modal verbs

Copula: be with a predicate

She is a student. I am hungry. She was five years old.

Modal verbs: can, may, must, might, shall,...

She can swim. You must come

Auxiliary verbs:

- Be, have, will when used to form complex tenses:

He was being followed. She has seen him. We will have been gone.

- Do in questions, negation:

Don't go. Did you see him?

Penn Treebank tags:

MD: modal verbs

Prepositions

Prepositions occur before noun phrases to form a prepositional phrase (PP):

on/in/under/near/towards the wall,
with(out) milk,
by the author,
despite your protest

PPs can modify nouns, verbs or sentences:

I drink [coffee [with milk]]
I [drink coffee [with my friends]]

Penn Treebank tags:

IN: preposition

TO: 'to' (infinitival 'to eat' and preposition 'to you')

Conjunctions

Coordinating conjunctions conjoin two elements:

X and/or/but X

[[John]NP and [Mary]NP] NP,

[[Snow is cold]S but [fire is hot]S]S.

Subordinating conjunctions introduce a subordinate (embedded) clause:

[He thinks that [snow is cold]S]S

[She wonders whether [it is cold outside]S]S

Penn Treebank tags:

CC: coordinating

IN: subordinating (same as preposition)

Particles

Particles resemble prepositions (but are not followed by a noun phrase) and appear with verbs:

come on

he brushed himself off

turning the paper over

turning the paper down

Phrasal verb: a verb + particle combination that has a different meaning from the verb itself

Penn Treebank tags:

RP: particle

Pronouns

Many pronouns function like noun phrases, and refer to some other entity:

- Personal pronouns: I, you, he, she, it, we, they
- Possessive pronouns: mine, yours, hers, ours
- Demonstrative pronouns: this, that,
- Reflexive pronouns: myself, himself, ourselves
- Wh-pronouns (question words):
what, who, whom, how, why, whoever, which

Relative pronouns introduce relative clauses
the book that [he wrote]

Penn Treebank tags:

PRP: personal pronoun PRP\$ possessive WP: wh-pronoun

Determiners

Determiners precede noun phrases:

the/that/a/every book

- Articles: the, an, a
- Demonstratives: this, these, that
- Quantifiers: some, every, few,...

Penn Treebank tags:

DT: determiner