

CS 519: Scientific Visualization

Volume Visualization


Eric Shaffer

Some slides adapted Alexandru Telea, *Data Visualization Principles and Practice*

Some slides adapted from Travis Gorkin, *Volume Rendering using Graphics Hardware*



Volume Rendering Definition

- ❑ Generate 2D projection of 3D data set
- ❑ Visualization of medical and scientific data
- ❑ Rendering natural effects - fluids, smoke, fire
- ❑ Direct Volume Rendering (DVR)
 - ❑ Done without extracting any surface geometry



Volumetric Data

- ▣ 3D Data Set
 - ▣ Discretely sampled on regular grid in 3D space
 - ▣ 3D array of samples
- ▣ Voxel – volume element
 - ▣ One or more constant data values
 - ▣ Scalars – density, temperature, opacity
 - ▣ Vectors – color, normal, gradient
 - ▣ Spatial coordinates determined by position in the grid



8x8x8 resolution orthogonal (uniform) voxel grid

Volume Visualization

1. Motivation

- how to see through 3D scalar volumes?

2. Methods and techniques

- ray function (MIP, average intensity, distance to value, isosurface)
- classification
- compositing
- volumetric shading


3. Advanced points

- sampling and interpolation
 - classification and interpolation order
 - performance issues
-

Volume Visualization

Scalar volume $s : \mathbf{R}^3 \rightarrow \mathbf{R}$

How to visualize this?



direct color mapping

slicing

contouring

- see only outer surface
- all details on slice
- all details on contour
- no info outside slice
- no info outside contour


How to visualize this so we see through the volume

Idea

- use known techniques (slices and contours)
- use transparency

First try

- draw several contours C_i for several values s_i
- transparency α_i proportional to scalar value s




We start seeing a little bit through the volume...

...But this won't work for too many contours

isovalue = 65
isovalue = 127


Second try

- draw several parallel slices S_i
- transparency α_i inversely proportional to number of slices $\alpha_i = \frac{1}{\|S\|}$



axis-aligned slices

- not OK if we view volume across slicing direction



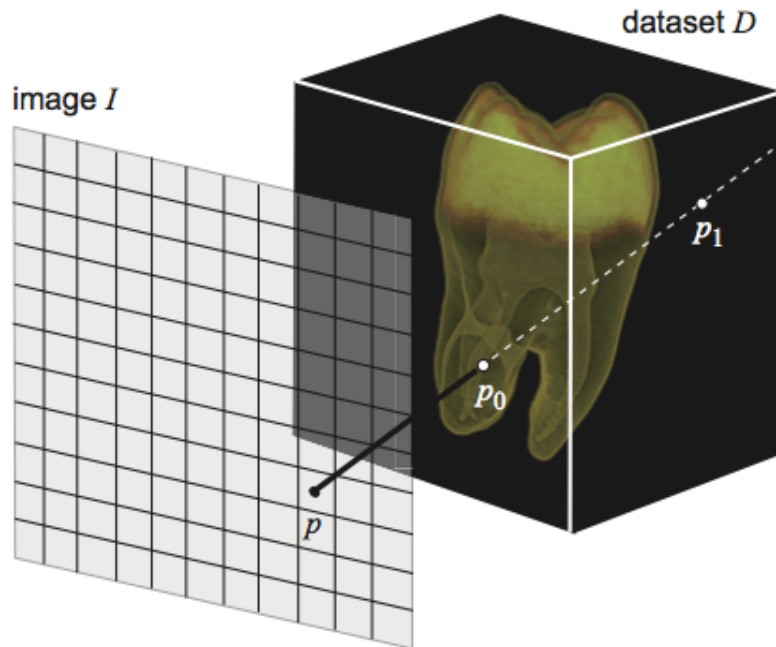
view direction-aligned slices

- any viewing direction OK
- must reslice when changing viewpoint

Volume Rendering Basics

Main idea

- consider a scalar signal $s : D \rightarrow \mathbf{R}$ to be drawn on the screen image I
- for each pixel $p \in I$
 - construct a ray \mathbf{r} orthogonal to I passing through p
 - compute intersection points p_0 and p_1 of \mathbf{r} with D
 - express $I(p)$ as function of s along \mathbf{r} between p_0 and p_1



1. Parameterize ray

$$p(t) = (1-t)p_0 + tp_1, \quad t \in [0,1]$$

1. Compute pixel color

$$I(p) = f(F(s(t))), \quad t \in [0,1]$$

transfer function

ray function

Volume Rendering

Define a **ray function**

$$F : \{s(t) \mid t \in [0,1]\} \rightarrow \mathbf{R}$$

all scalar values along ray

a single resulting scalar value

The ray function 'aggregates' all scalar values along a ray

Next, define a **transfer function** $f : \mathbf{R} \rightarrow [0,1]^4$

a single scalar value

an RGBA color

- same concept as color mapping (see Module 2)

Idea

- **ray function**: says how to combine all scalar values along a ray into a single value
- **transfer function**: says how to map a single scalar value to a color
- The process of computing all rays for an image I is called **ray casting**


Maximum Intensity Projection (MIP)

First example of ray function

- find maximum scalar along ray, then apply transfer function to its value

$$I(p) = f\left(\max_{t \in [0, T]} s(t)\right)$$

- useful to emphasize high-value points in the volume



Example MIP of human head CT

- white = low density (air)
- black = high density (bone)

OK, but gives no depth cues

Average Intensity Projection

Second example of ray function

- compute average scalar along ray, then map it to color

$$I(p) = f \left(\frac{\int_{t=0}^T s(t) dt}{T} \right)$$

- useful to emphasize average tissue type (e.g. density in a CT scan)



Example Human torso CT

- black = low density (air)
- white = high density (bone)

Average intensity projection
is equivalent to an X-ray

maximum intensity projection average intensity projection

Distance to Value Function

Third example of ray function

- compute distance along ray until a specific scalar value σ

$$I(p) = f \left(\min_{t \in [0, T], s(t) \geq \sigma} t \right)$$

- useful to emphasize depth where some specific tissue is located



distance to value
20



distance to value
50

Example
Human head CT

- black = low distance
- white = high distance

Isosurface Function

Fourth example of ray function

- compute whether a given isovalue σ exists along ray

$$I(p) = \begin{cases} f(\sigma), & \exists t \in [0, T], s(t) = \sigma, \\ I_0, & \text{otherwise.} \end{cases}$$

- produces same result as marching cubes, but with a higher accuracy



isosurface
(marching cubes)



isosurface
(software ray
casting)



isosurface
(hardware ray
casting)

Composite Function

Fifth example of ray function

- compute a color at each point along the ray (apply transfer function *first*)
- blend (compose) all colors to get the final pixel color (ray function=alpha blending)

$$I(p) = F(\{f(s(t)) \mid t \in [0, 1]\})$$

transfer function (applied to all pixels along ray)
ray function (blends all colors produced by transfer function along ray)

Idea

- **transfer function**: controls color+transparency of all material types
- **ray function**: blends together all material colors+transparencies along ray
- most powerful (but most computationally expensive) ray function
- allows huge range of effects (depending on type of transfer function)
- designing 'good' transfer functions is however non-trivial

Compositing Color Samples

- Over operator – back-to-front order

$$\hat{C}_i = C_i + (1 - A_i)\hat{C}_{i+1}$$

$$\hat{A}_i = A_i + (1 - A_i)\hat{A}_{i+1}$$

- Under operator – front-to-back order

$$\hat{C}_i = (1 - \hat{A}_{i-1})C_i + \hat{C}_{i-1}$$

$$\hat{A}_i = (1 - \hat{A}_{i-1})A_i + \hat{A}_{i-1}$$

Volumetric Shading

Shading

- is required if we compute e.g. isosurfaces
- but can also be useful for composite ray function

Method

- instead of simply using the colors $I(t) = f(s(t))$
- composite the shaded colors (see Chapter 2, Phong lighting)

$$I(t) = c_{\text{amb}} + c_{\text{diff}}(t) \max(-\mathbf{L} \cdot \mathbf{n}(t), 0) + c_{\text{spec}}(t) \max(-\mathbf{r} \cdot \mathbf{v}, 0)^\alpha$$


How to implement

- lighting coefficients c and light vector \mathbf{L} : user sets them as desired
- surface normal \mathbf{n} : compute from gradient of scalar value

(we did the same for isosurfaces, see Module 3)

Volumetric Shading

Results



no shading

diffuse lighting

diffuse and
specular
lighting

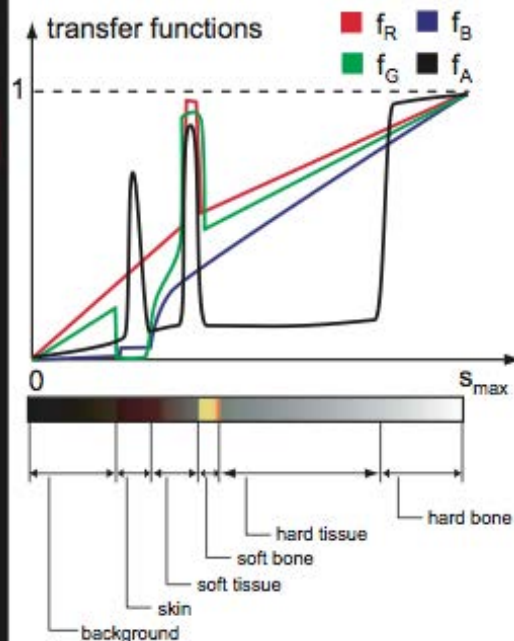

Shading

- gives very good cues of depth and shape structure
- is quite cheap and simple to compute

Transfer Functions

Extremely powerful modeling tool (mainly when using composite ray function)

- design four functions f_R, f_G, f_B, f_A
- use color and transparency to emphasize desired material properties (e.g. tissue type)
- use any ray function described so far




Example
Human head CT

- emphasize bone
- show also muscles

Volume Visualization of Vector Fields


Volume rendering can be used to visualize also **vector datasets**




Volume rendering of fluid flow vector field magnitude

- red = slow flow
- green = more rapid flow
- blue = fastest flow



Transfer Function Examples



(a)



(b)



- a) electron density
- b) car engine part
- c) bonsai tree (scanned)
- d) fish

Transfer Function Example



Volume rendering of human MRI dataset

- shading: mimics natural lighting
- backdrop added for extra effect

Implementation Issues


Sampling density

- recall the ray parameterization $q(t) = (1-t)q_0 + tq_1, \quad t \in [0,1]$
- we need to sample along the ray (e.g. integrate, compute min/max, etc)
- how small should we take the sampling step $\delta = dt$?



(a) $\delta = 0.1$

(b) $\delta = 0.5$



(c) $\delta = 1.0$


(d) $\delta = 2.0$

Human head CT, four different δ values

- smaller δ : more accuracy
- too small δ : slow rendering

Practical guideline

- δ should never exceed a voxel size (otherwise we skip voxels while traversing the ray...)



Volume Visualization Summary


Volume visualization (book Chapter 10)

- Extends classical scalar visualization to 'see through' 3D volumes
 - ray functions and transfer functions
- Evaluation
 - produces highly realistic, easy to interpret images
 - requires quite some computational power
 - can be easily accelerated using GPUs (e.g. pixel shaders, CUDA)
 - good transfer function design: critical, application-dependent, hard

Volume Rendering in More Depth


We take a closer look at how two fundamental problems are addressed

- How can we render at interactive speeds?
 - ▣ Image-order techniques (i.e. ray-casting)
 - ▣ Object-order techniques (using rasterization)
 - ▣ What optimizations can be made for each approach?
- How can we generate high-quality images?
 - ▣ What mathematical model do we use to render?
 - ▣ How does the model work with the approaches listed above?




A Volume Rendering Optical Model

- Light interacts with volume contents through:
 - ▣ Absorption
 - ▣ Emission
 - ▣ Scattering
- Sample volume along viewing rays
- Accumulate optical properties



Transfer Function


- ▣ Maps voxel data values to optical properties



- ▣ Glorified color maps
- ▣ Emphasize or classify features of interest in the data
- ▣ Piecewise linear functions, Look-up tables, 1D, 2D
- ▣ GPU – simple shader functions, texture lookup tables

Image Order Technique: Volume Ray Marching

1. Raycast – once per pixel
2. Sample – uniform intervals along ray
3. Interpolate – trilinear interpolate, apply transfer function
4. Accumulate – integrate optical properties



Ray Marching Accumulation Equations

- ▣ Accumulation = Integral
- ▣ Color

$$\overline{C} = \int_0^{\infty} \overline{C}_i T_i ds$$

Transmissivity = 1 - Opacity

$$T = 1 - A$$

Total Color = Accumulation (Sampled Colors x Sampled Transmissivities)

Ray Marching Accumulation Equations

- Discrete Versions
- Accumulation = Sum
- Color

$$\bar{C} = \sum_{i=1}^n \bar{C}_i T_i$$

Transmissivity = 1 - Opacity

$$T = 1 - A$$

- Opacity

$$A = 1 - \prod_{j=1}^n (1 - A_j)$$

$$C = \sum_{i=1}^n C_i \prod_{j=1}^{i-1} (1 - A_j)$$

Compositing Color Samples

- Over operator – back-to-front order

$$\hat{C}_i = C_i + (1 - A_i)\hat{C}_{i+1}$$

$$\hat{A}_i = A_i + (1 - A_i)\hat{A}_{i+1}$$

- Under operator – front-to-back order

$$\hat{C}_i = (1 - \hat{A}_{i-1})C_i + \hat{C}_{i-1}$$

$$\hat{A}_i = (1 - \hat{A}_{i-1})A_i + \hat{A}_{i-1}$$

CPU Based Volume Rendering

- Raycast and raymarch for each pixel in scene

- Camera (eye) location: x_C

- For Each Pixel

- Look Direction: \hat{n}

- Cast Ray Along: $x_C + \hat{n}s$

- Accumulate Color Along Line

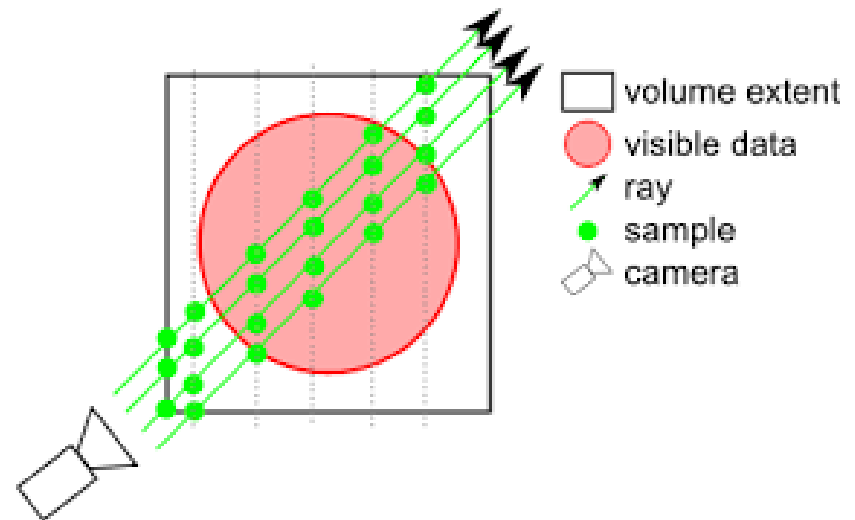
- Sequential or coarse-grained parallel process

- Minutes or Hours per frame

- Optimizations


- Space Partitioning

- Early Ray Termination

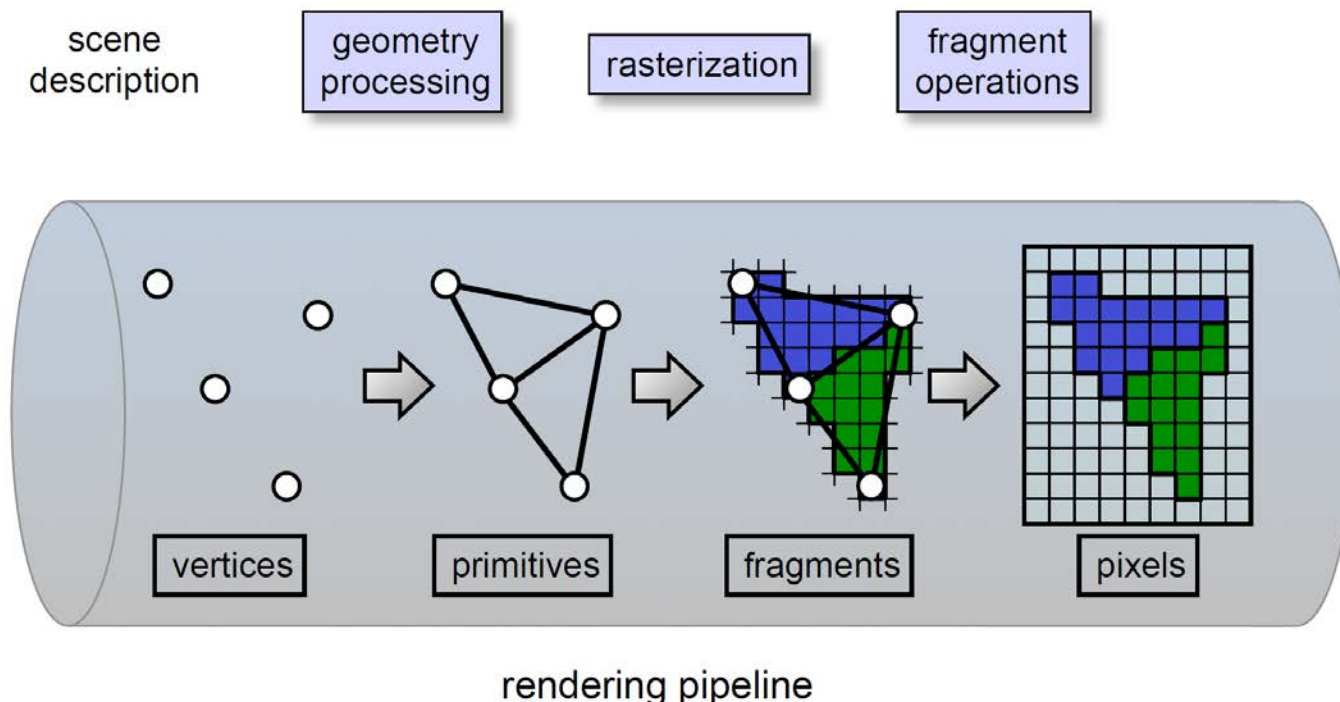


Using Rasterization for Volume Rendering

- ▣ How can we speed things up?
- ▣ Modern rasterization engines are designed for interactive speed
 - ▣ OpenGL, D3D, Vulkan
- ▣ It's possible to use that technology for volume rendering
 - ▣ We can render polygons and composite them
 - ▣ Each polygon will be a slice through the volume
 - ▣ So is this object-order or image-order rendering?



A Modern Rasterization Engine in Two Slides

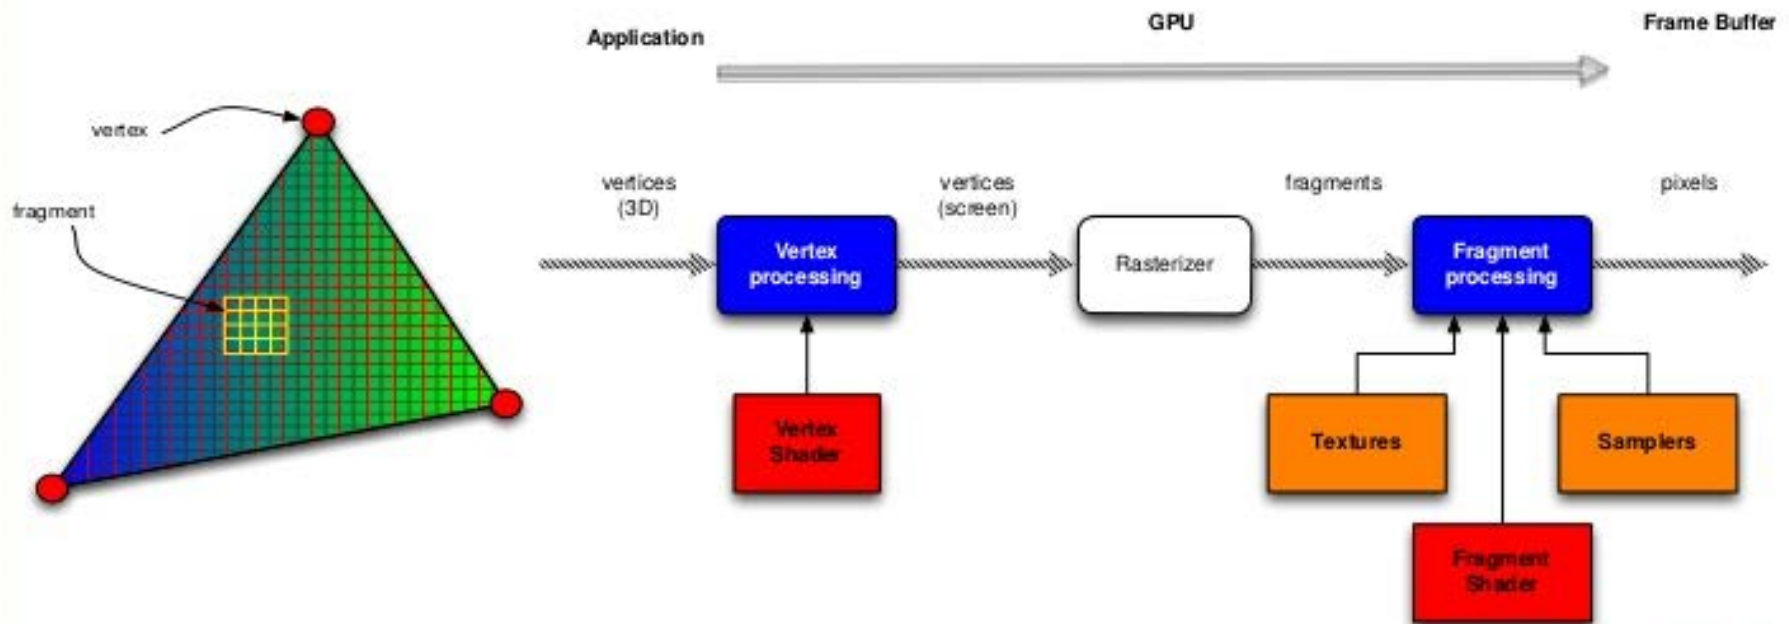


A Modern Rasterization Engine in Two Slides


WebGL pipeline



- Programmable vertex & fragment shaders




How 2D Texture Mapping Works in One Slide



- Each fragment (pixel) generated by polygon is colored using a 2D image as a lookup table
- Some APIs support 3D textures...similar techniques used to map colors onto polygons


Slice-Based Volume Rendering (SBVR)

- ❑ No volumetric primitive in graphics API
- ❑ Proxy geometry - polygon primitives as slices through volume
- ❑ Texture polygons with volumetric data
- ❑ Draw slices in sorted order – back-to-front
- ❑ Use fragment shader to perform compositing (blending)



Volumetric Data


- ❑ Voxel data sent to GPU memory as
 - ❑ Stack of 2D textures
 - ❑ 3D texture
- ❑ Leverage graphics pipeline




Proxy Geometry

Object-Aligned Slices

- Fast and simple
- Three stacks of 2D textures – x, y, z principle directions



- Texture stack swapped based on closest to viewpoint




Proxy Geometry

- ▣ Issues with Object-Aligned Slices
 - ▣ 3x memory consumption
 - ▣ Data replicated along 3 principle directions
 - ▣ Change in viewpoint results in stack swap
 - ▣ Image popping artifacts
 - ▣ Lag while downloading new textures
 - ▣ Sampling distance changes with viewpoint
 - ▣ Intensity variations as camera moves


Proxy Geometry

- View-Aligned Slices
 - Slower, but more memory efficient
 - Consistent sampling distance

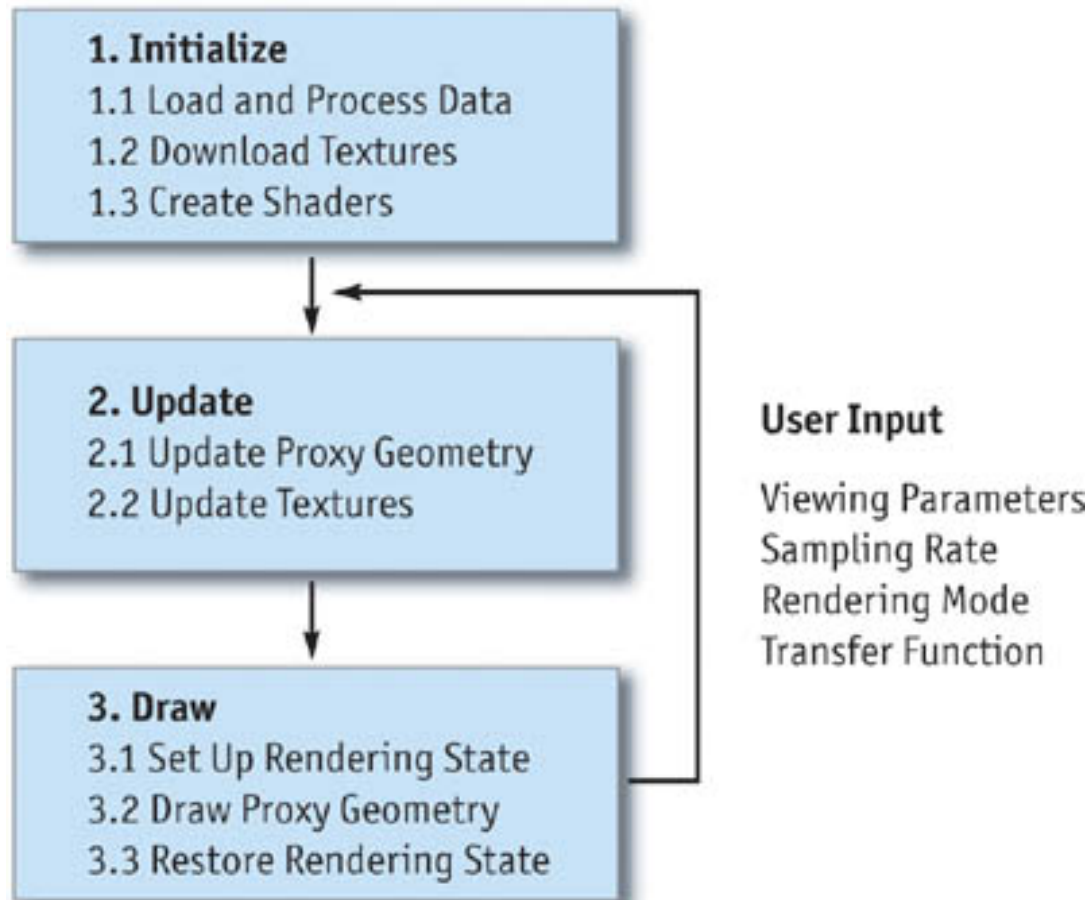


Proxy Geometry


- View-Aligned Slices Algorithm
 - Intersect slicing planes with bounding box
 - Sort resulting vertices in (counter)clockwise order
 - Construct polygon primitive from centroid as triangle fan




Sliced-Based Volume Rendering Steps



Slice-based Volume Rendering



2D textures
axis-aligned



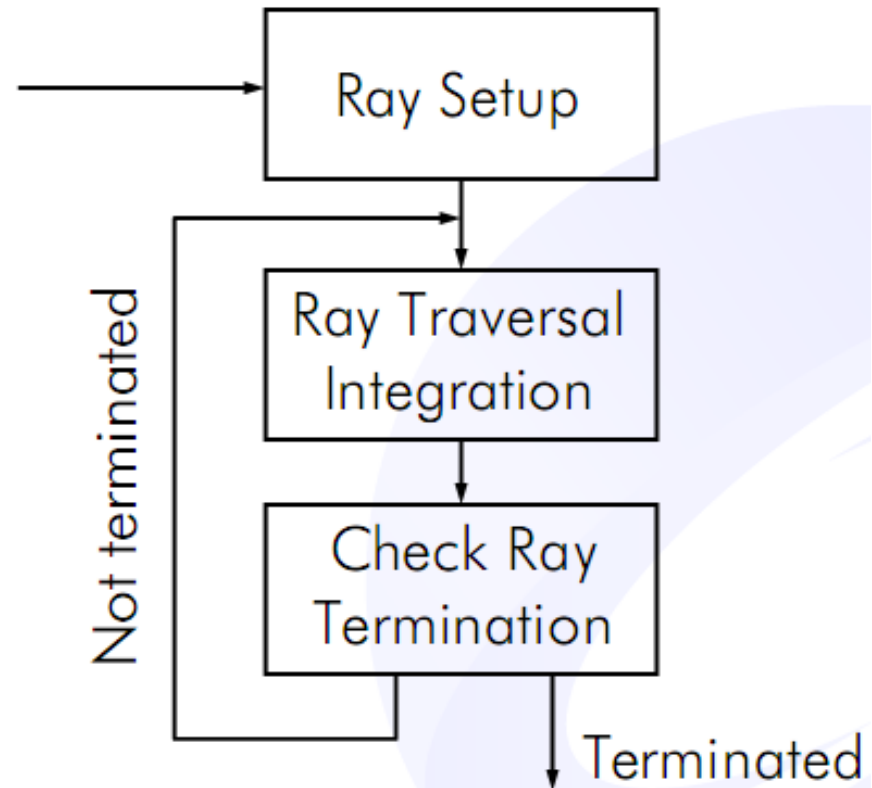
3D texture
view-aligned

Volume Raycasting on GPU

- ▣ Alternative: Use the GPU to cast rays through volume
 - ▣ Image or Object order?
- ▣ Raymarching implemented in fragment shader
 - ▣ Cast rays of through volume
 - ▣ Accumulate color and opacity
 - ▣ Terminate when opacity reaches threshold


Volume Raycasting on GPU

- ▣ Multi-pass algorithm
- ▣ Initial passes
 - ▣ Precompute ray directions and lengths
- ▣ Additional passes
 - ▣ Perform raymarching in parallel for each pixel
 - ▣ Split up full raymarch to check for early termination




Step 1: Ray Direction Computation

- ▣ Ray direction computed for each pixel
- ▣ Stored in 2D texture for use in later steps
- ▣ Pass 1: Front faces of volume bounding box
- ▣ Pass 2: Back faces of volume bounding box
- ▣ Vertex color components encode object-space principle directions



Step 1: Ray Direction Computation

- ▣ Subtraction blend two textures
- ▣ Store normalized direction – RGB components
- ▣ Store length – Alpha component



Fragment Shader Raymarching

- ▶ $DIR[x][y]$ – ray direction texture
 - ▶ 2D RGBA values
- ▶ P – per-vertex float3 positions, front of volume bounding box
 - ▶ Interpolated for fragment shader by graphics pipeline
- ▶ s – constant step size
 - ▶ Float value
- ▶ d – total raymarched distance, $s \times \text{\#steps}$
 - ▶ Float value

- ▶ Parametric Ray Equation

$$\mathbf{r} = \mathbf{P} + d \cdot \mathbf{DIR}[x][y]$$

- ▶ \mathbf{r} – 3D texture coordinates used to sample voxel data

Fragment Shader Raymarching

- ▶ Ray traversal procedure split into multiple passes
 - ▶ M steps along ray for each pass
 - ▶ Allows for early ray termination, optimization
- ▶ Optical properties accumulated along M steps
 - ▶ Simple compositing/blending operations
 - ▶ Color and alpha(opacity)
- ▶ Accumulation result for M steps blended into 2D result texture
 - ▶ Stores overall accumulated values between multiple passes
- ▶ Intermediate Pass – checks for early termination
 - ▶ Compare opacity to threshold
 - ▶ Check for ray leaving bounding volume

Improving Image Quality

- ▣ Local illumination using Blinn-Phong illumination model
- ▣ Volumetric shadows

Local Illumination

- Blinn-Phong Shading Model

$$I = k_a + I_L k_d (\hat{l} \cdot \hat{n}) + I_L k_s (\hat{h} \cdot \hat{n})^N$$

Resulting = Ambient + Diffuse + Specular

- Requires surface normal vector
 - Whats the normal vector of a voxel?

Local Illumination

□ Blinn-Phong Shading Model

$$I = k_a + I_L k_d (\hat{l} \cdot \hat{n}) + I_L k_s (\hat{h} \cdot \hat{n})^N$$

Resulting = Ambient + Diffuse + Specular

□ Requires surface normal vector

▣ Whats the normal vector of a voxel? **Gradient**

▣ Central differences between neighboring voxels

$$\text{grad}(I) = \nabla I = \frac{(\text{right} - \text{left})}{2x}, \frac{(\text{top} - \text{bottom})}{2x}, \frac{(\text{front} - \text{back})}{2x}$$

Local Illumination

- ❑ Compute on-the-fly within fragment shader
 - ❑ Requires 6 texture fetches per calculation
- ❑ Precalculate on host and store in voxel data
 - ❑ Requires 4x texture memory
 - ❑ Pack into 3D RGBA texture to send to GPU

Voxel Data

- X Gradient
- Y Gradient
- Z Gradient
- Value




3D Texture

- R
- G
- B
- A

Local Illumination

- ▣ Improve perception of depth
- ▣ Amplify surface structure



Volumetric Shadows


- Light attenuated as passes through volume
- 'Deeper' samples receive less illumination
- Second raymarch from sample point to light source
 - Accumulate illumination from sample's point of view
 - Same accumulation equations
- Precomputed Light Transmissivity
 - Precalculate illumination for each voxel center
 - Trilinearly interpolate at render time
 - View independent, scene/light source dependent

Volumetric Shadows on GPU

- Light attenuated from light's point of view
- CPU – Precomputed Light Transfer
 - Secondary raymarch from sample to light source
- GPU
 - Two-pass algorithm
 - Modify proxy geometry slicing
 - Render from both the eye and the light's POV
 - ▣ Two different frame buffers

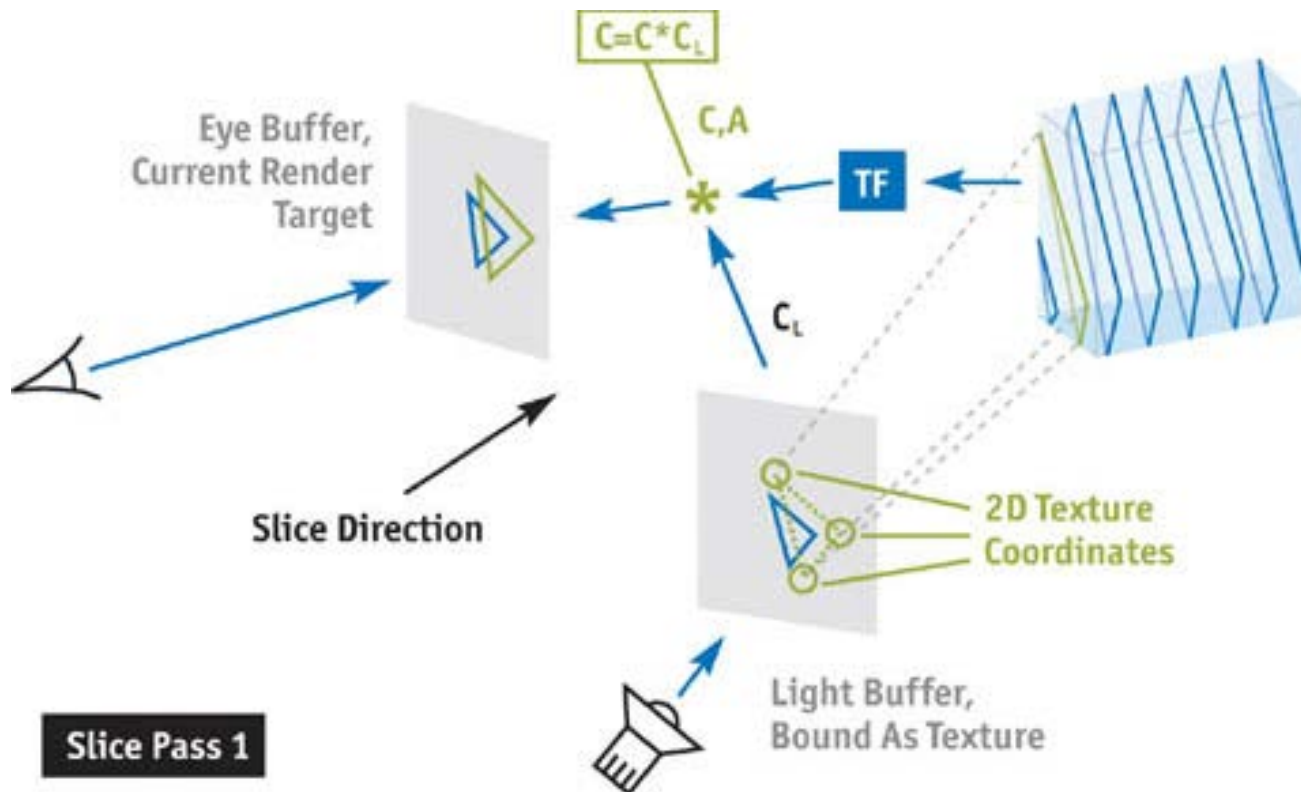
Two Pass Volume Rendering with Shadows

- ❑ Slice axis set half-way between view and light directions
 - ❑ Allows each slice to be rendered from eye and light POV
- ❑ Render order for light – front-to-back
- ❑ Render order for eye – (a) front-to-back
(b) back-to-front




First Pass

- Render from eye
- Fragment shader
 - Look up light color from light buffer bound as texture
 - Multiply material color * light color



Second pass

- Render from light
- Fragment shader
 - Only blend alpha values – light transmissivity



Volumetric Shadows



(a)



(b)

Volume Rendering in CUDA

- ❑ 3D Slicer – www.slicer.org
- ❑ Open source software for visualization and image analysis
- ❑ Funded by NIH, medical imaging, MRI data
- ❑ Currently integrating CUDA volume rendering into Slicer 3.2

