

Sparsity, Compressive Sensing and Random Projections

3 November 2017

Today's lecture

- Sparsity
- Compressive sensing
- Quantization, randomness and high dimensions

What is sparsity?

- Depends who you ask
- Basic idea: We want most numbers in a collection to be zero
- Too many ways to express that

A starting point

- Linear equation with multiple solutions:

$$y = \mathbf{a} \cdot \mathbf{x} \Rightarrow 2 = \begin{bmatrix} 1 & 1 \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

- Which solution would you pick?

$$\mathbf{x} = \begin{bmatrix} 2 \\ 0 \end{bmatrix}$$

A sparse answer

$$\mathbf{x} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$


What MATLAB gives

$$\mathbf{x} = \begin{bmatrix} 4 \\ -2 \end{bmatrix}$$

This one is fine too!

Infinite solutions

- All solutions lie on a line
- Which one do we pick?
 - Why did MATLAB pick $[1, 1]$?
- Does it make a difference?



The generic answer

- Least squares problem:

$$\mathbf{y} = \mathbf{A} \cdot \mathbf{x} \Rightarrow \mathbf{x} = \mathbf{A}^+ \cdot \mathbf{y}$$
$$\Rightarrow \arg \min_{\mathbf{x}} \left(\left\| \mathbf{A} \cdot \mathbf{x} - \mathbf{y} \right\|_2 + \left\| \mathbf{x} \right\|_2 \right)$$

- Find the minimum-norm \mathbf{x} that minimizes the error
 - But why $\left\| \mathbf{x} \right\|_2$?

Least squares, pseudoinverse, and ℓ_2

- Use Lagrangian multipliers:

$$\left\| \mathbf{x} \right\|_2^2 + \lambda^\top \cdot (\mathbf{A} \cdot \mathbf{x} - \mathbf{y}) \Rightarrow \hat{\mathbf{x}} = -\frac{1}{2} \mathbf{A}^\top \cdot \lambda$$

- Put back in original equation:

$$\mathbf{A} \cdot \hat{\mathbf{x}} = -\frac{1}{2} \mathbf{A} \cdot \mathbf{A}^\top \cdot \lambda = \mathbf{y} \Rightarrow \lambda = -2 \left(\mathbf{A} \cdot \mathbf{A}^\top \right)^{-1} \cdot \mathbf{y}$$

$$\Rightarrow \hat{\mathbf{x}} = \mathbf{A}^\top \cdot \left(\mathbf{A} \cdot \mathbf{A}^\top \right)^{-1} \cdot \mathbf{y} = \mathbf{A}^+ \cdot \mathbf{y}$$

Many more norms


- p -Norm (or L_p / L_p / ℓ_p)

$$\|\mathbf{x}\|_p = \sqrt[p]{\sum_i |x_i|^p}$$


- ℓ_2 norm is the Euclidean norm
- ℓ_1 norm is sum of absolute values
- ℓ_0 norm is the number of non-zero values
- ℓ_∞ norm is max of all values

How do they look?

- Unit norms in 2D





- Unit norms in 3D







ℓ_2 -based pseudoinverse

- All possible solutions lie on a hyperplane
 - Minimum ℓ_2 solution will be the point where the smallest possible ℓ_2 -ball touches the solutions hyperplane



Other ℓ_p -norm solutions

- With different norms the solution will change
 - Larger p will produce a “busier” solution
 - Smaller p will produce sparser solution



Which one to use?

- For sparsity we ideally we want minimum ℓ_0
 - Directly results in smallest number of non-zero values
- But, this is an inconvenient form ...




$$\ell_0(\mathbf{x}) = \sum_i \left[x_i \neq 0 \right] \quad \leftarrow \begin{array}{l} \text{Iverson bracket:} \\ \text{if content evaluates} \\ \text{to true it returns 1} \\ \text{otherwise it returns 0} \end{array}$$

- Discontinuous, no derivative, not convex, etc ...

$$\frac{\partial \|\mathbf{x}\|_0}{\partial \mathbf{x}} = ? \longrightarrow \boxed{\|\mathbf{x}\|_0} + \lambda^\top \cdot (\mathbf{A} \cdot \mathbf{x} - \mathbf{y})$$


Let's try something simpler then

- How about using the ℓ_1 instead?
 - Seems to produce the same solution



Why does this work?

- The ℓ_1 case is (sort of) convex
 - ℓ_1 ball shrinks as we move towards the ideal sparse solution
 - There's one ill-defined scenario, but it isn't a big problem
 - Which is it?
- So let's solve that instead



The problem to solve

- We now have:

$$\arg \min_{\mathbf{x}} \left(\left\| \mathbf{A} \cdot \mathbf{x} - \mathbf{y} \right\|_2 + \left\| \mathbf{x} \right\|_1 \right)$$

- Minor glitch: We can't differentiate the absolute values in the ℓ_1 norm!

$$\sum |x_i| + \lambda^\top \cdot (\mathbf{A} \cdot \mathbf{x} - \mathbf{y})$$

- But we can use other tools

Linear programming

- A linear program is defined as:

$$\text{minimize } \mathbf{c}^T \cdot \mathbf{x}$$

$$\text{subject to } \mathbf{A} \cdot \mathbf{x} \leq \mathbf{y}$$

$$\text{and } \mathbf{x} \geq 0$$

- A Nobel-prize staple of optimization theory, resource allocation, economics, etc.

Doesn't exactly match the ℓ_1 problem

- We would like to change our problem definition to fit the linear programming formulation

What we have

$$\begin{array}{ll} \text{minimize} & \|\mathbf{x}\|_1 \\ \text{subject to} & \mathbf{A} \cdot \mathbf{x} = \mathbf{y} \end{array}$$

What we can solve

$$\begin{array}{ll} \text{minimize} & \mathbf{c}^\top \cdot \mathbf{x} \\ \text{subject to} & \mathbf{A} \cdot \mathbf{x} \leq \mathbf{y} \\ \text{and} & \mathbf{x} \geq 0 \end{array}$$

With some shuffling around

- We rewrite the unknown vector \mathbf{x} as a difference of positive-valued vectors:

$$\mathbf{x} = \mathbf{u} - \mathbf{v}, \quad u_i, v_i \geq 0, \quad \mathbf{z} = \begin{bmatrix} \mathbf{u} \\ \mathbf{v} \end{bmatrix}$$

- Now our problem can be written as:

$$\mathbf{A} \cdot \mathbf{x} = \mathbf{y} \Rightarrow \begin{bmatrix} \mathbf{A} & -\mathbf{A} \end{bmatrix} \cdot \begin{bmatrix} \mathbf{u} \\ \mathbf{v} \end{bmatrix} = \mathbf{y} \Rightarrow \begin{bmatrix} \mathbf{A} & -\mathbf{A} \end{bmatrix} \cdot \mathbf{z} = \mathbf{y}$$

Now it is a linear program

- And we can solve our problem

Minimum ℓ_1 problem


$$\begin{array}{ll}\text{minimize} & \|\mathbf{x}\|_1 \\ \text{subject to} & \mathbf{A} \cdot \mathbf{x} = \mathbf{y}\end{array}$$

Equivalent linear program

$$\begin{array}{ll}\text{minimize} & \|\mathbf{z}\|_1 = \mathbf{1}^\top \cdot \mathbf{z} \\ \text{subject to} & [\mathbf{A}, -\mathbf{A}] \cdot \mathbf{z} = \mathbf{y} \\ \text{and} & \mathbf{z} \geq 0\end{array}$$

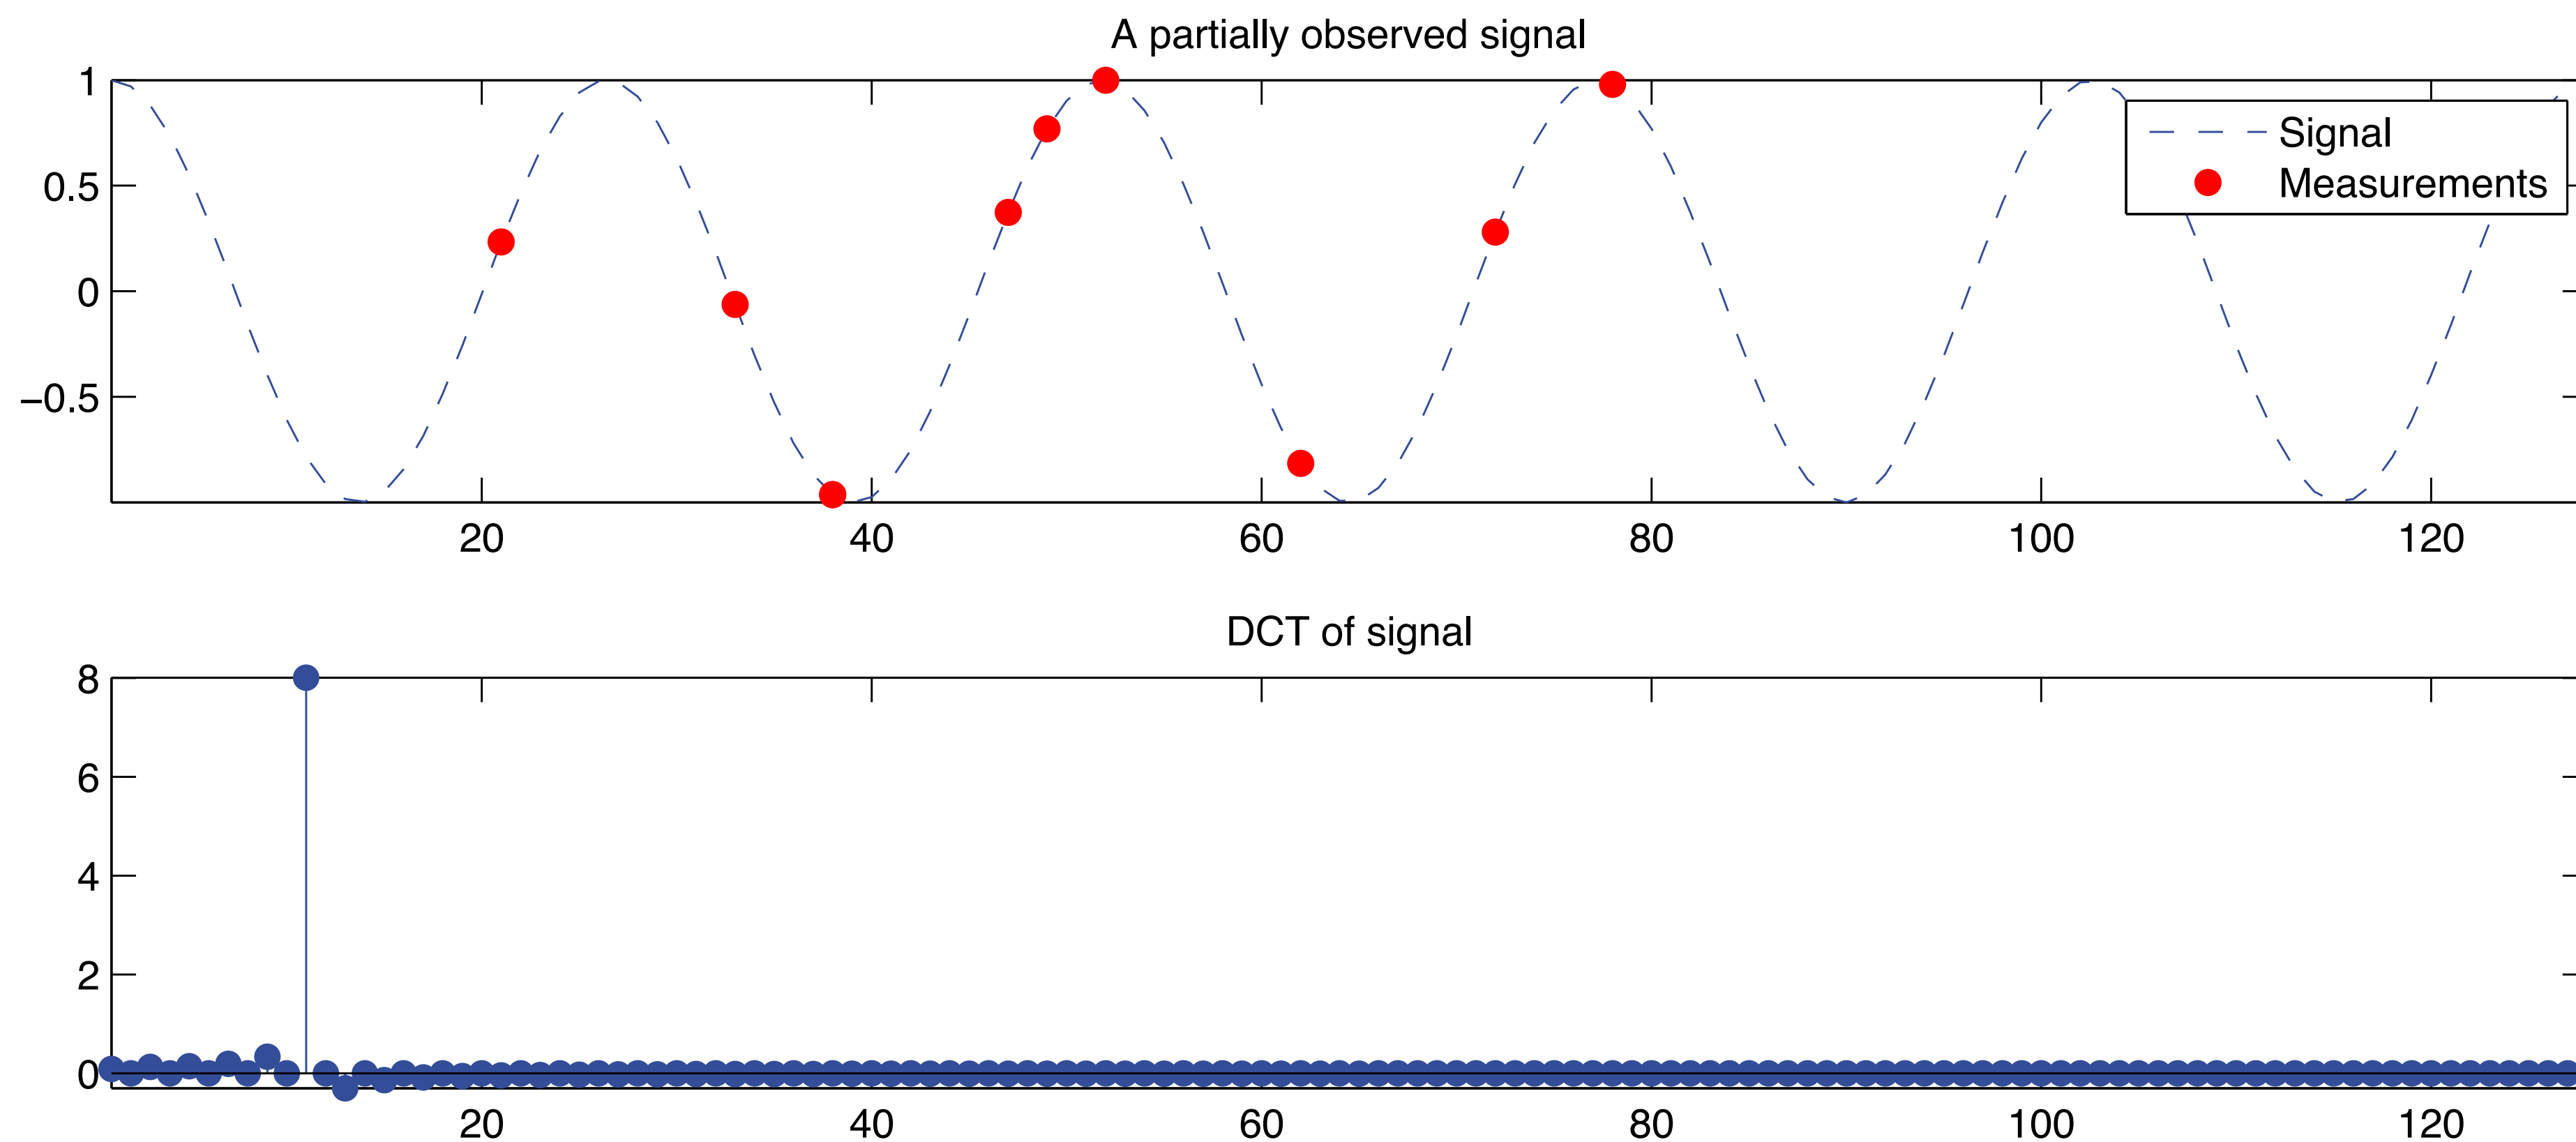
A simple example

- Suppose you measure this sinusoid
 - Can you recover the original signal using only the small number of measurements that we made?



Key observation

- The original signal is sparse in the frequency domain
 - How about we use that to construct a problem?



The problem to solve

- Find a set of small coefficients \mathbf{x} in the DCT domain, and make sure that they explain all our data \mathbf{y} , i.e.:


$$\mathbf{A} \cdot \mathbf{C}^{-1} \cdot \mathbf{x} = \mathbf{y}$$

- Where matrix \mathbf{A} selects only the indices that we observe
- Simple least squares problem using minimum ℓ_2 \mathbf{x}

$$\mathbf{x} = \left(\mathbf{A} \cdot \mathbf{C}^{-1} \right)^+ \cdot \mathbf{y}$$

And it doesn't really do much

- But you expected that!




What happened?

- Minimizing the ℓ_2 resulted in adding more frequencies in the signal
 - ℓ_2 doesn't give sparsity
 - Small coefficients \neq min ℓ_2
- We instead should find a minimal ℓ_1 solution
 - Because it actually enforces sparsity

And the result

- A much better reconstruction



A realization



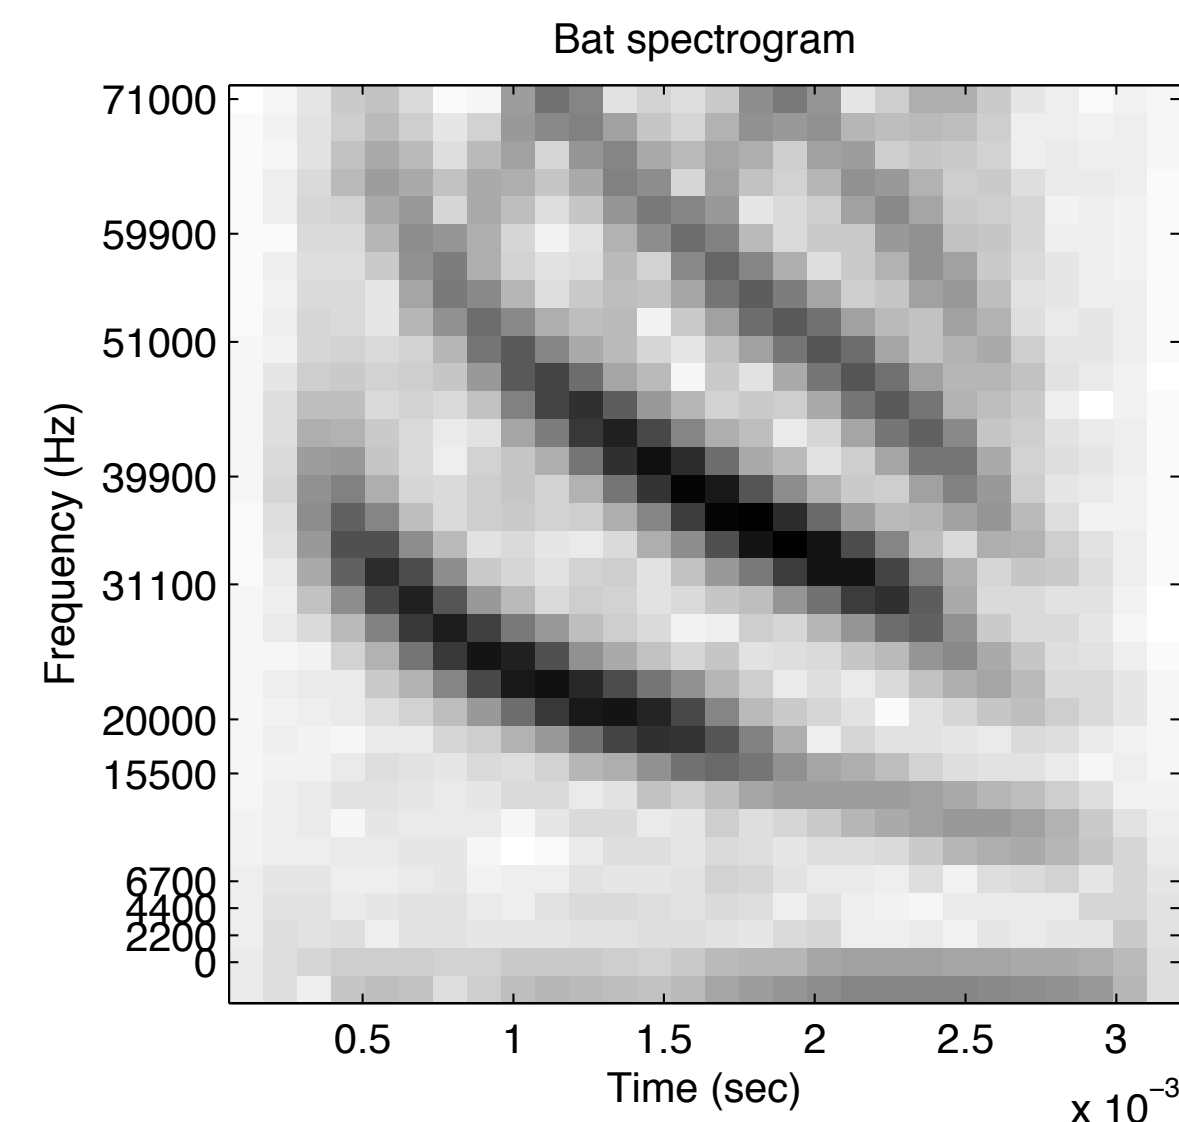
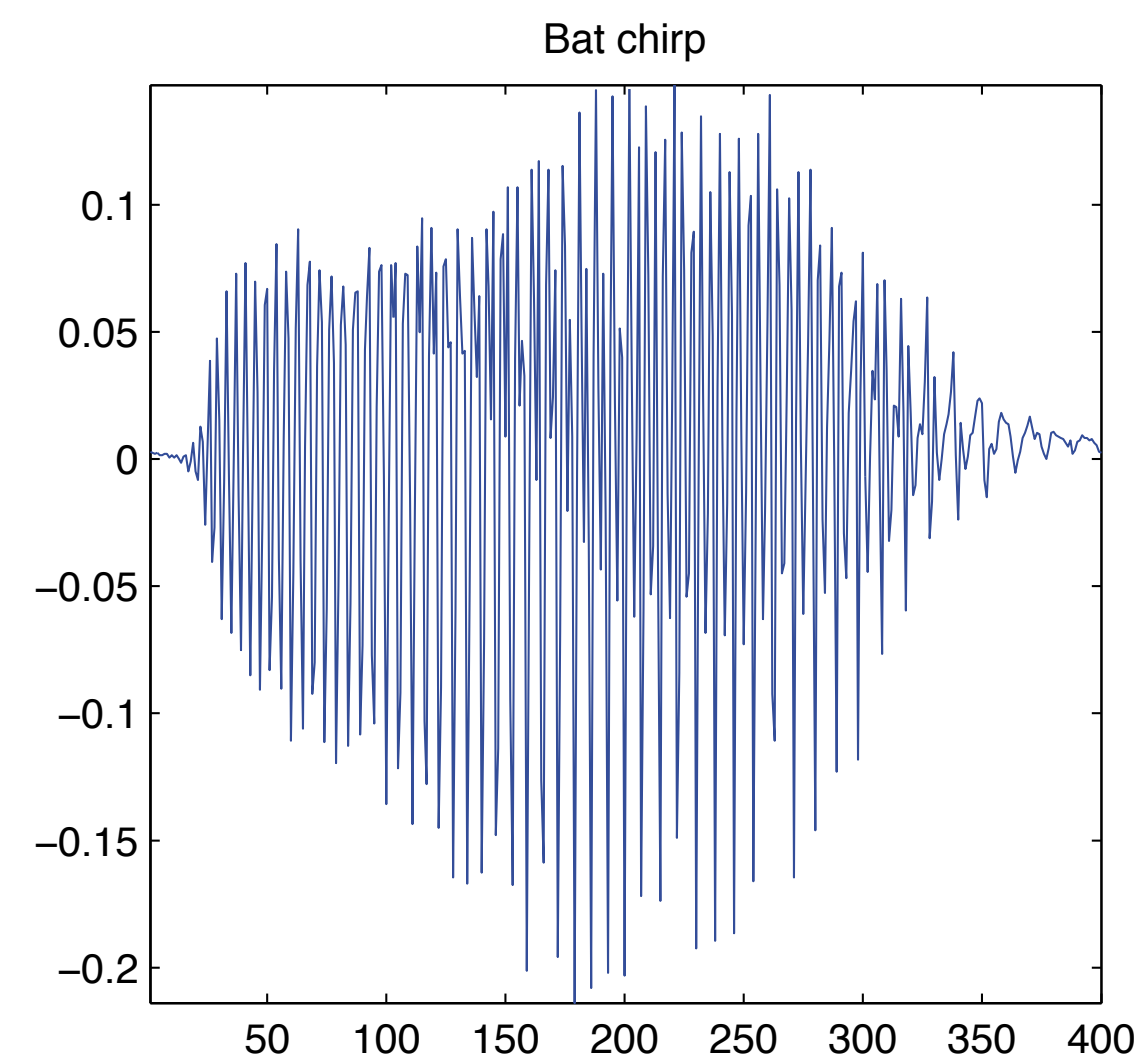
- According to the rules of sampling this is impossible!
- What is the magic taking place here?
 - Why is sparsity special?

Why sparsity?

- Sparsity implies structure, and structure is everywhere
- Signals often exhibit sparsity after undergoing the right transformation

Sparsity in signals


- Many signals are sparse in certain domains
 - e.g. sound spectra
 - Image wavelets
 - ...
- This is an opportunity to measure signals better
 - or to describe them easier




Vector spaces of signals

- Signals can be sparse in various ways
 - And some are “compressible”


1-sparse signal space



2-sparse signal space



Compressible signal space



Sparse approximations

- Represent signals using:

$$\mathbf{f} = \sum_k \overset{\text{Coefficients}}{a_k} \underset{\text{Bases / Dictionary}}{\mathbf{b}_k}$$


- Two goals:
 - Analysis: Study \mathbf{f} through structure of \mathbf{a} and \mathbf{b} (seen that)
 - Approximation: Reconstruct \mathbf{f} with a minimal number of terms

Exposing sparsity via dictionaries

- Can we use dictionaries that produce sparse coefficients?
- Why would they be useful and how would we implement them?


Example case: Wavelets

- Note how most coefficients are zero
 - A sparse representation



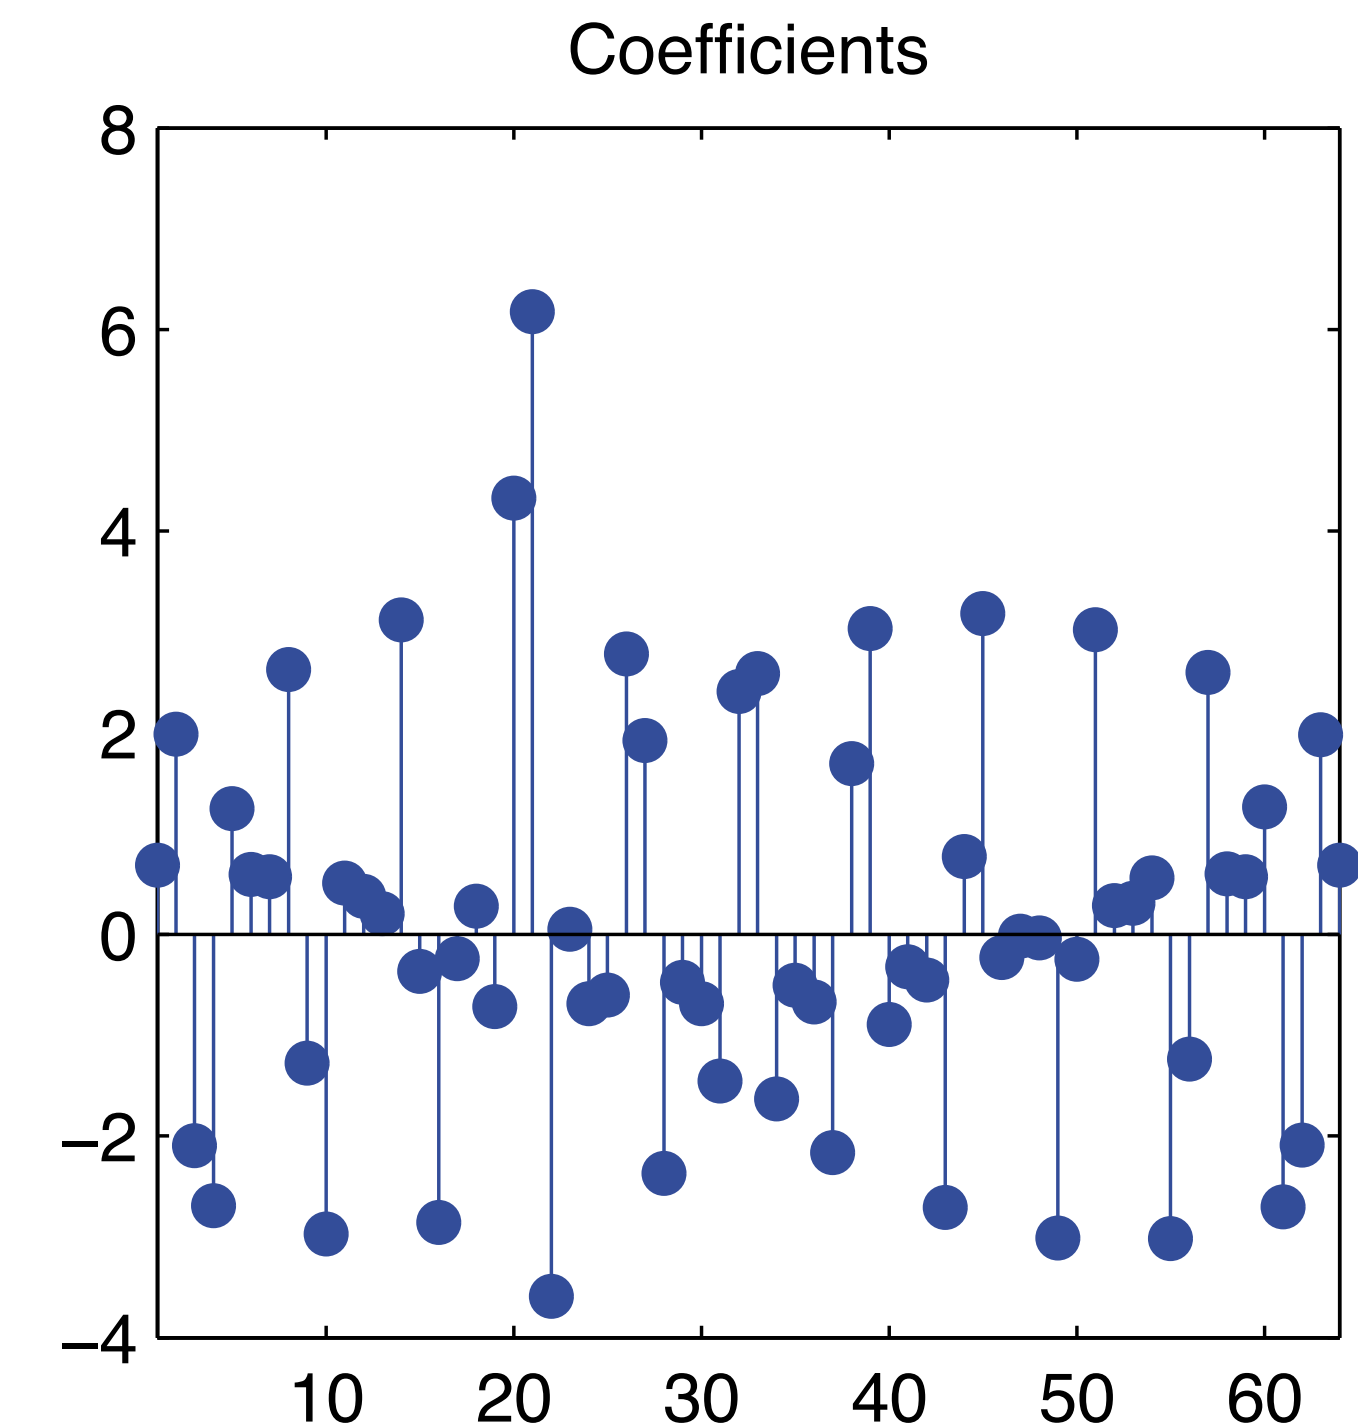

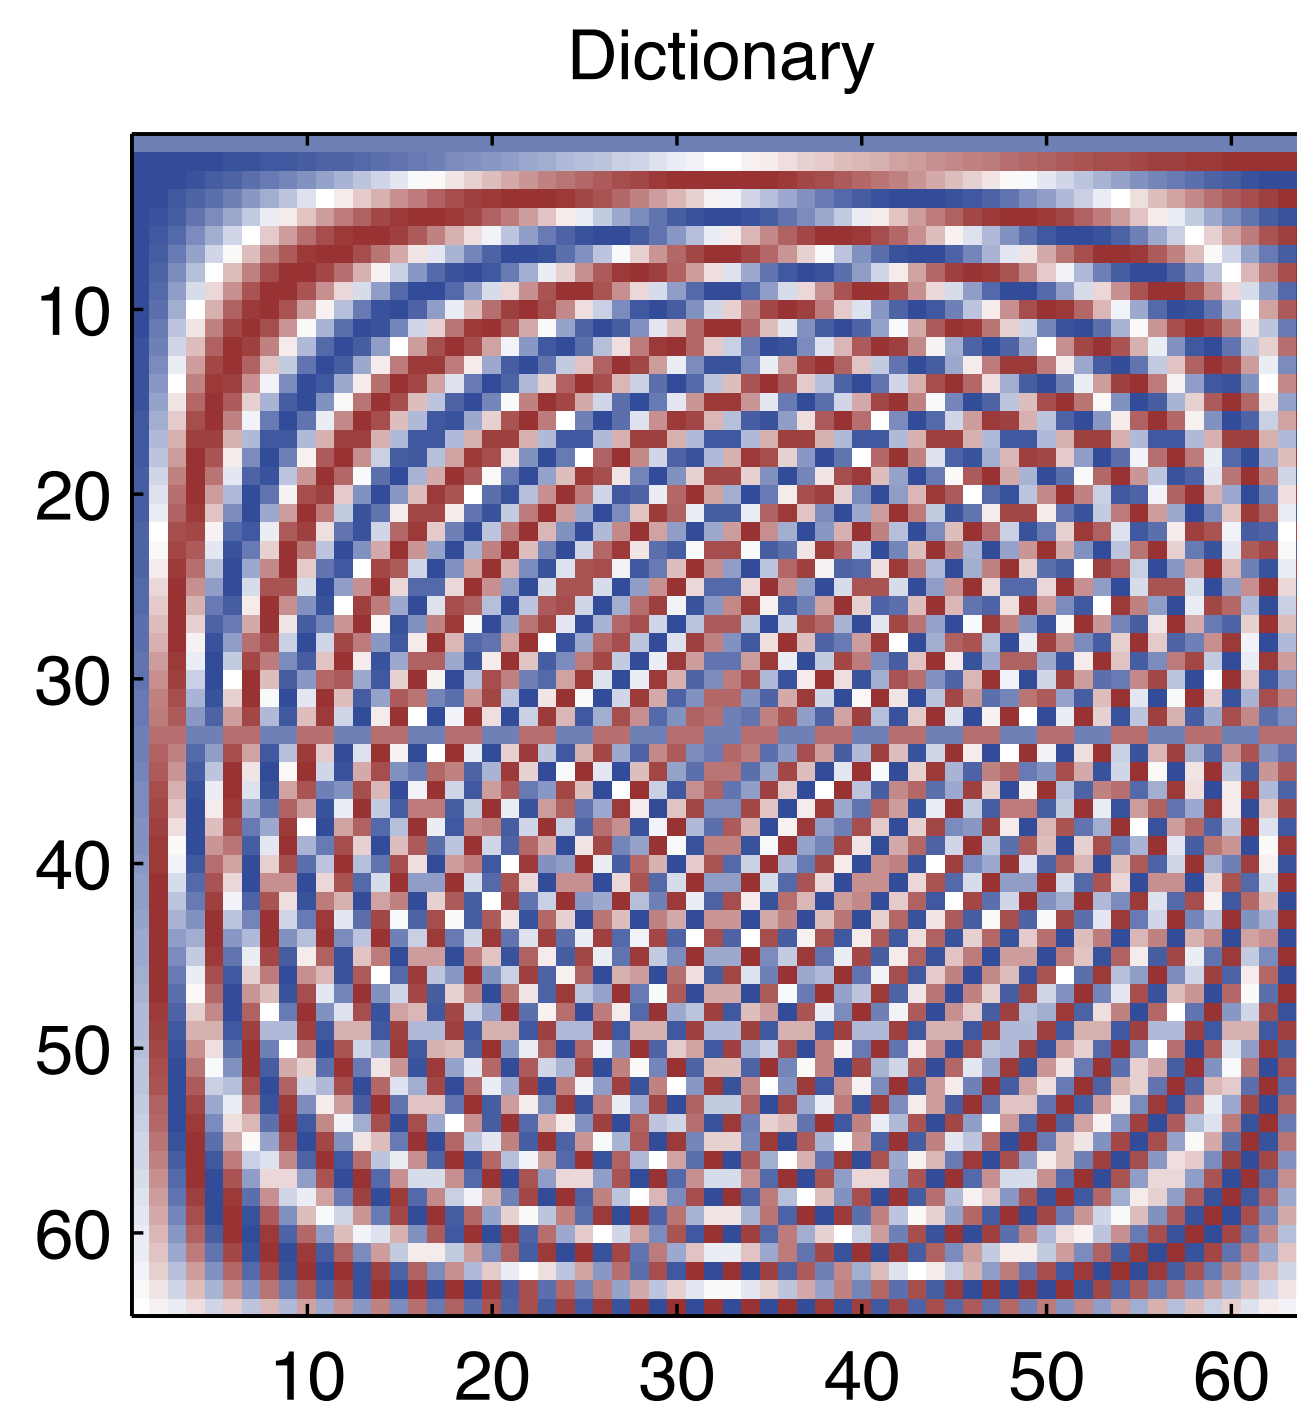
A simple example

- A sinusoid with a couple of spikes




Using a generic dictionary

- Analyzed via the DCT
 - Resulting coefficients are not sparse
 - Multiple sines are used to approximate the spikes






A “better” dictionary

- Use both sinusoids and spikes!
 - Now we won't use as many sines to represent the spikes






Applying the dictionary

- Using the straightforward decomposition won't help
 - Spike elements are not utilized
 - Minimal ℓ_2 cost penalizes the bases describing the loud spikes



Doing it the right way

- This time we ask for minimum ℓ_1 coefficients
 - And we get a perfect description of the input!



Overcomplete dictionaries


- Use dictionaries that contain “everything”!
 - Use compact descriptions of elements
- Some problems
 - Large size/computations
 - Lack of fast algorithms (e.g. FFT)
 - Problems with “coherence”
 - This is a big one

Dictionary coherence


- Make sure that the dictionary elements don't result in ambiguous coefficients

- A measure of coherence: $\mu = \max_{i,j} \left| \langle \mathbf{d}_i, \mathbf{d}_j \rangle \right|$

Good




Bad!




Examples of incoherent dictionaries


Fourier-like bases




Impulse bases



Gaussian noise bases



Binary noise bases




Getting greedy

- The linear programming approach will fail now
 - It is slow for large dictionaries
 - It looks for exact equality, not approximation
- We can instead use a greedy approach to resolving sparse approximations

Matching pursuit (MP)


- Family of many approaches based on successive fits
 - Each new fit explains what's the previous ones couldn't

Measure input
against dictionary $\langle \mathbf{d}_k, \mathbf{f} \rangle = \rho_k$




On a familiar example

- Each iteration knocks off an element
 - by 4th iteration there's nothing significant left to represent
 - Faster! For $N = 1024$, MP: 0.005 sec, LP: 63 sec



CoSaMP (Compressive Sensing MP)

- A useful variation to get K non-zero coefficients



Revisiting sampling

- Traditional acquisition samples uniformly
 - e.g. constant sample rates in audio, CCD grids in camera
- Foundation: Nyquist/Shannon sampling theory
 - Sample at twice the highest frequency
 - Projects to a complete basis that spans all the signal space

A redundancy in the loop

- Take a picture
 - Using dense sampling → lots of data
- Transform to a sparse domain and quantize
 - i.e. MPEG/JPEG compression → fewer data
- Process, transmit, view, etc.

Compressive sensing

- Why sample and then compress?
 - Do both at once!
- Sample fewer samples and use signal sparsity
 - Helps in finding a unique and plausible sparse signal

The compressive sensing pipeline

- Acquire signal using underdetermined measurements
 - e.g. linear combinations of a few samples: $\mathbf{y}_i = \mathbf{P}_i \cdot \mathbf{x}$
 - Don't sample densely, don't sample all the data
- Reconstruct signal assuming sparsity
 - Sparsity constraints signal space w.r.t. measurements
 - Allows for a plausible reconstruction

What's a good measurement matrix?

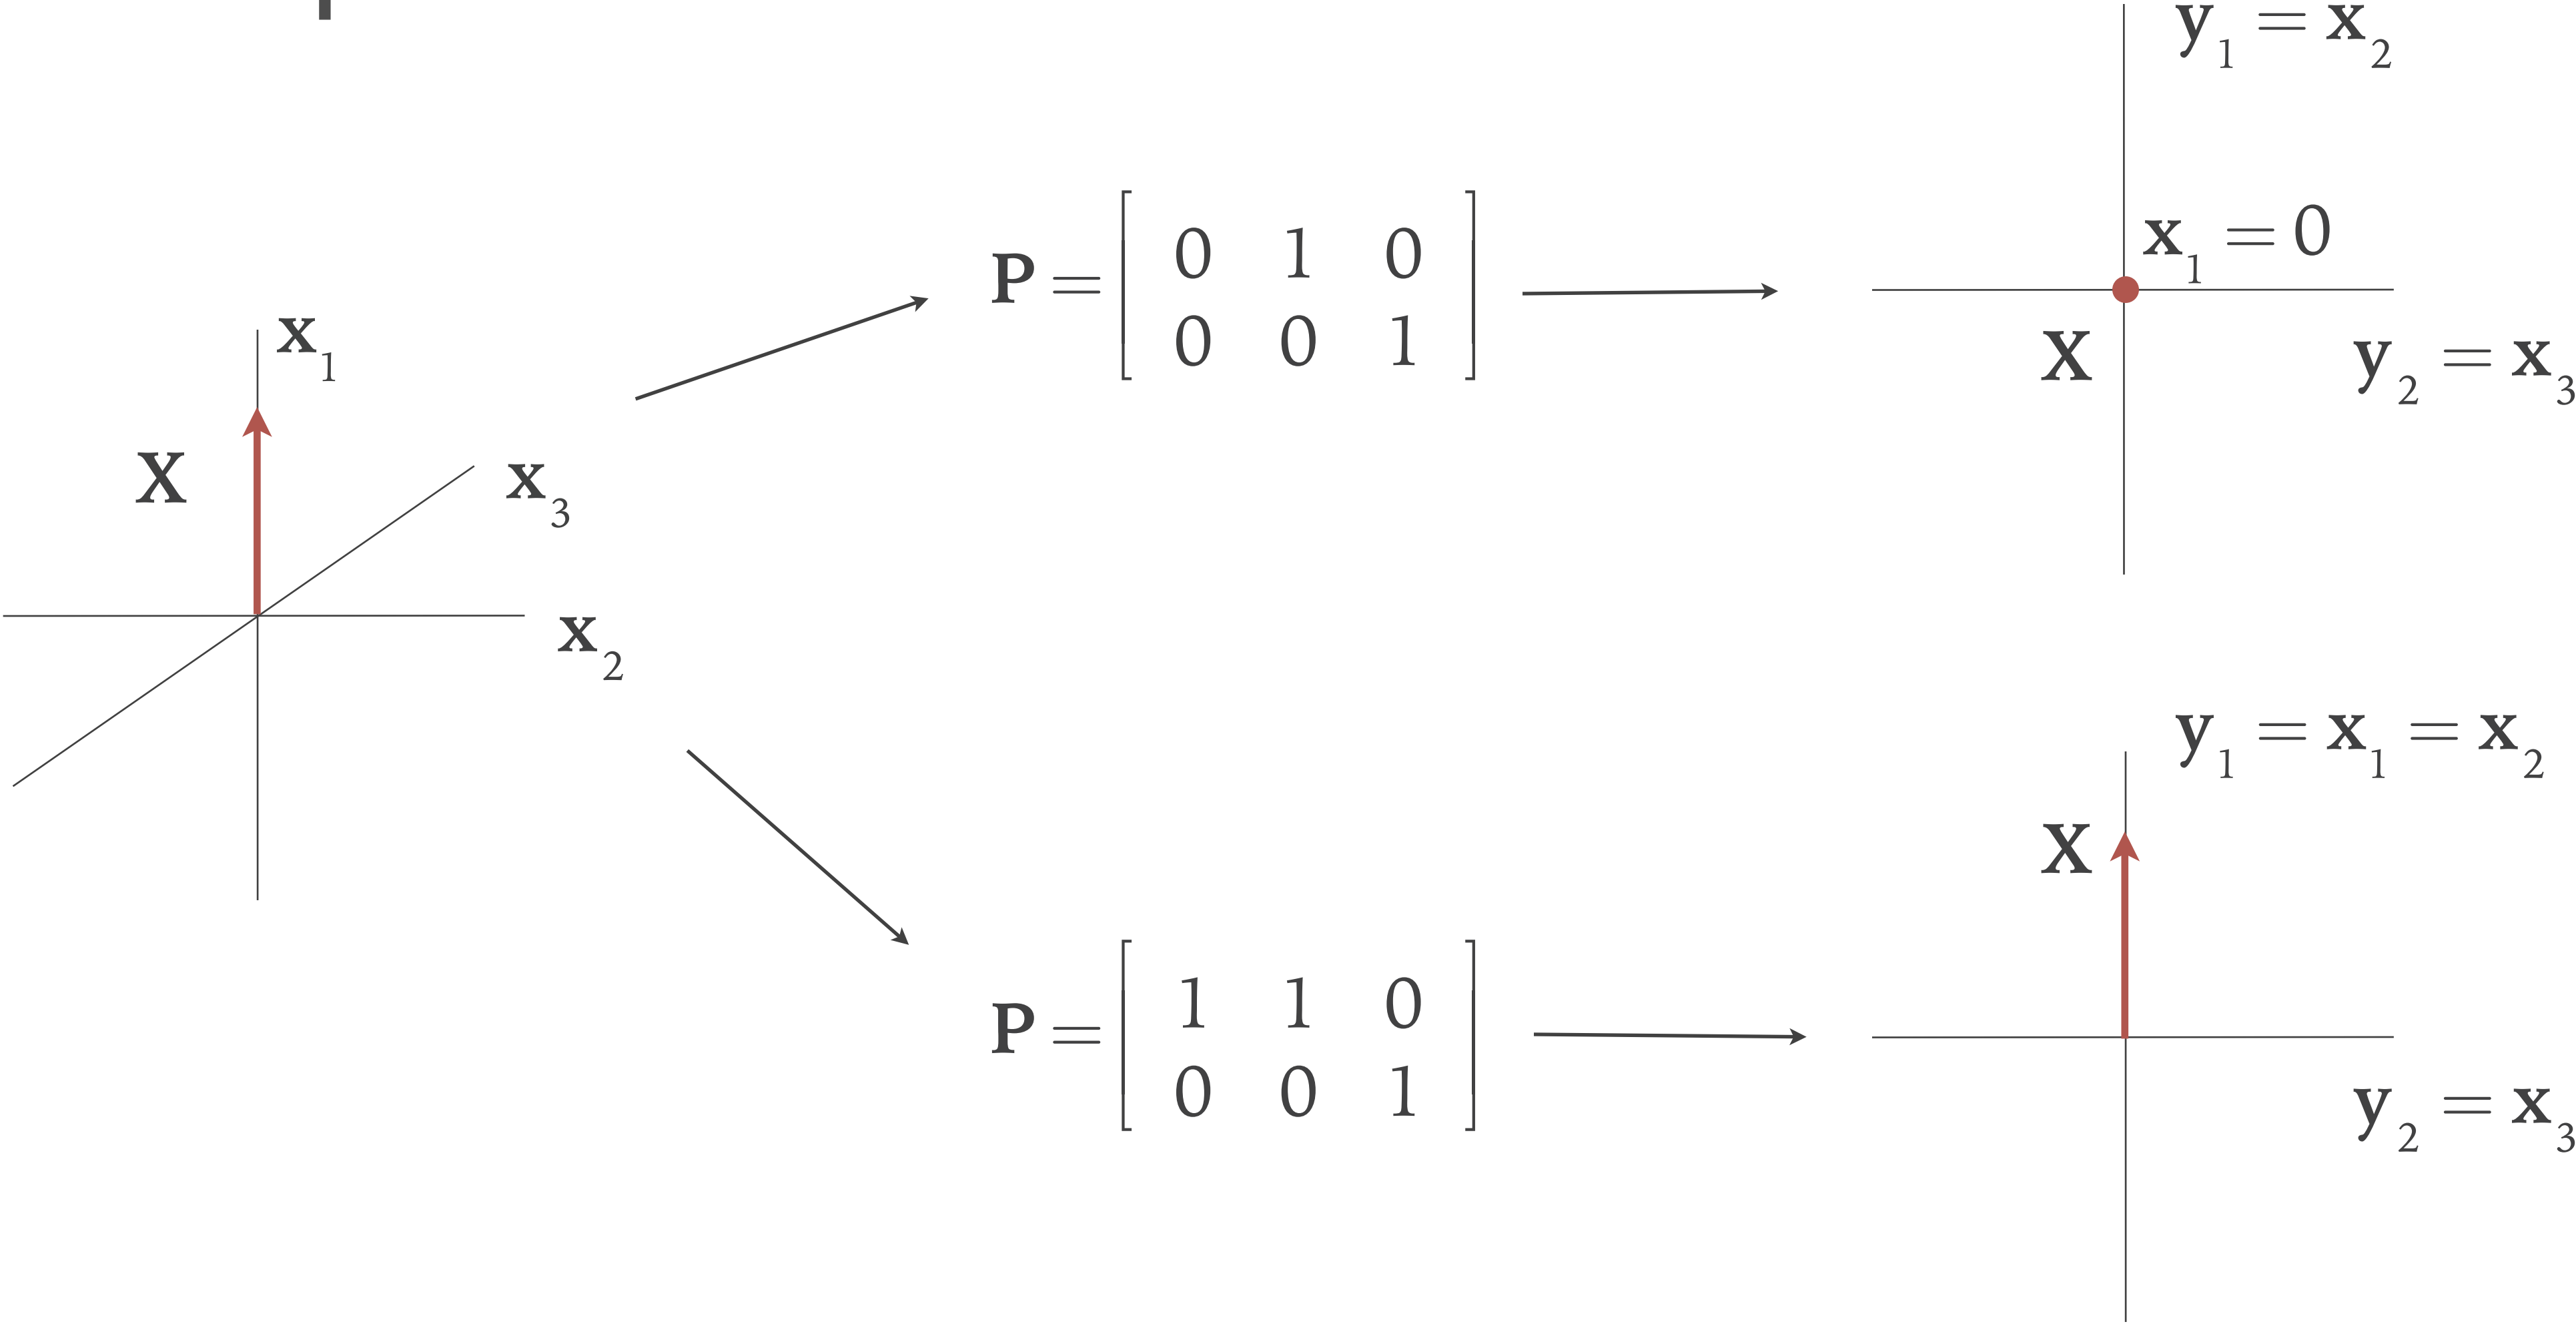
- We need:

$$\mathbf{P} \cdot \mathbf{x}_1 \neq \mathbf{P} \cdot \mathbf{x}_2 \text{ for all } K\text{-sparse } \mathbf{x}_1 \neq \mathbf{x}_2$$

- i.e. ensure that we can distinguish different inputs
- Necessary condition: \mathbf{P} must have at least $2K$ rows
 - Assuming noiseless and well-behaved data

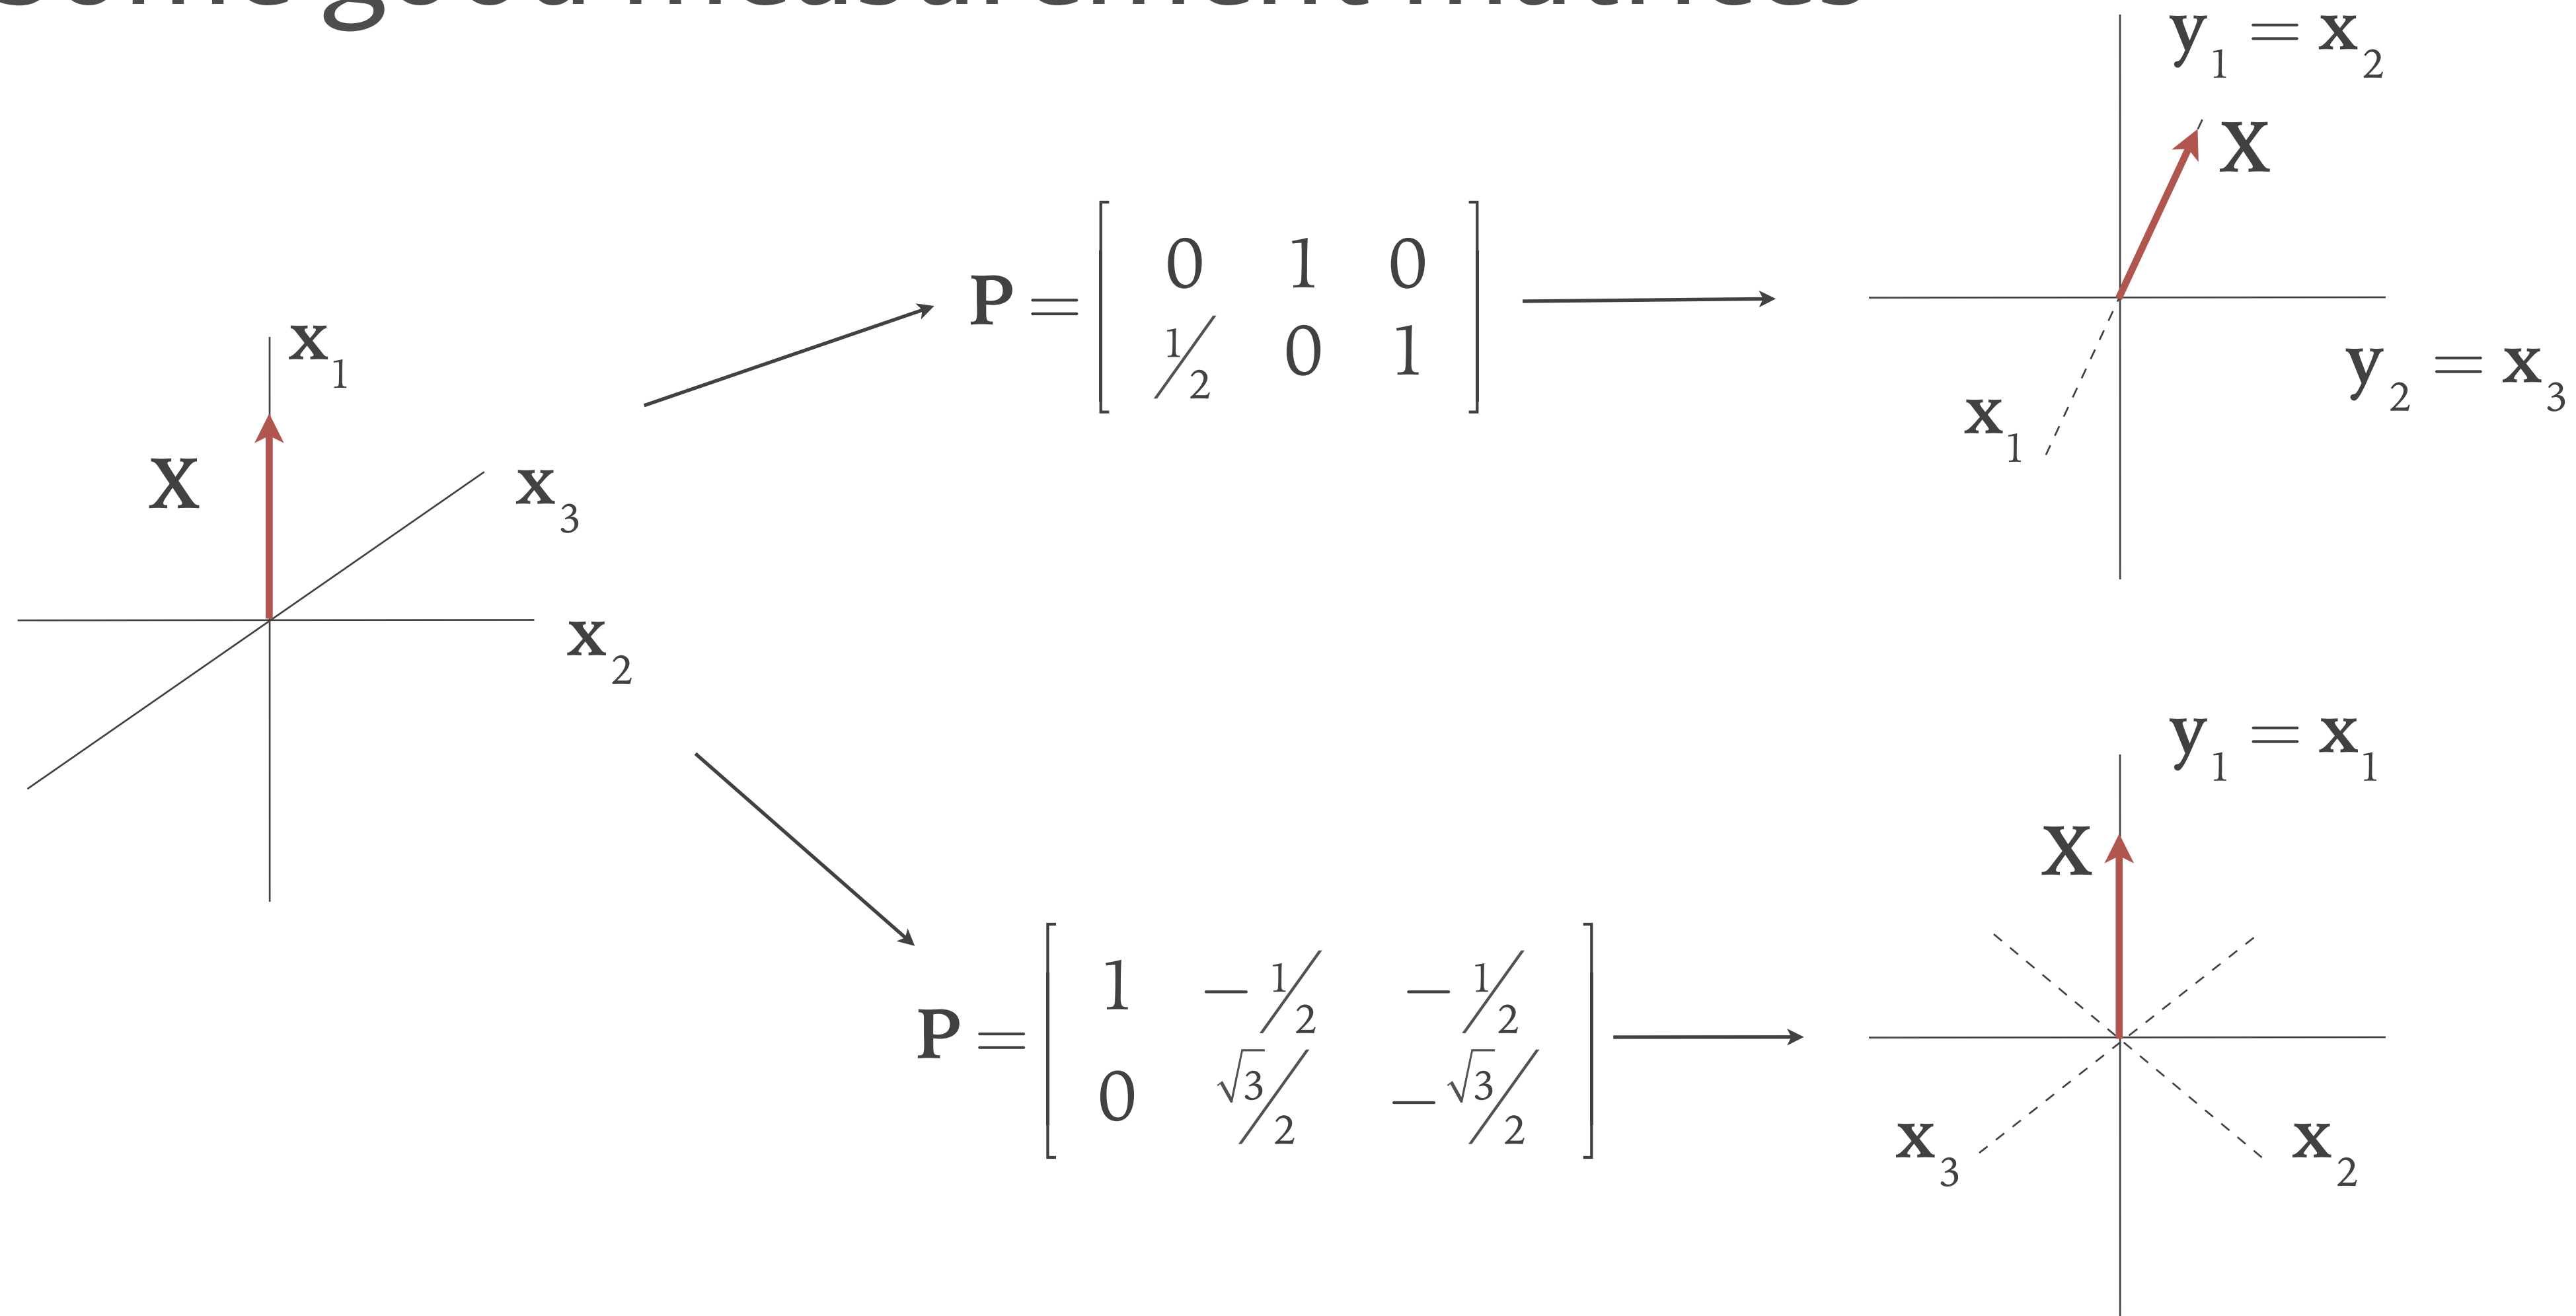
Example 1-sparse case

- Some poor measurement matrices



Example 1-sparse case

- Some good measurement matrices



Embedding viewpoint

- This is similar to the subspace/manifold methods
 - Find low-rank projection that preserves cluster characteristics
- Special case: Random projection
 - For n points in p -dims there exists a q -dim projection that preserves distances by a factor of $1 + \epsilon$, $q \geq O(\epsilon^{-2} \log n)$

Restricted Isometry Property (RIP)

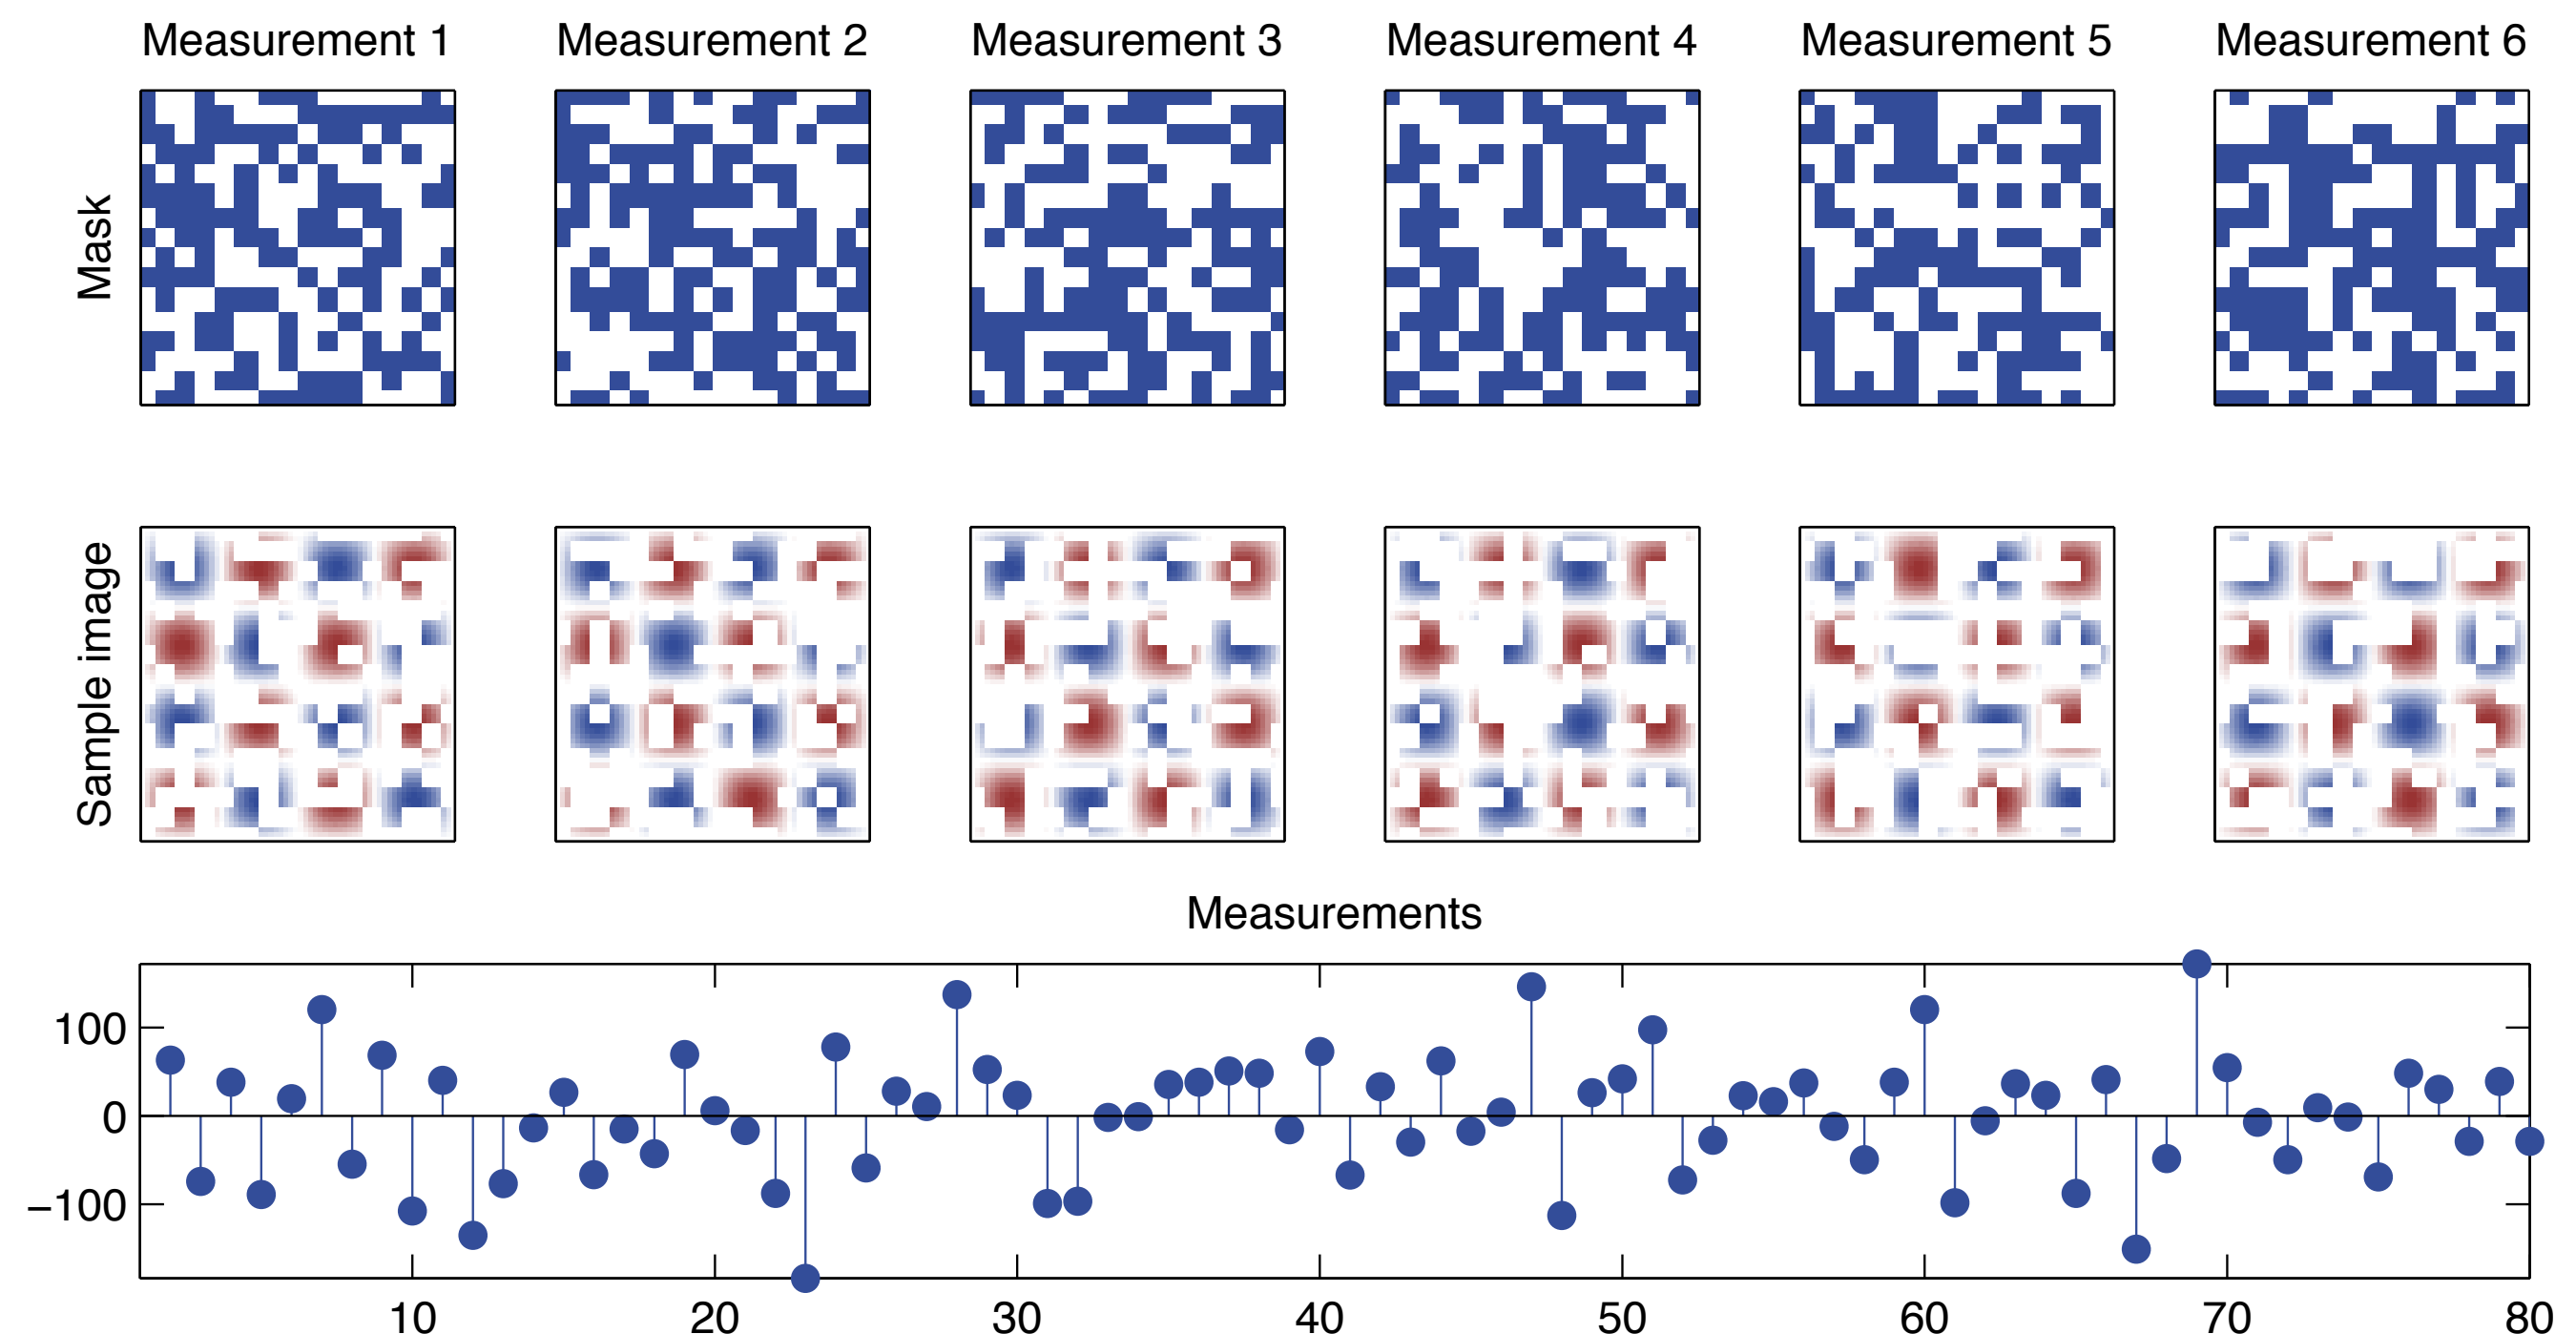

- A measurement matrix \mathbf{P} has order- K RIP if:

$$\left(1 - \delta_K\right) \leq \frac{\|\mathbf{P} \cdot \mathbf{x}\|_2^2}{\|\mathbf{x}\|_2^2} \leq \left(1 + \delta_K\right), \quad \forall \text{ } K\text{-sparse } \mathbf{x}$$

- How do we check? Difficult ...
 - But we have some good choices
 - Gaussian, Bernoulli, Fourier, etc.

An simple compressed sensing case

- Simple 64×64 input
 - Take a set of masked measurements, store only average value
 - Measure using: $y_i = \mathbf{p}_i^\top \cdot \text{vec}(\mathbf{x})$, $\mathbf{p}_i \in \{0,1\}$



Resolving via the DCT

- Model is:

$$\mathbf{y} = \mathbf{P}^\top \cdot \mathbf{x} = \mathbf{P}^\top \cdot (\mathbf{C}^\top \cdot \mathbf{z})$$

- Assume sparsity in \mathbf{z}

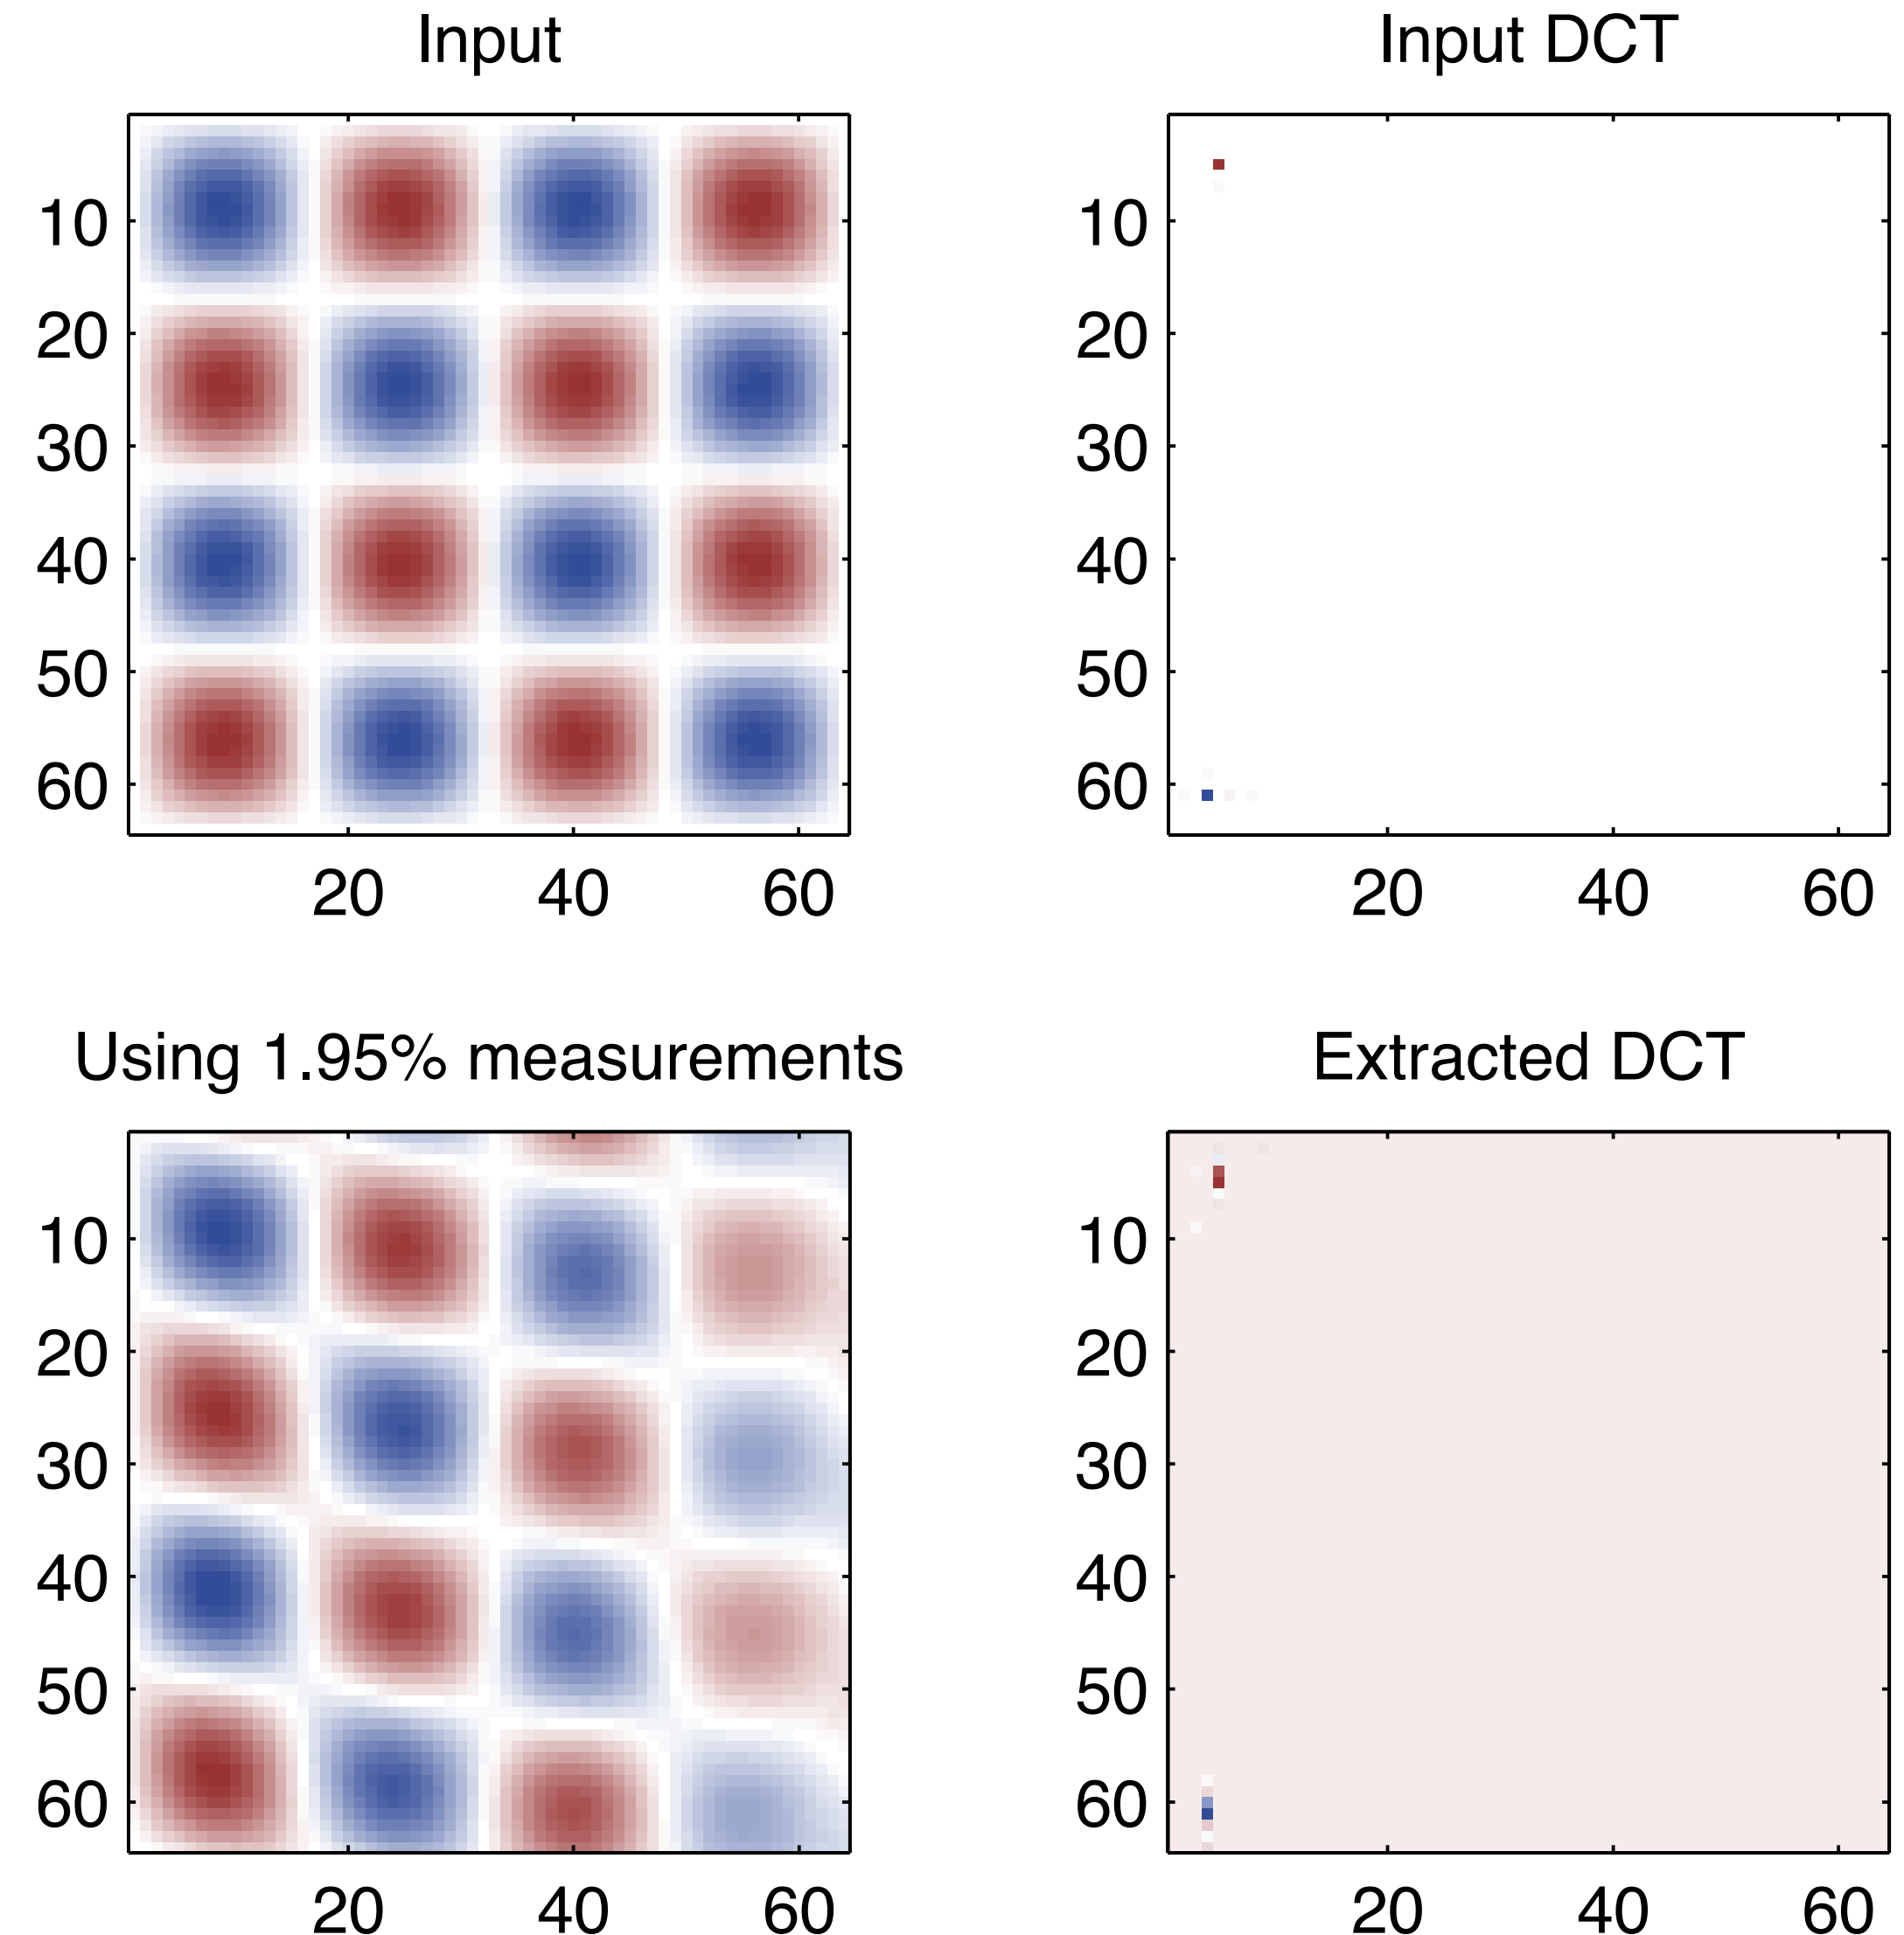
- Resolve:

$$\min \|\mathbf{z}\|_1 + \|\mathbf{y} - \mathbf{P}^\top \cdot \mathbf{C}^\top \cdot \mathbf{z}\|$$

- Reconstruct:

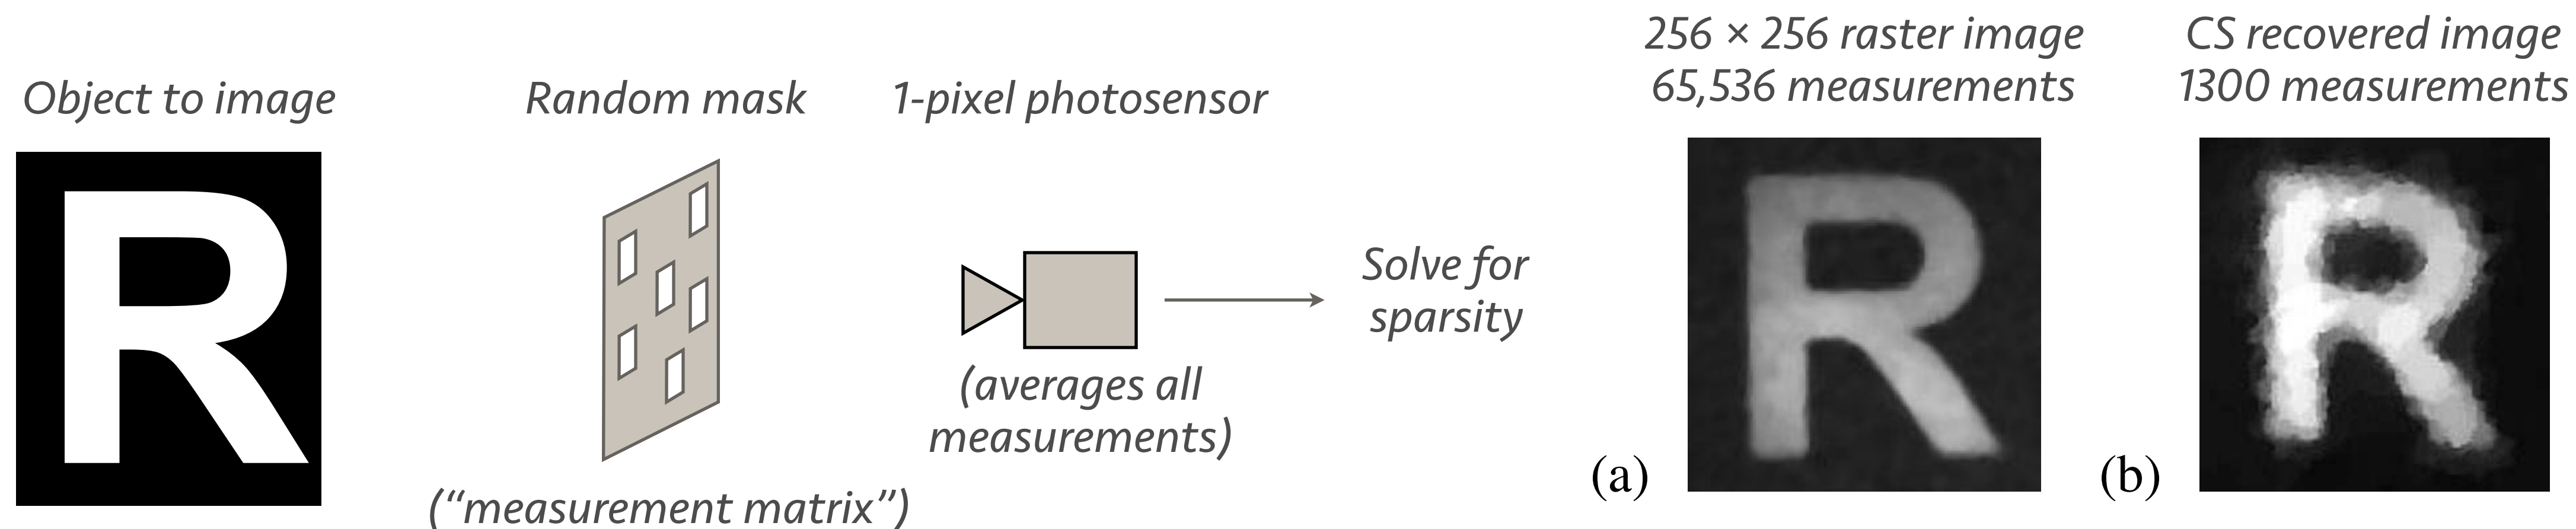
$$\hat{\mathbf{x}} = \mathbf{C}^\top \cdot \mathbf{z}$$

- Using ~2% of samples!



Putting compressive sensing to work

- The single-pixel camera
 - Measure image intensity using multiple random masks
 - Reconstruct assuming sparsity
 - In some domain ...




Other kinds of sparsity

- Up to now we talked about a simple form of sparsity
 - Sparsity over all coefficients separately
- Some problems need more elaborate definitions
 - e.g. block structure, joint structure, temporal structure, etc.


Block sparsity

- Have sparsity appear in non-overlapping blocks
 - Useful for some imaging operations



Joint sparsity

- Obtain multiple solutions with the same sparsity
 - Useful for multimodal/multichannel data



Learning and sparsity

- How about other models?
 - Add sparsity as an extra “regularizer”
- Sparse decompositions
 - Sparse NMF, sparse PCA, etc ...
- ICA is already (sort of) sparse!

Recap

- Sparsity and ℓ_p norms
 - Different definition of sparsity
- Minimum- ℓ_1 coefficient algorithms
 - Linear programming
 - Greedy methods
- Compressive sensing and random projections
 - 1-pixel camera

Reading material

- Compressive sensing
 - <http://dsp.rice.edu/sites/dsp.rice.edu/files/cs/baraniukCSlecture07.pdf>
- Compressive Sensing page
 - <http://dsp.rice.edu/cs>
- Experiments with Random Projections
 - <http://dimacs.rutgers.edu/Research/MMS/PAPERS/rp.pdf>

Next lecture

- Deep learning!
 - aka neural nets v2.0
- Also Problem Set 4 is out
 - Optional, due at last day of classes
 - Use it to perk up your grade if you need to
 - Would also help if your final project is floundering